

18-447 Lecture 8:

Data Hazard and Resolution

James C. Hoe

Department of ECE

Carnegie Mellon University

Housekeeping

- Your goal today
 - detect and resolve data hazards in in-order instruction pipelines
 - control dependence next time
- Notices
 - HW 2, due Mon 2/21
 - Lab 2, status check wk6, due wk7 (Handout #7)
 - HW 3, due Mon 2/28 (Handout #8)
- Readings
 - P&H Ch 4

Instruction Pipeline Reality

- Not identical tasks
 - coalescing instruction types into one “multi-function” pipe
 - external fragmentation (some idle stages)
- Not uniform suboperations
 - group or sub-divide steps into stages to minimize variance
 - internal fragmentation (some too-fast stages)
- Not independent tasks
 - dependency detection and resolution
 - next lecture(s)

Recall

Data Dependence

Data dependence

x3 ← x1 op x2
 ...
 x5 ← **x3** op x4

Read-after-Write (RAW)

Anti-dependence

x3 ← **x1** op x2
 ...
x1 ← x4 op x5

Write-after-Read (WAR)

Output-dependence

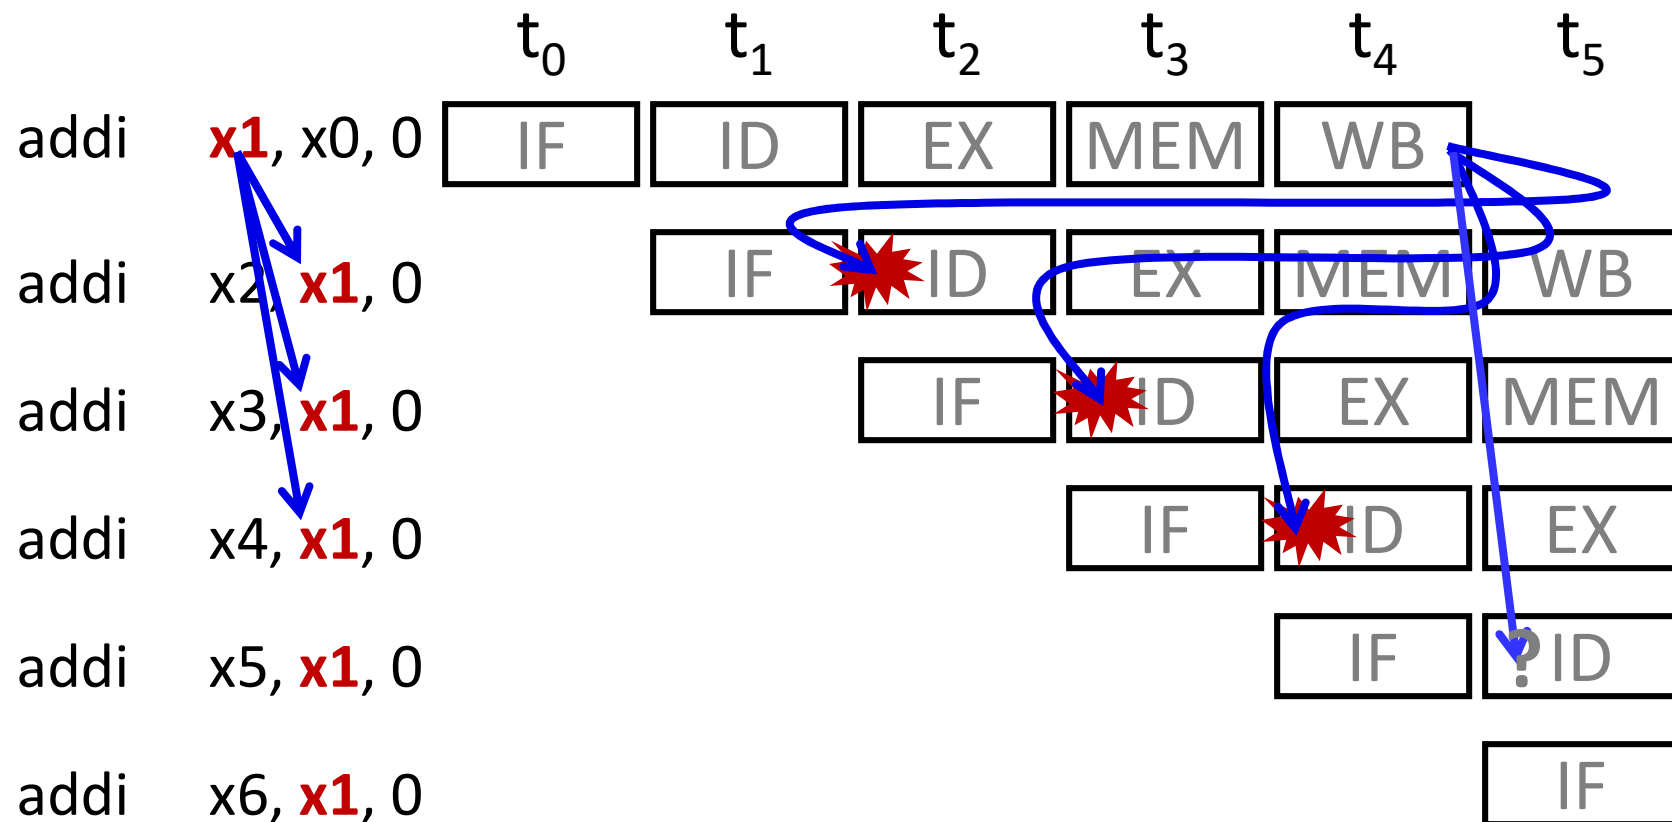
x3 ← x1 op x2
 ...
x3 ← x6 op x7

Write-after-Write (WAW)

false dependence

Don't forget memory instructions

Dependency vs Hazard: e.g. RAW



Dependence is property of program;
hazards specific to microarchitecture

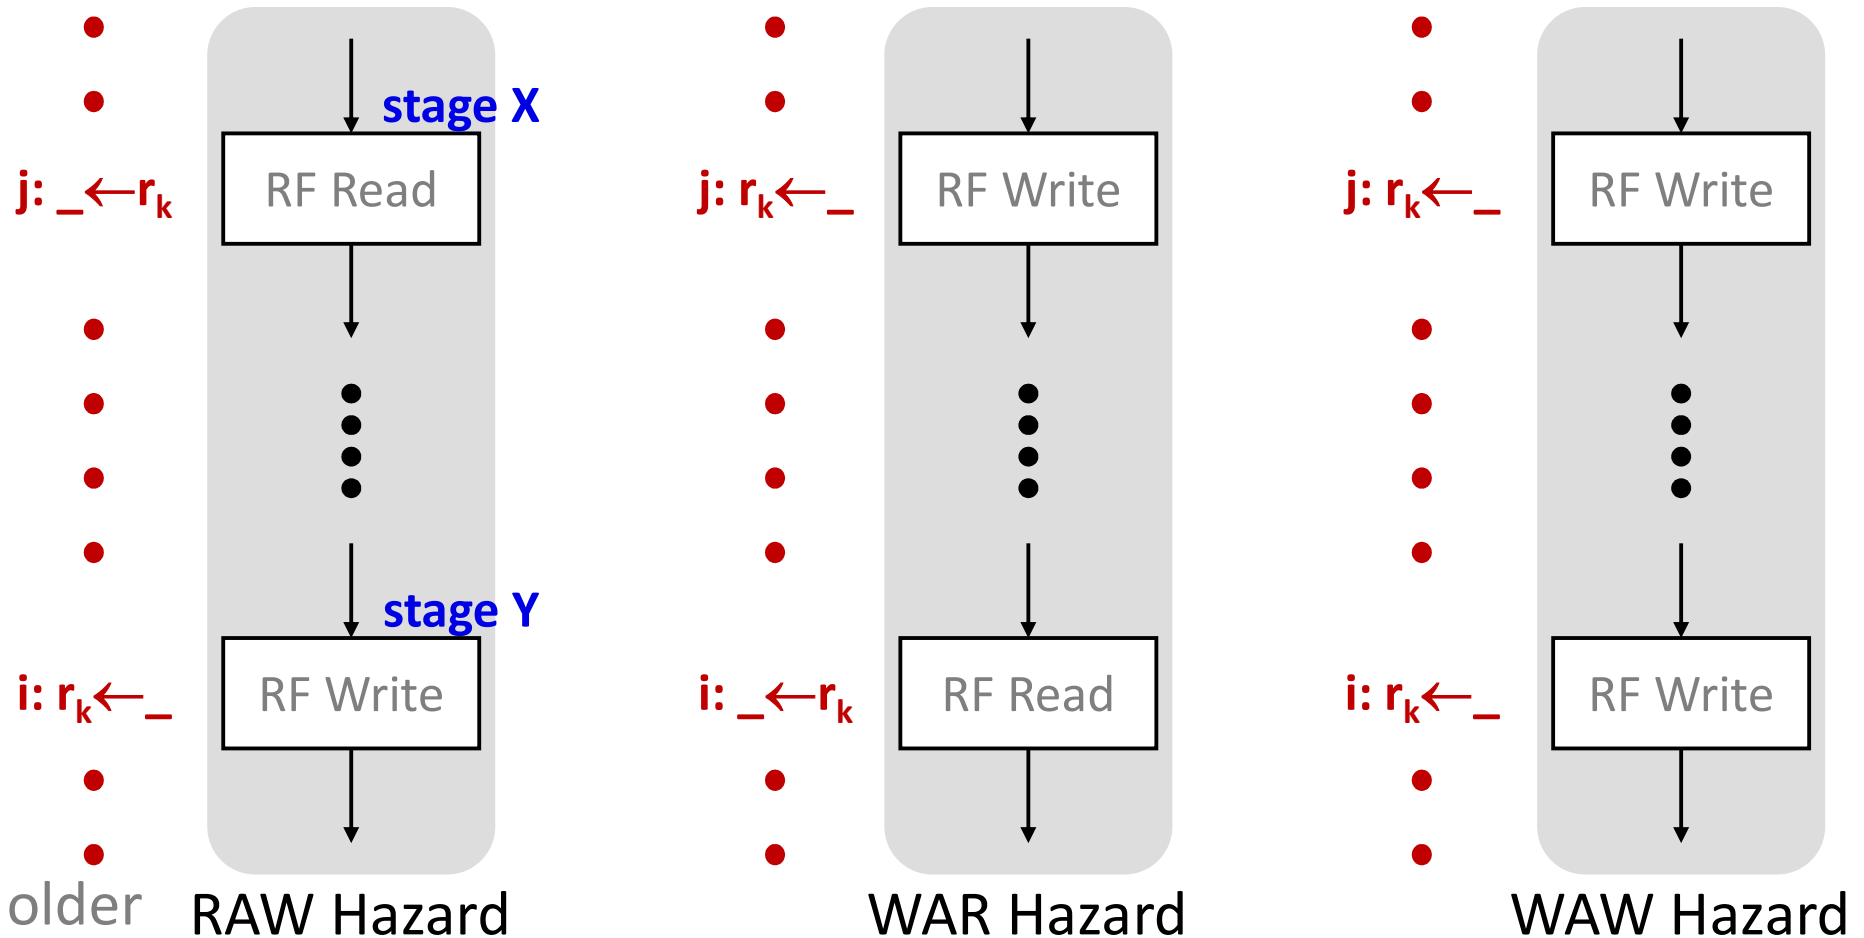
Register Data Hazard Analysis

	R/I-Type	LW	SW	Bxx	Jal	Jalr
IF						
ID	read RF	read RF	read RF	read RF		read RF
EX						
MEM						
WB	write RF	write RF			write RF	write RF

- For a given pipeline, when is there a register data hazard between 2 dependent instructions?
 - dependence type: RAW, WAR, WAW?
 - instruction types involved?
 - distance between the two instructions?

Hazard in In-order Pipeline

younger



$$\text{dist}_{\text{dependence}}(i, j) \leq \text{dist}_{\text{hazard}}(X, Y) \Rightarrow \text{Hazard!!}$$

$$\text{dist}_{\text{dependence}}(i, j) > \text{dist}_{\text{hazard}}(X, Y) \Rightarrow \text{Safe}$$

RAW Hazard Analysis Example


	R/I-Type	LW	SW	Bxx	Jal	Jalr
IF						
ID	read RF	read RF	read RF	read RF		read RF
EX						
MEM						
WB	write RF	write RF			write RF	write RF

- Older I_A and younger I_B have RAW hazard iff
 - I_B (R/I, LW, SW, Bxx or JALR) reads a register written by I_A (R/I, LW, or JAL/R)
 - $\text{dist}(I_A, I_B) \leq \text{dist}(\text{ID}, \text{WB}) = 3$

What about WAW and WAR hazard?

What about memory data hazard?

Pipeline Stall: universal hazard resolution



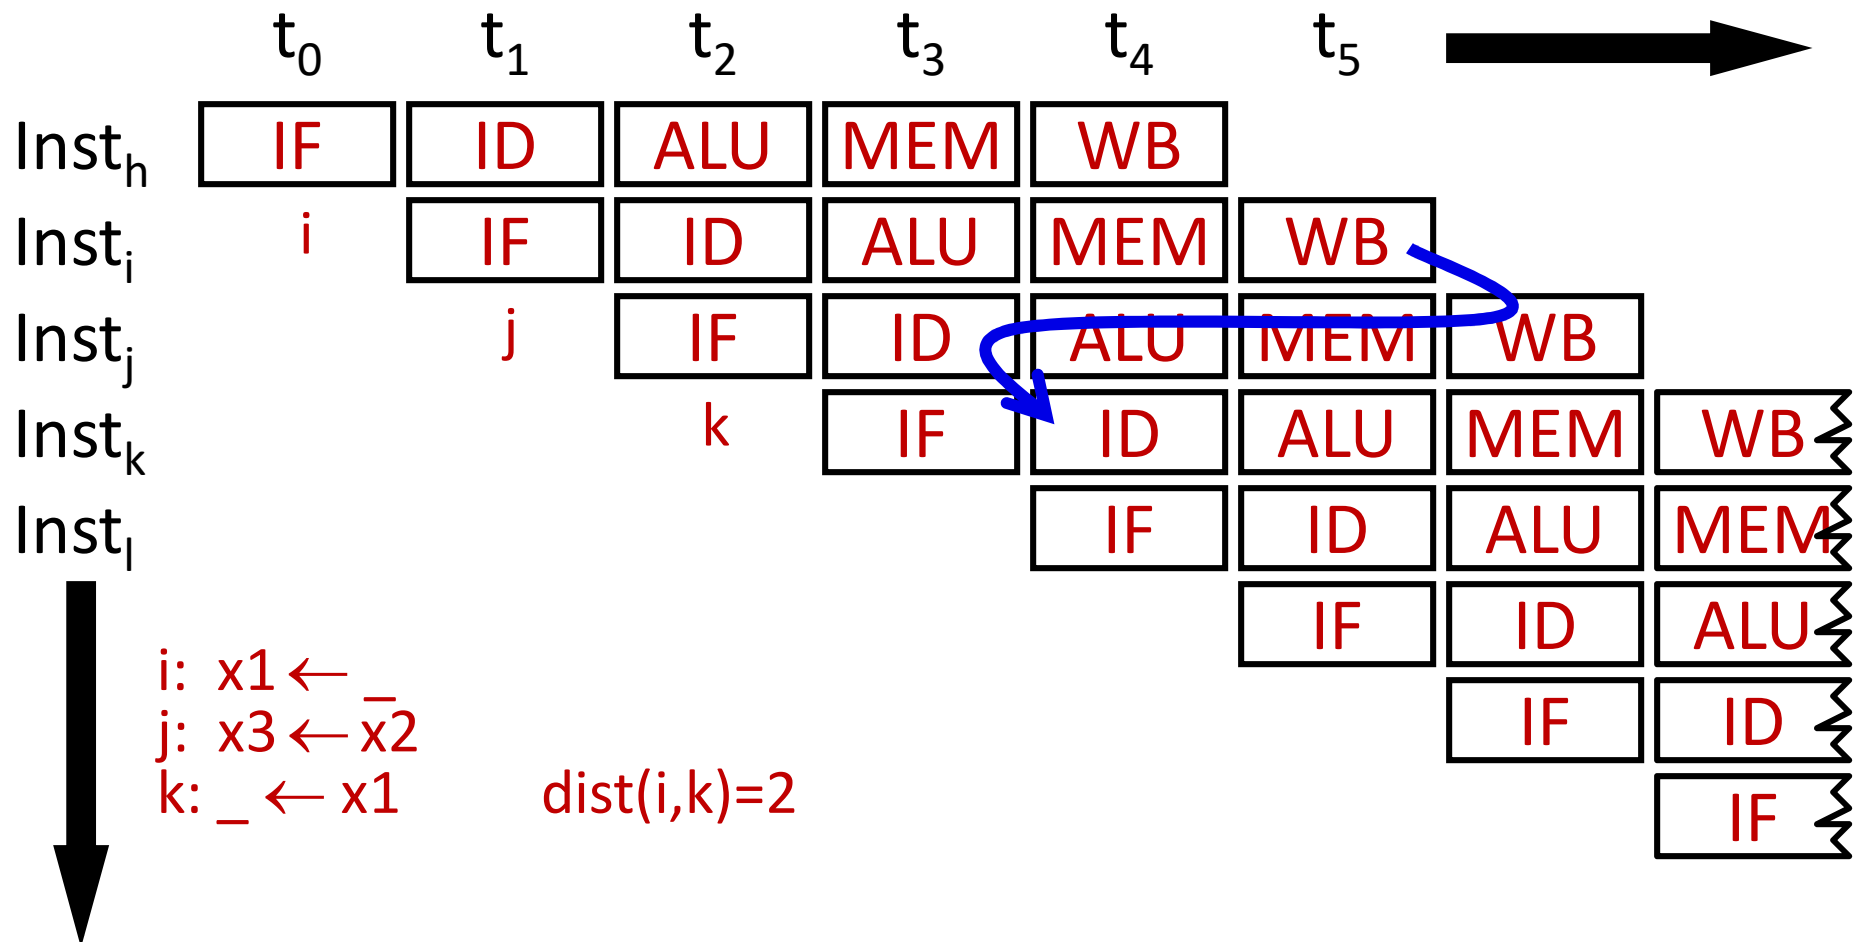
```
i: x1 ← bubble
bubble
bubble
j: _ ← x1
```

$$\text{dist}(i,j)=4$$

Stall==make younger instruction
wait until hazard passes

1. stop all up-stream stages
2. drain all down-stream stages

Pop Quiz: What happens in this case?

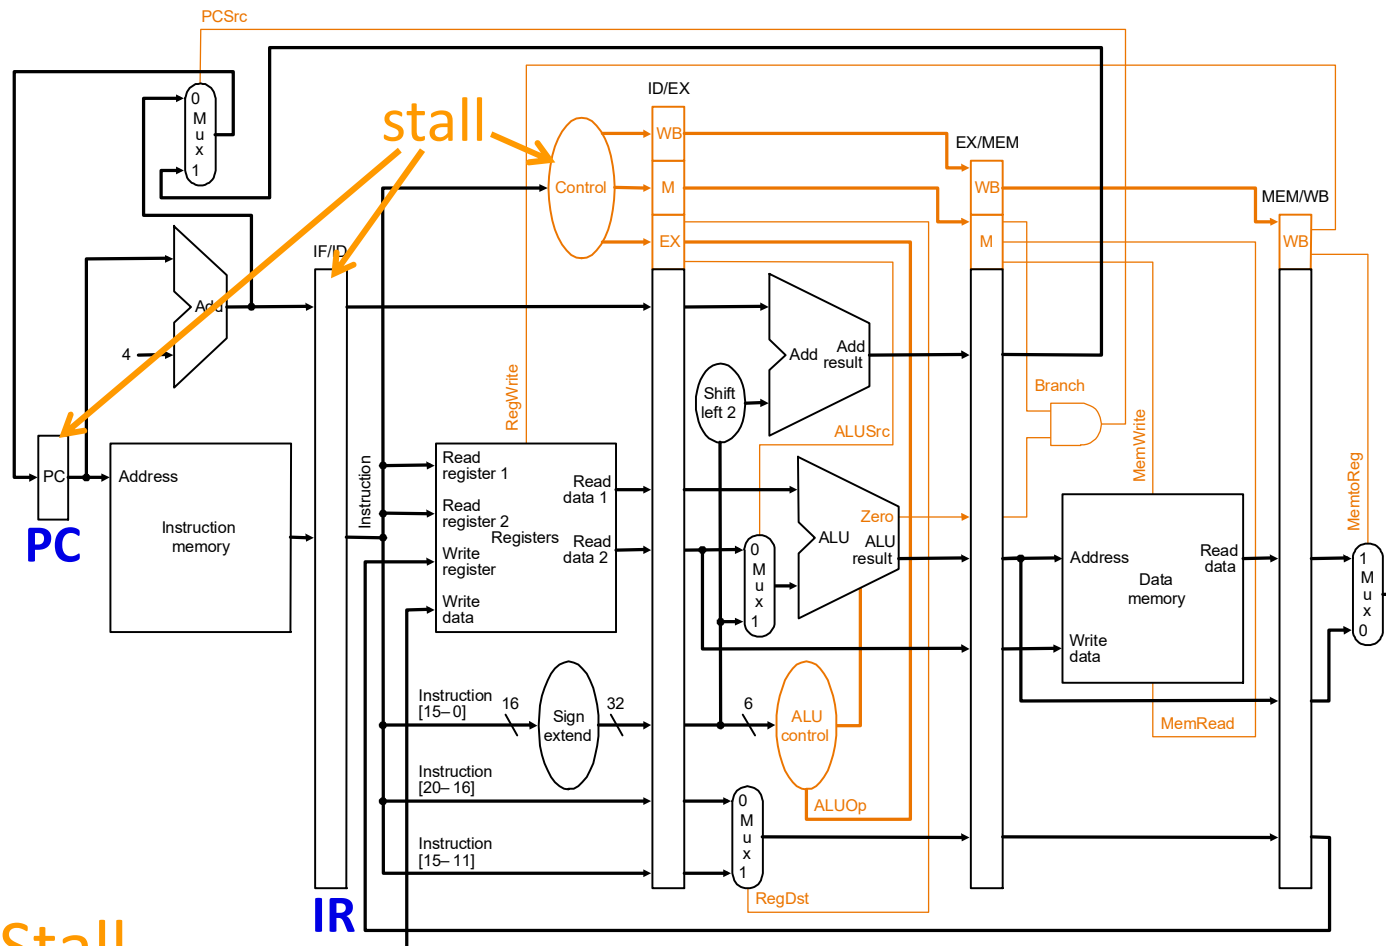


Pipeline Stall

	t ₁	t ₂	t ₃	t ₄	t ₅	t ₆	t ₇	t ₈	t ₉	t ₁₀	t ₁₁
IF	i	j	k	k	k	k	l				
ID	h	i	j	j	j	j	k	l			
EX		h	i	bub	bub	bub	j	k	l		
MEM			h	i	bub	bub	bub	j	k	l	
WB				h	i	bub	bub	bub	j	k	l

i: x1 ← _

j: _ ← x1



- Stall
 - disable **PC** and **IR** latching
 - set $\text{RegWrite}_{\text{ID}}=0$ and $\text{MemWrite}_{\text{ID}}=0$

When to Stall

- Older I_A and younger I_B have RAW hazard iff
 - I_B (R/I, LW, SW, Bxx or JALR) reads a register written by I_A (R/I, LW, or JAL/R)
 - $\text{dist}(I_A, I_B) \leq \text{dist}(\text{ID}, \text{WB}) = 3$

Above is about existence of hazard

- Operationally, to detect hazard in time to prevent:
 - before I_B in ID reads a register, I_B needs to check if any I_A in EX, MEM or WB is going to update it

(if so, value in RF is “stale”)

Watch out for x0!!

Stall Condition

- Helper function
 - *waitForRs1*(*l*) returns true if *l* uses *rs1* && *rs1*!=x0
- Stall IF and ID when
 - $(rs1_{ID} == rd_{EX}) \ \&\& \text{RegWrite}_{EX} \ \&\& \text{waitForRs1}(IR_{ID})$ or
 - $(rs1_{ID} == rd_{MEM}) \ \&\& \text{RegWrite}_{MEM} \ \&\& \text{waitForRs1}(IR_{ID})$ or
 - $(rs1_{ID} == rd_{WB}) \ \&\& \text{RegWrite}_{WB} \ \&\& \text{waitForRs1}(IR_{ID})$ or
 - $(rs2_{ID} == rd_{EX}) \ \&\& \text{RegWrite}_{EX} \ \&\& \text{waitForRs2}(IR_{ID})$ or
 - $(rs2_{ID} == rd_{MEM}) \ \&\& \text{RegWrite}_{MEM} \ \&\& \text{waitForRs2}(IR_{ID})$ or
 - $(rs2_{ID} == rd_{WB}) \ \&\& \text{RegWrite}_{WB} \ \&\& \text{waitForRs2}(IR_{ID})$

It is crucial that EX, MEM and WB
continue to advance during stall

Impact of Stall on Performance

- Each stall cycle corresponds to 1 lost ALU cycle
- A program with N instructions and S stall cycles:

$$\text{average IPC} = N / (N + S)$$

- S depends on
 - frequency of hazard-causing dependencies
 - distance between hazard-causing instruction pairs
 - distance between hazard-causing dependencies

(suppose i_1, i_2 and i_3 all depend on i_0 , once i_1 's hazard is resolved by stalling, i_2 and i_3 do not stall)

Sample Assembly [P&H]


for (j=i-1; j>=0 && v[j] > v[j+1]; j-=1) { }

	addi	\$s1, \$s0, -1	3 stalls
for2tst:	slti	\$t0, \$s1, 0	3 stalls
	bne	\$t0, \$zero, exit2	
	sll	\$t1, \$s1, 2	3 stalls
	add	\$t2, \$a0, \$t1	3 stalls
	lw	\$t3, 0(\$t2)	
	lw	\$t4, 4(\$t2)	3 stalls
	slt	\$t0, \$t4, \$t3	3 stalls
	beq	\$t0, \$zero, exit2	
		
	addi	\$s1, \$s1, -1	
	j	for2tst	

exit2:

Data Forwarding (or Register Bypassing)

- What does “**ADD** r_x r_y r_z ” mean? Get inputs from $RF[r_y]$ and $RF[r_z]$ and put result in $RF[r_x]$?
- But, RF is just a part of an abstraction
 - a way to connect dataflow between instructions
 - “operands to **ADD** are resulting values of the last instructions to assign to $RF[r_y]$ and $RF[r_z]$ ”
 - RF doesn't have to exist/behave as a literal object!!!
- If only dataflow matters, don't wait for WB . . .

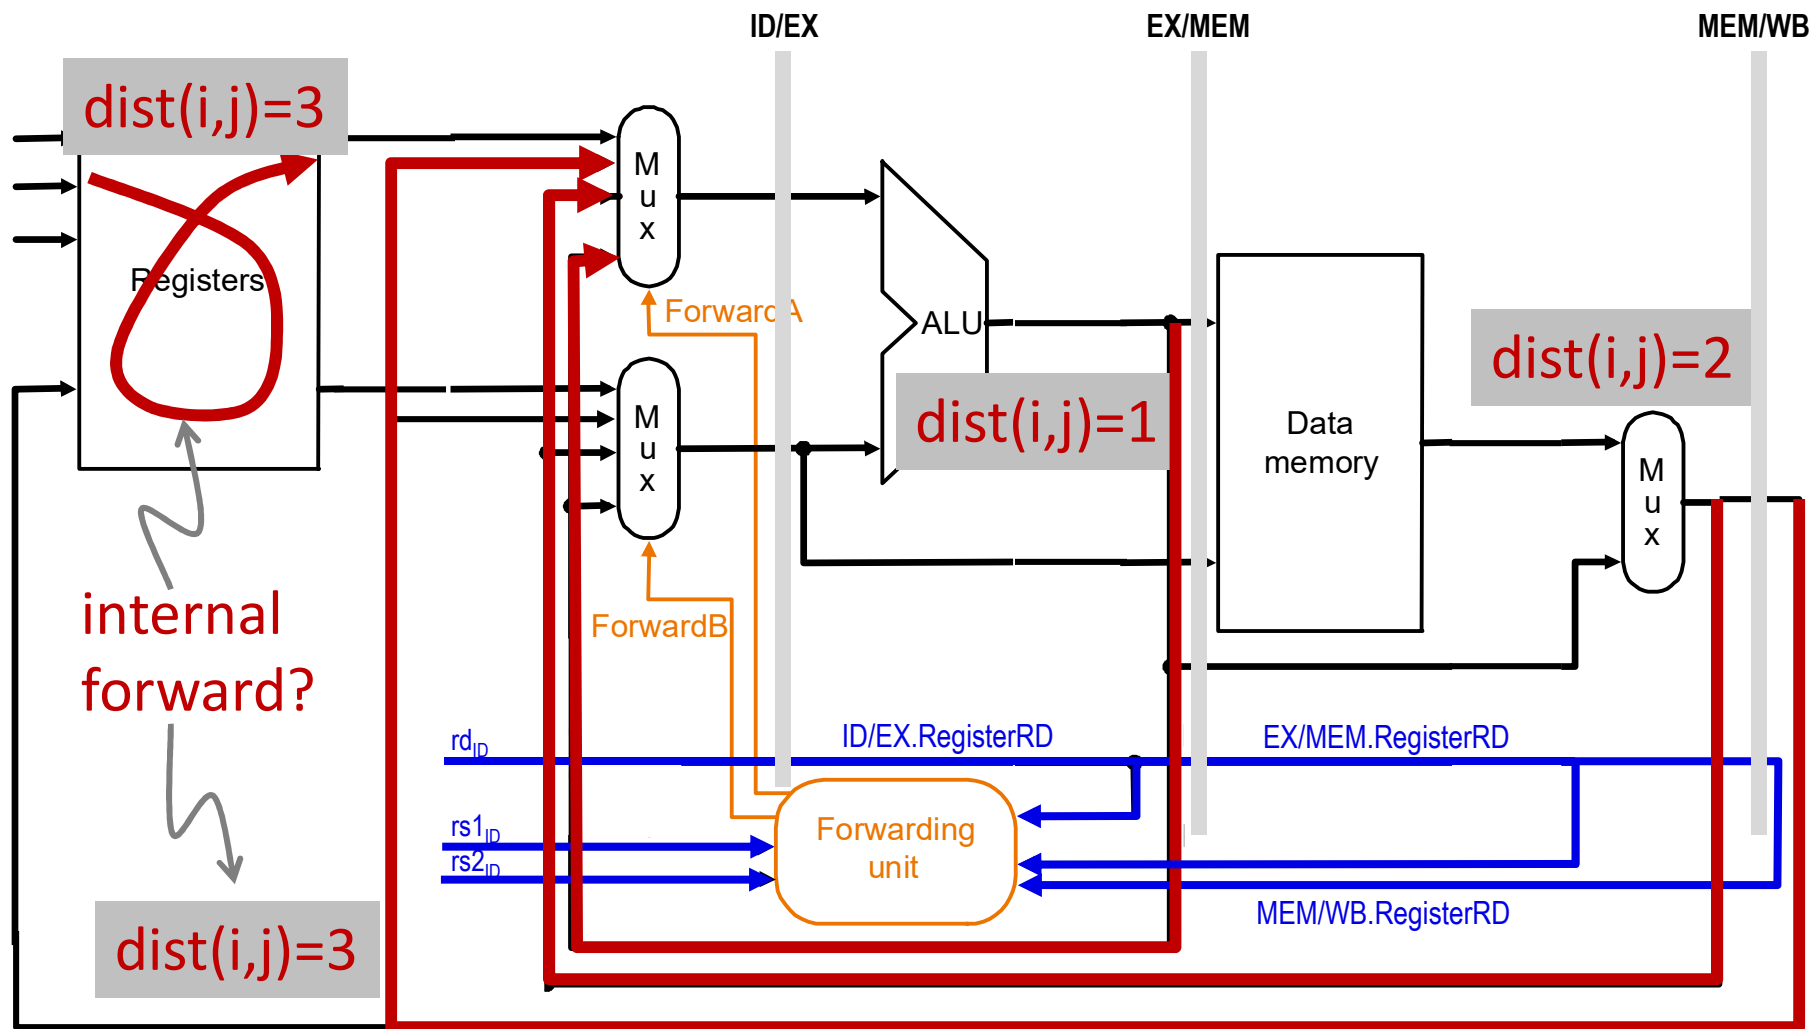


Resolving RAW Hazard by Forwarding

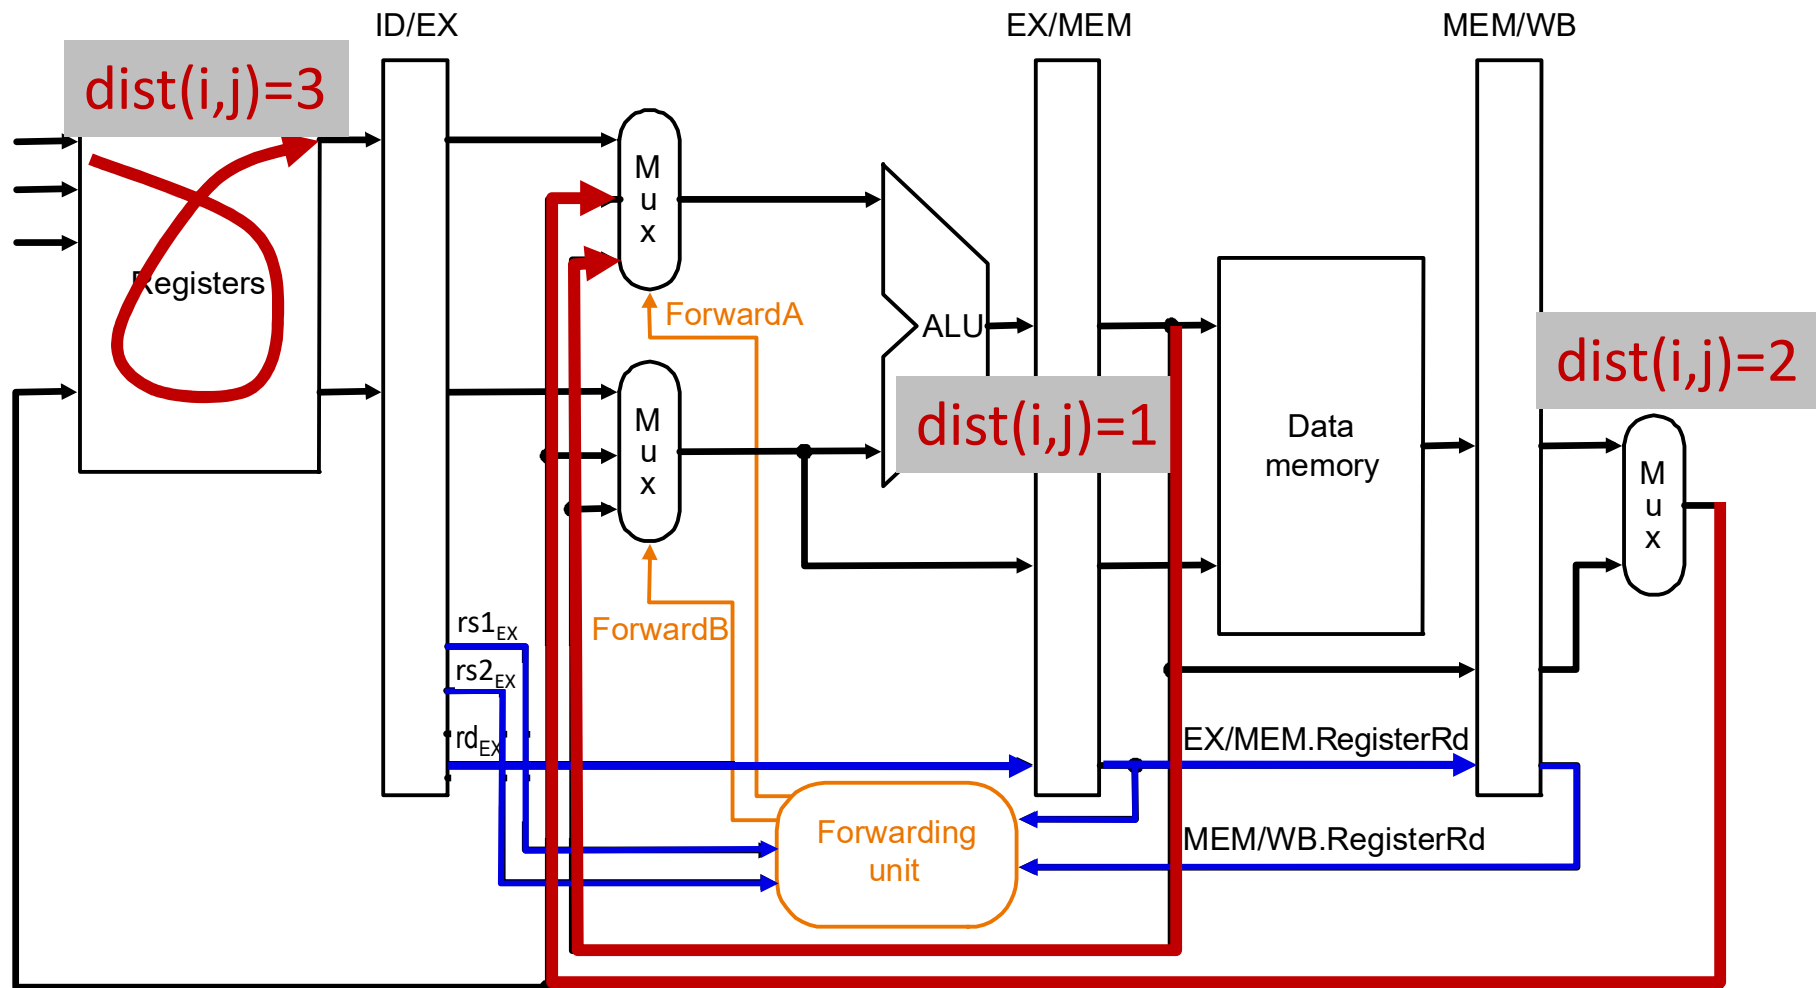
A hazard exists

- Older I_A and younger I_B have RAW hazard iff
 - I_B (R/I, LW, SW, Bxx or JALR) reads a register written by I_A (R/I, LW, or JAL/R)
 - $\text{dist}(I_A, I_B) \leq \text{dist}(\text{ID}, \text{WB}) = 3$
- To detect hazard in time to prevent, before I_B in ID reads a register, I_B needs to check if any I_A in EX, MEM or WB is going to update it
- Before: I_B need to stall for I_A to update RF
- Now: I_B need to stall for I_A to produce result
 - retrieve I_A result from datapath when ready
 - must retrieve from youngest if multiple hazards

Forwarding Paths (v1)



Forwarding Paths (v2)



better if EX is the fastest stage

Forwarding Logic (for v1)

```

if ( $rs1_{ID} \neq 0$ ) && ( $rs1_{ID} == rd_{EX}$ ) &&  $RegWrite_{EX}$  then
    forward writeback value from EX           // dist=1
else if ( $rs1_{ID} \neq 0$ ) && ( $rs1_{ID} == rd_{MEM}$ ) &&  $RegWrite_{MEM}$  then
    forward writeback value from MEM           // dist=2
else if ( $rs1_{ID} \neq 0$ ) && ( $rs1_{ID} == rd_{WB}$ ) &&  $RegWrite_{WB}$  then
    forward writeback value from WB           // dist=3
else
    use  $A_{ID}$                                 // dist > 3
  
```

Must prioritize young-to-old
 Why doesn't *waitForRs1*() appear?
 Isn't it bad to forward from LW in EX?

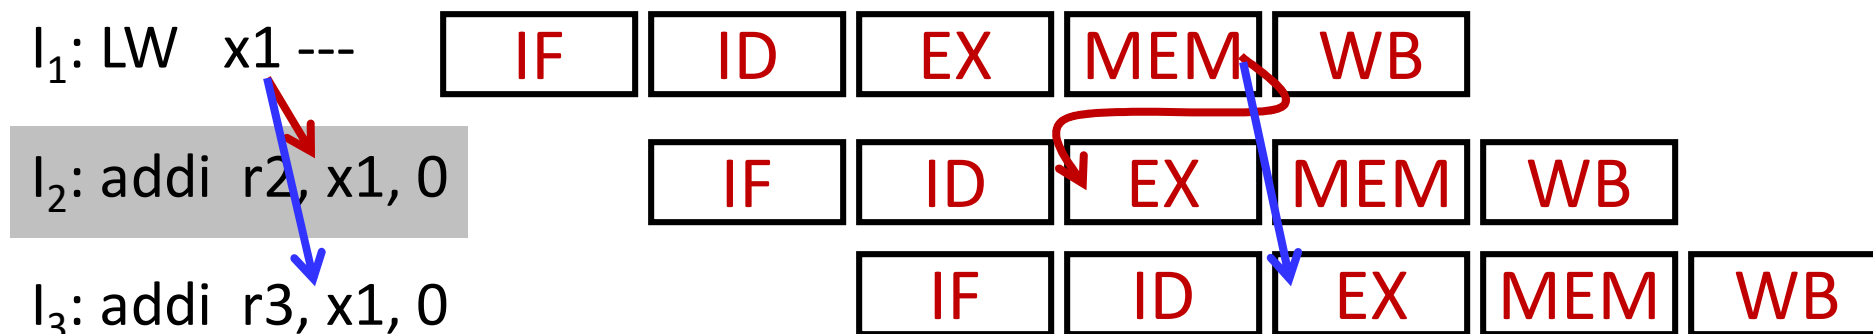
Data Hazard Analysis (with Forwarding)

	R/I-Type	LW	SW	Bxx	Jal	Jalr
IF						
ID						
EX	use produce	use	use	use	produce	use produce
MEM		produce	(use)			
WB						

- Even with forwarding, RAW dependence on immediate preceding LW results in hazard
- $$\text{Stall} = \{[(rs1_{ID} == rd_{EX}) \ \&\& \ \text{waitForRs1}(IR_{ID})] \ || \ i.e., \ op_{EX}=Lx$$

$$[(rs2_{ID} == rd_{EX}) \ \&\& \ \text{waitForRs2}(IR_{ID})]\} \ \&\& \ \text{MemRead}_{EX}$$

Historical: MIPS Load “Delay Slot”



- R2000 defined LW with arch. latency of 1 inst
 - invalid for I_2 (in LW's delay slot) to ask for LW's result
 - any dependence on LW at least distance 2
- Delay slot vs dynamic stalling
 - fill with an independent instruction (no difference)
 - if not, fill with a NOP (no difference)
- Can't lose on 5-stage . . . good idea?

Hint: 1. non-atomic instruction; 2. μ arch influence

Sample Assembly [P&H]

for (j=i-1; j>=0 && v[j] > v[j+1]; j-=1) { }

```

                addi    $s1, $s0, -1
for2tst:        slti    $t0, $s1, 0
                bne     $t0, $zero, exit2
                sll     $t1, $s1, 2
                add     $t2, $a0, $t1
                lw      $t3, 0($t2)
                lw      $t4, 4($t2)
                slt     $t0, $t4, $t3
                beq     $t0, $zero, exit2
                .....
                addi    $s1, $s1, -1
j               for2tst

```

1 stall or
1 nop (MIPS)


exit2:

Why not very deep pipelines?


- With only 5 stages, still plenty of combinational logic between registers
- “Superpipelining” \Rightarrow increase pipelining such that even intrinsic operations (e.g. ALU, RF access, memory access) require multiple stages
- What’s the problem?

Inst₀: addi x1, x0, 0

Inst₁: addi x2, x1, 0



Intel P4's Superpipelined Adder Hack



32-bit addition pipelined over 2 stages, $BW=1/\text{latency}_{16\text{-bit-add}}$
 No stall between back-to-back dependencies

Terminology

- Dependency
 - property of program
 - ordering requirement between instructions
- Pipeline Hazard:
 - property of uarch when interacting with program
 - (potential) violation of dependencies in program
- Hazard Resolution:
 - **static** \Rightarrow schedule instructions at compile time to avoid hazards
 - **dynamic** \Rightarrow detect hazard and adjust pipeline operation
Stall, Flush or Forward

Dependencies and Pipelining

(architecture vs. microarchitecture)

Sequential and atomic
instruction semantics



Defines what is correct;
doesn't say do it this way

True dependence between two
instructions may only require
ordering of certain sub-operations

