

18-447 Lecture 9: Control Hazard and Resolution

James C. Hoe
Department of ECE
Carnegie Mellon University

Housekeeping

- Your goal today
 - “simple” control flow resolution in in-order pipelines
 - there is more fun to come on this
- Notices
 - HW 2, due Mon 2/21
 - Lab 2, status check wk6, due wk7 (Handout #7)
 - HW 3, due Mon 2/28 (Handout #8)
 - Midterm 1, Wed 3/2 (look for practice midterm 1)
- Readings
 - P&H Ch 4

Format of the Midterm

- Covers lectures (L1~L10), HW, labs, assigned readings (from textbooks and papers)
- Types of questions
 - freebies: remember the materials
 - >> **probing: understand the materials** <<
 - applied: apply the materials in original interpretation
- ****70 minutes, 70 points****
 - point values calibrated to time needed
 - closed-book, one 8½x11-in² hand-written cribsheet
 - no electronics
 - use pencil or black/blue ink only


Control Dependence

- C-Code


```

{ code A }
if X==Y then
    { code B }
else
    { code C }
{ code D }
  
```

Control Flow Graph



Assembly Code
(linearized)




At ISA-level, control dependence == “data dependence on PC”

Applying Hazard Analysis on PC

	R/I-Type	LW	SW	Bxx	Jal	Jalr
IF	use	use	use	use	use	use
ID	produce	produce	produce			
EX				produce	produce	produce
MEM						
WB						


- All instructions read and write PC
- PC dependence distance is exactly 1
- PC hazard distance in 5-stage is at least 1
 - ⇒ Yes, there is RAW hazard
 - ⇒ Can't eliminate by forwarding; so must stall

Resolve Control Hazard by Stalling




Note: this is if decoding to non-control-flow; BR resolves in EX

Only 1 way to beat “true” dependence



Resolve Control Hazard by Guessing



Control Speculation for Dummies

- Guess $\text{nextPC} = \text{PC} + 4$ to keep fetching every cycle

Is this a good guess?

- ~20% of the instruction mix is control flow
 - ~50 % of “forward” control flow taken (if-then-else)
 - ~90% of “backward” control flow taken (end-of-loop)

Over all, typically ~70% taken and ~30% not taken


[Lee and Smith, 1984]

- Expect “ $\text{nextPC} = \text{PC} + 4$ ” ~86% of the time, but what about the remaining 14%?

What do you do when wrong?

What do you lose when wrong?


Control Speculation: PC+4



When $inst_h$ branch resolves

- branch target ($Inst_k$) is fetched
- flush instructions fetched since $inst_h$ ("wrong-path")

Pipeline Flush on Misprediction



$Inst_h$ is a taken branch; $Inst_i$ and $Inst_j$ fetched but not executed

Pipeline Flush on Misprediction

	t ₀	t ₁	t ₂	t ₃	t ₄	t ₅	t ₆	t ₇	t ₈	t ₉	t ₁₀
IF	h	i	j	k	l	m	n				
ID		h	i	bub	k	l	m	n			
EX			h	bub	bub	k	l	m	n		
MEM				h	bub	bub	k	l	m	n	
WB					h	bub	bub	k	l	m	n

branch resolved

Performance Impact

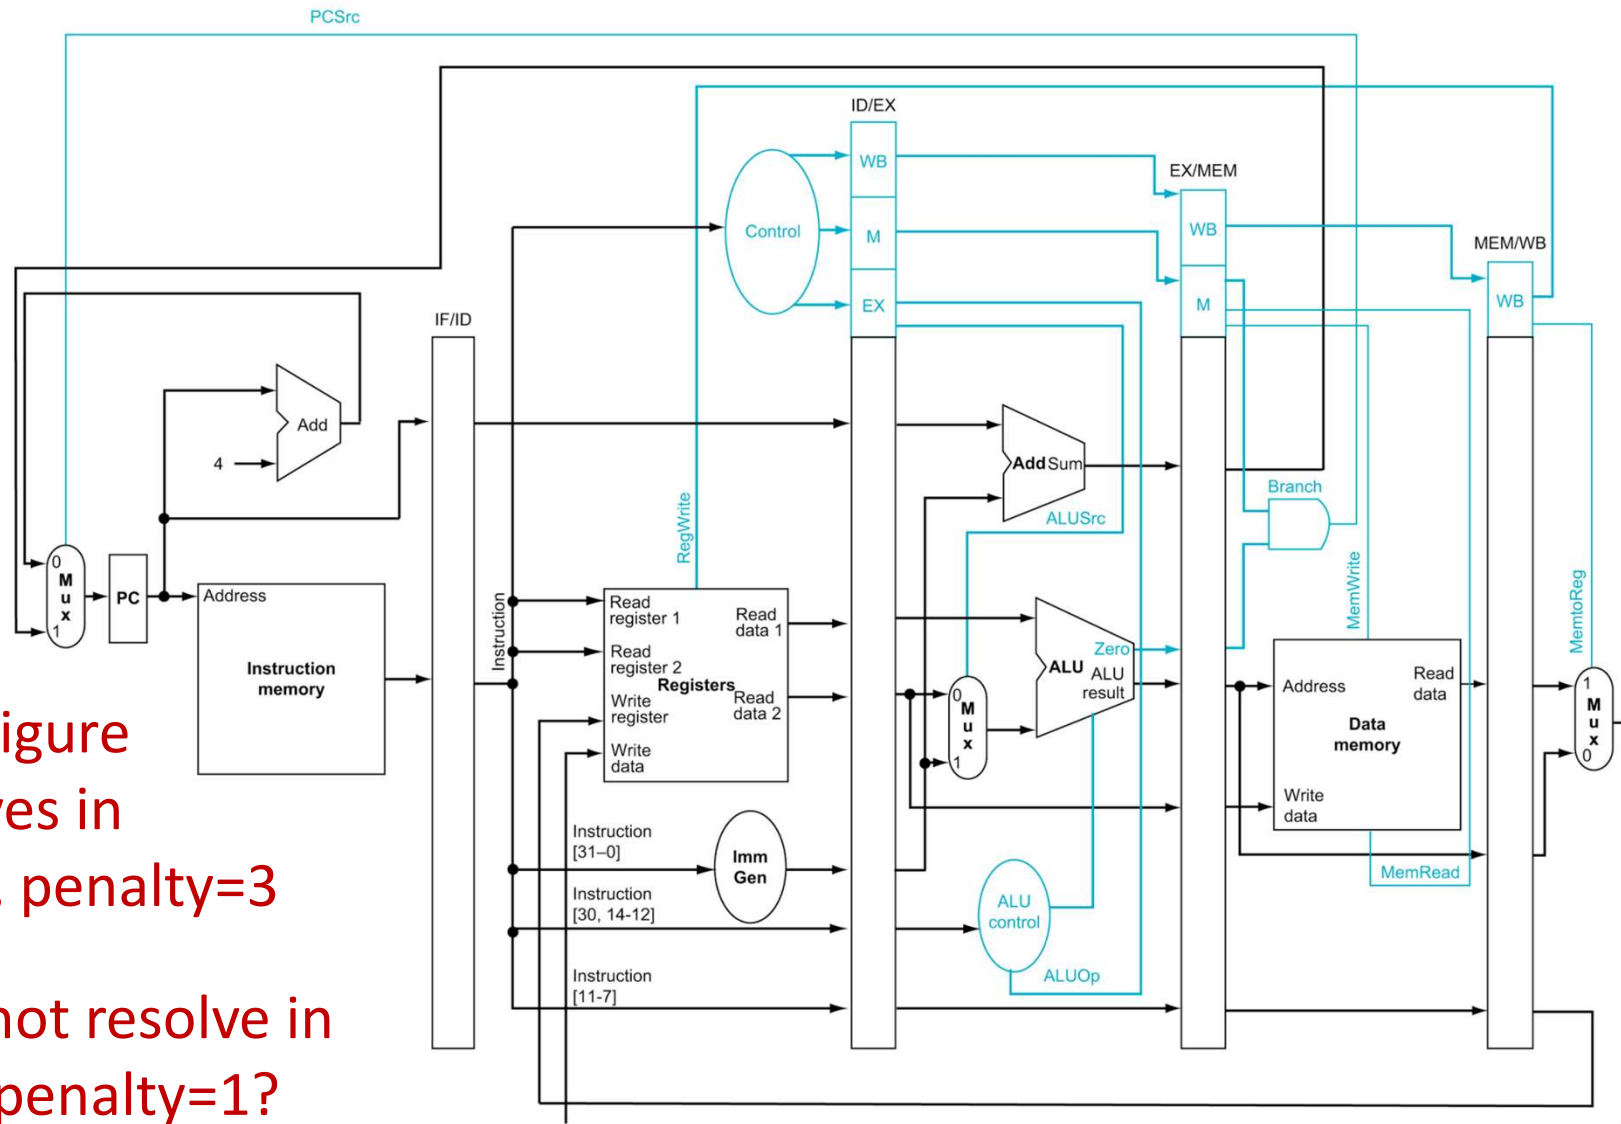
- Correct guess \Rightarrow no penalty most of the time!!
- Incorrect guess \Rightarrow 2 bubbles
- Assume
 - no data hazard stalls
 - 20% control flow instructions
 - 70% of control flow instructions are taken
 - $IPC = 1 / [1 + (0.20 * 0.7) * 2] =$
 $= 1 / [1 + 0.14 * 2] = 1 / 1.28 = 0.78$

*misprediction
rate*

*misprediction
penalty*

How to reduce the two penalty terms?

Reducing Mispredict Penalty



P&H figure
resolves in
MEM, penalty=3

Why not resolve in
ID so penalty=1?

MIPS R2000 ISA Control Flow Design


- Simple address calculation based on IR only
 - branch PC-offset: 16-bit full-addition
+ 14-bit half-addition
 - jump PC-offset: concatenation only
- Simple branch condition based on RF
 - one register relative ($>$, $<$, $=$) to 0
 - equality between 2 registers

No addition/subtraction necessary!

***Explicit ISA design choices to make possible
branch resolution in ID of a 5-stage pipeline***

Branch Resolved in ID

*what
about
this?*

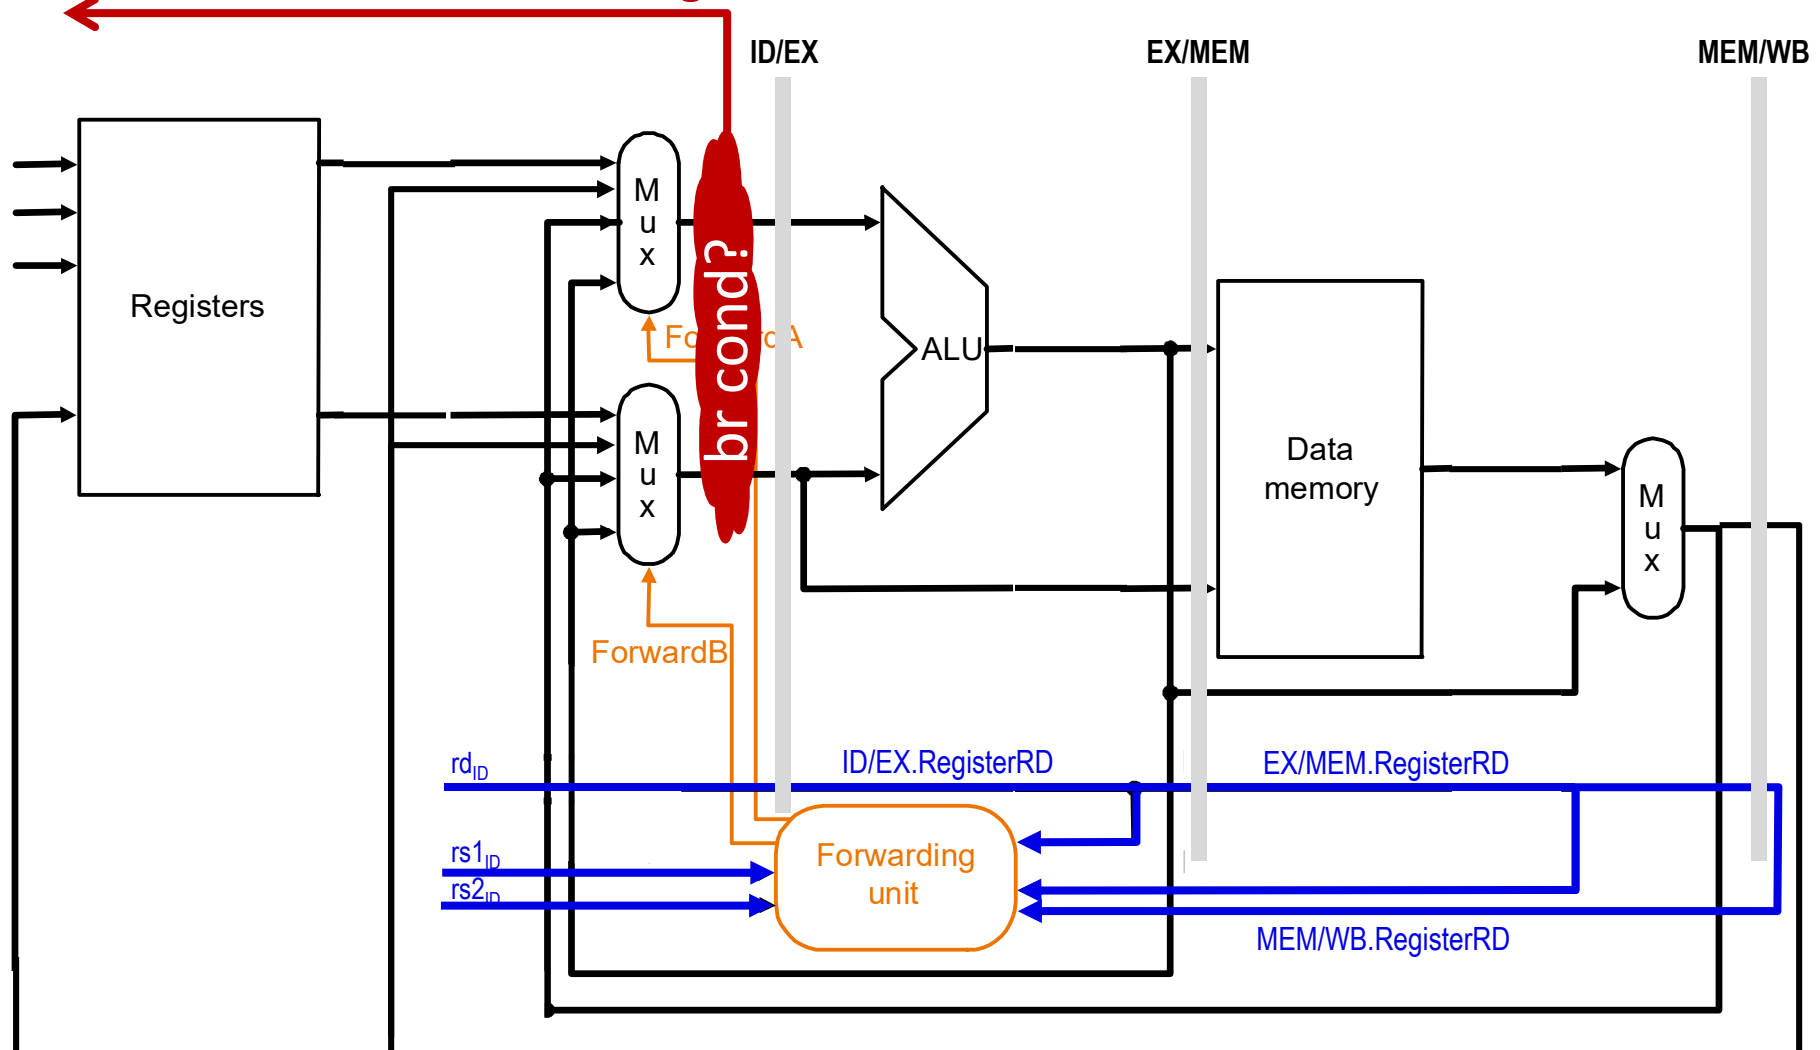


[Based on original figure from P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

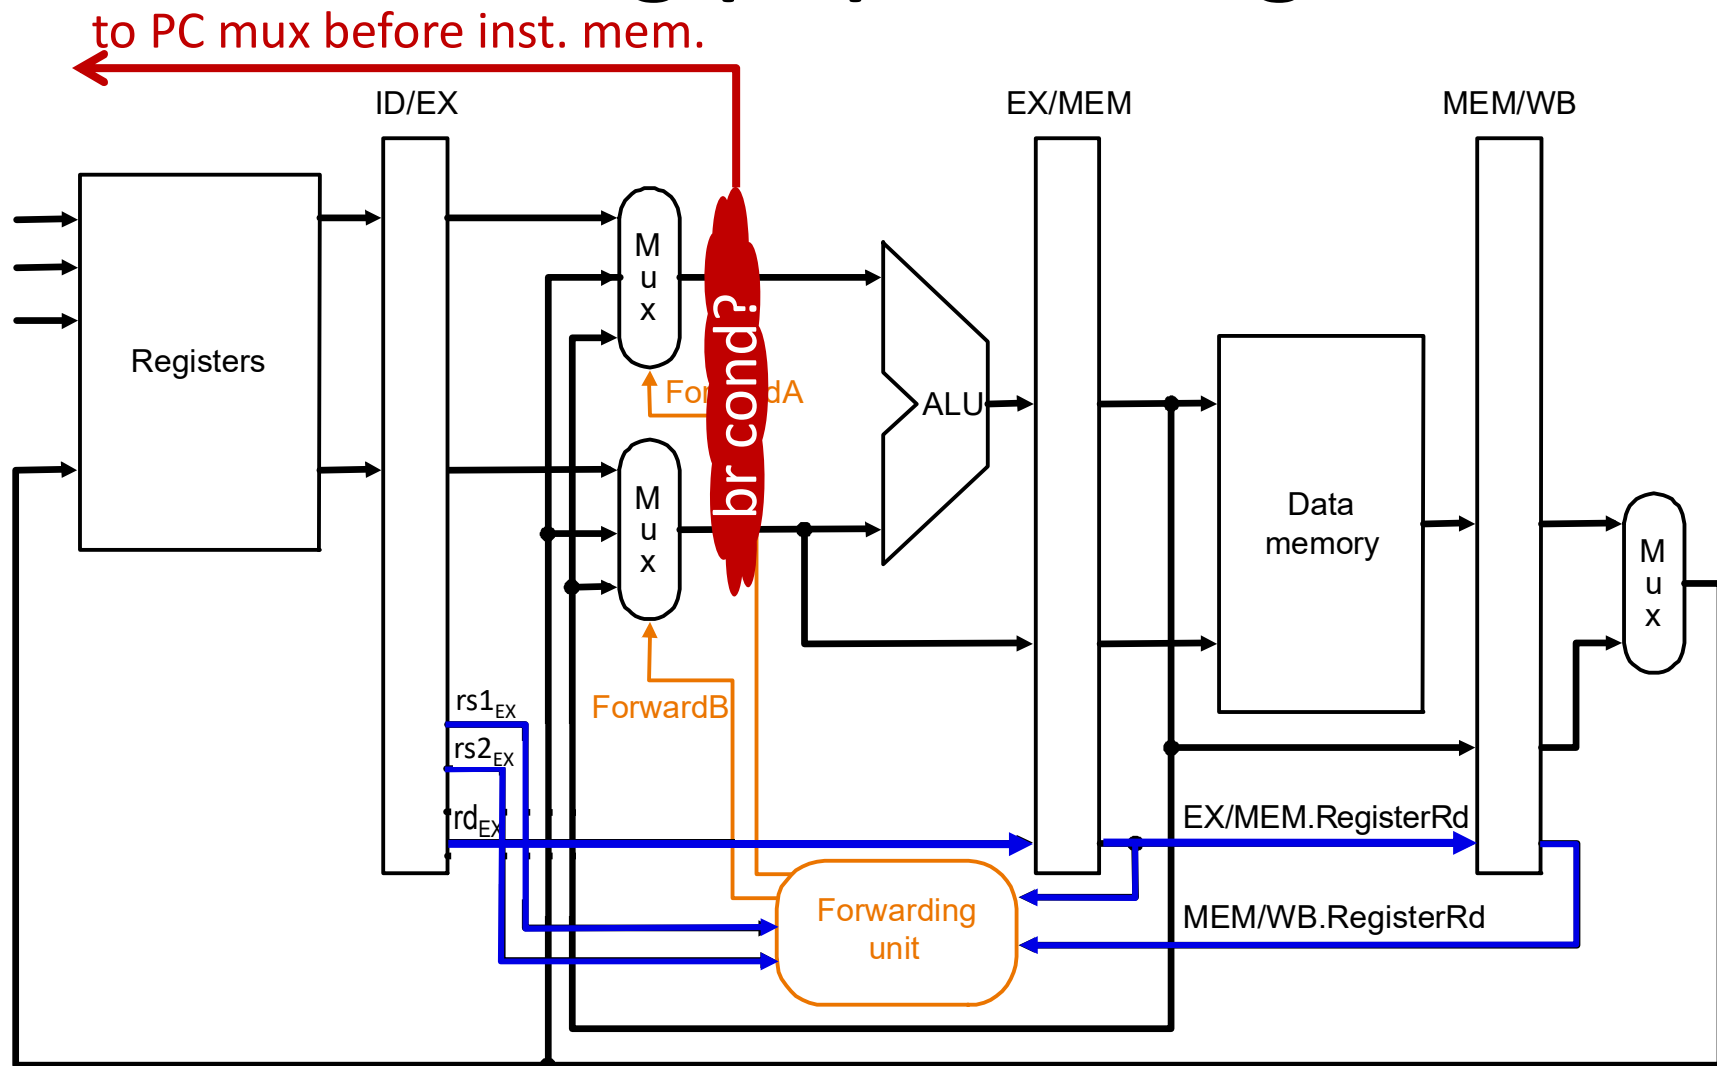
$$IPC = 1 / [1 + (0.2 * 0.7) * 1] = 0.88$$

Forwarding (v1): extend critical path


to nPC mux before the PC reg



Forwarding (v2): retiming hack



MIPS Branch Delay Slot



- Throwing PC+4 away cost 1 bubble; letting PC+4 finish won't hurt performance
- R2000 jump/branch has 1 inst. architectural latency
 - PC+4 after jump/branch always executed

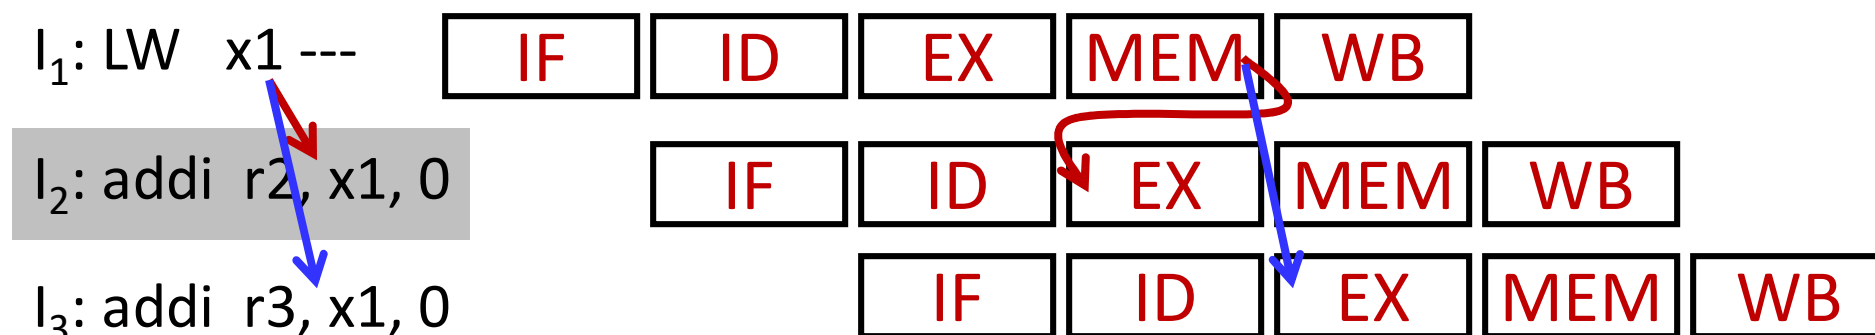
no need for pipeline flush logic

- if delay slot always do useful work, effective IPC=1
- ~80% of “delay slots” can be filled by compilers

$$IPC = 1 / [1 + (0.2 * \text{unfilled nop}) * 1] = 0.96$$

unfilled
nop

Also MIPS Load “Delay Slot”




- R2000 defined LW with arch. latency of 1 inst
 - invalid for I_2 (in LW's delay slot) to ask for LW's result
 - any dependence on LW at least distance 2
- Delay slot vs dynamic stalling
 - fill with an independent instruction (no difference)
 - if not, fill with a NOP (no difference)
- **MIPS**=**M**icroproc. without **I**nterlocked **P**ipeline **S**tages

Delay slots good idea? non-atomic, μ arch specific

Performance Impact

- Correct guess \Rightarrow no penalty most of the time!!
- Incorrect guess \Rightarrow 2 bubbles
- Assume
 - no data hazard stalls
 - 20% control flow instructions
 - 70% of control flow instructions are taken
 - $IPC = 1 / [1 + (0.20 * 0.7) * 2] =$
 $= 1 / [1 + 0.14 * 2] = 1 / 1.28 = 0.78$



How to reduce the two penalty terms?

In case you needed motivation

Basic Pentium III Processor Misprediction Pipeline

1	2	3	4	5	6	7	8	9	10
Fetch	Fetch	Decode	Decode	Decode	Rename	ROB Rd	Rdy/Sch	Dispatch	Exec

Basic Pentium 4 Processor Misprediction Pipeline


1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20
TC Nxt IP	TC Fetch	Drive	Alloc	Rename	Que	Sch	Sch	Sch	Disp	Disp	RF	RF	Ex	Flgs	Br Ck	Drive			

[The Microarchitecture of the Pentium 4 Processor,
Intel Technology Journal, 2001]

Can we make better guesses? (for when it is not MIPS or 5-stage)

- For control-flow instructions
 - why not always guess taken since 70% correct
 - need to know **taken target** to be helpful
- For non-control-flow instructions
 - can't do better than guessing **nextPC=PC+4**
 - still tricky since must guess before knowing it is control-flow or non-control-flow
- Guess **nextPC** from current **PC** alone, and fast!
- Fortunately
 - **instruction at same PC doesn't change**
 - **PC-offset target doesn't changes**
 - **okay to be wrong some of the time**

Branch Target Buffer (magic version)



- BTB
 - a giant table indexed by PC
 - returns the “guess” for nextPC
- When seeing a PC first time, after decoding, record in BTB . . .
 - PC + 4 if ALU/LD/ST
 - PC+offset if Branch or Jump
 - ?? if Jump Indirect
- Effectively guessing branches are always taken (and where to)

$$IPC = 1 / [1 + (0.20 * \underline{0.3}) * 2]$$

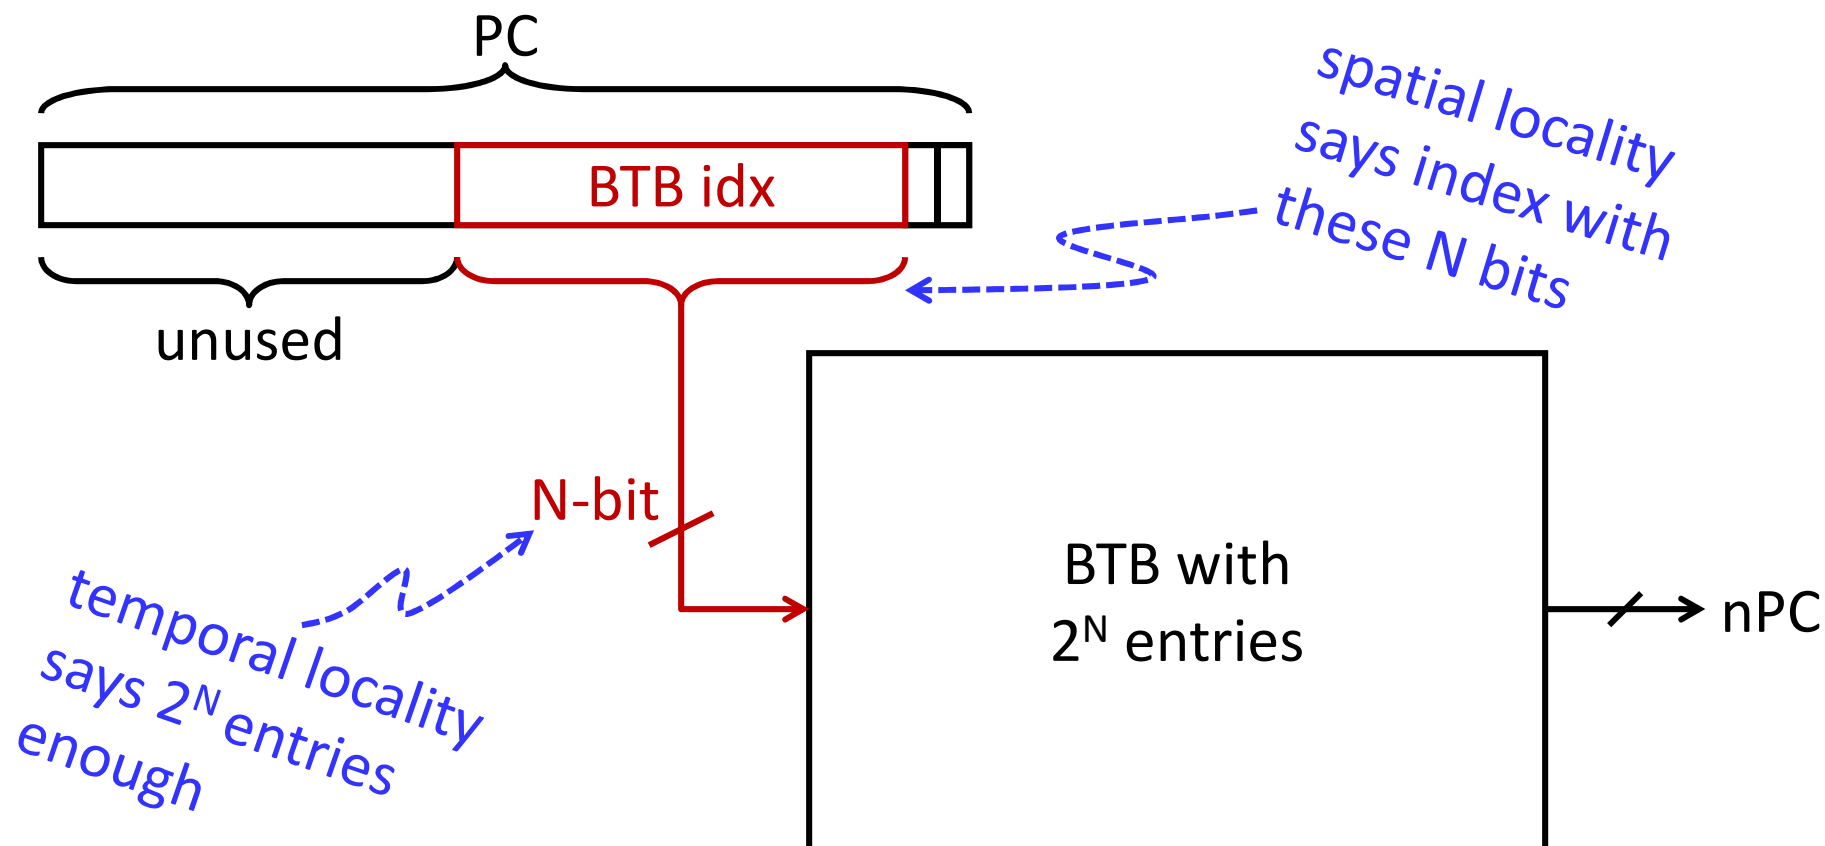
$$= 0.89$$

↖ If not taken

Locality Principle to the Rescue

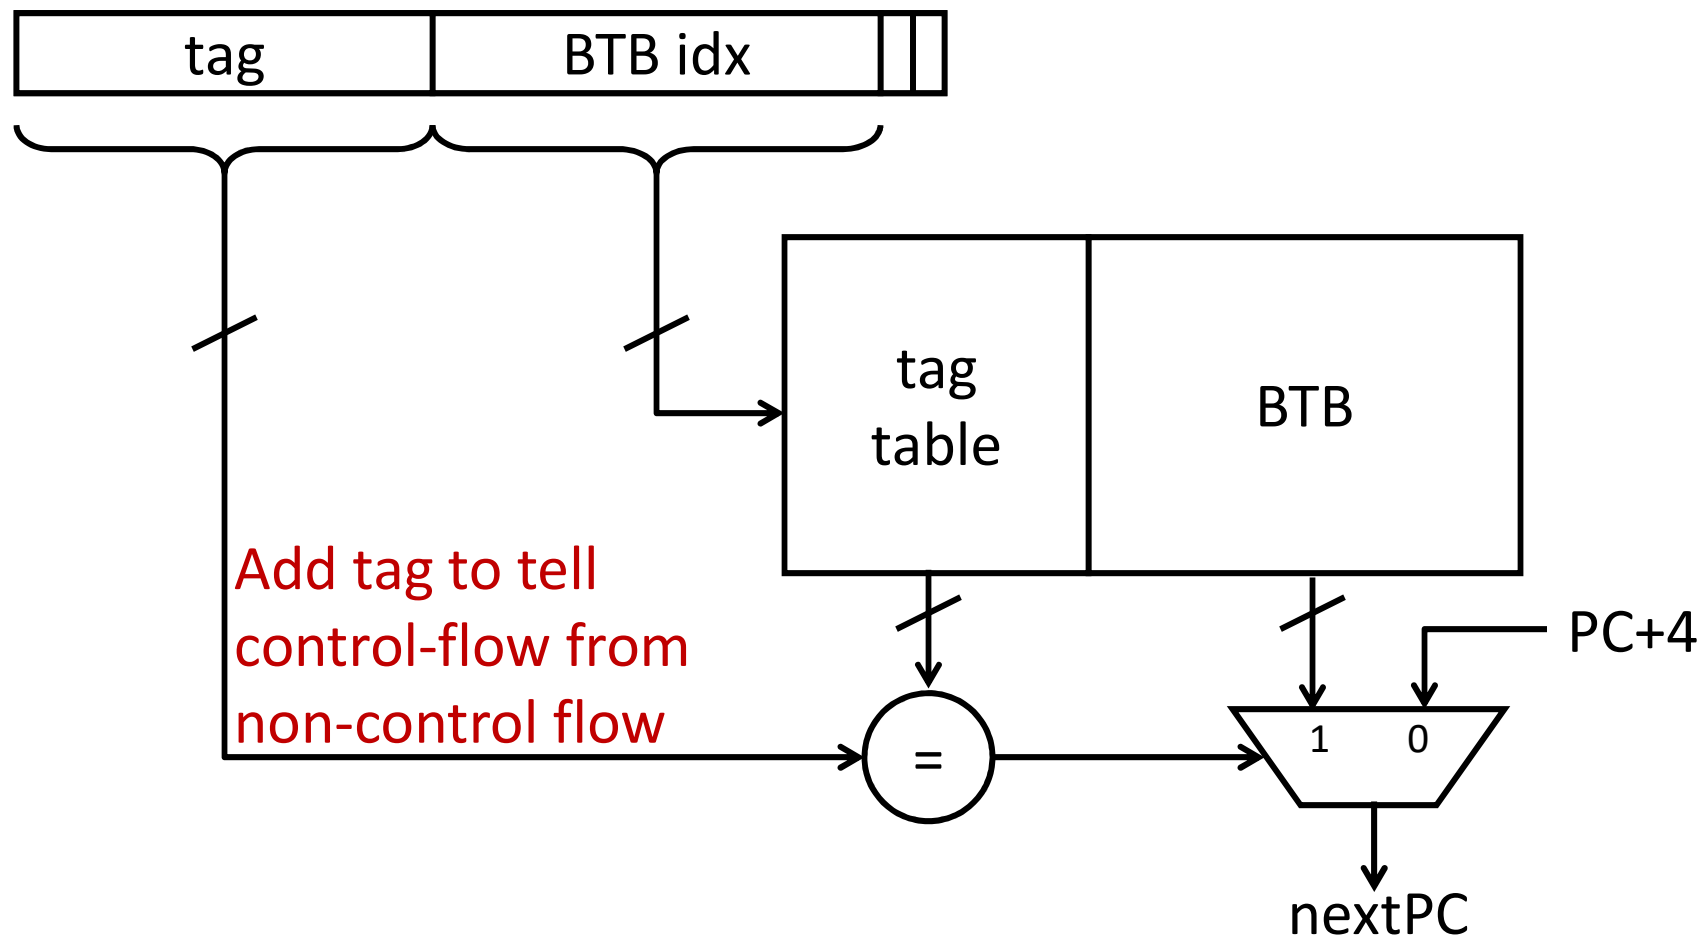
- **Temporal:** after accessing **A**, how many other distinct addresses before accessing **A** again?
 - **Spatial:** after accessing **A**, how many other distinct addresses before accessing **B**?
 - “Typical” programs have strong locality in memory references—instruction and data
we put them there ... BB, loops, arrays, structs ...
 - **Corollary:** a program with strong temporal and spatial locality access only a compact “**working set**” at any point in time
- ⇒ **just need BTB big enough for hot instructions**

Smaller BTB by Hashing



- “Hash” PC into a 2^N entry table
- What happens when two “hot” instructions collide? *No problem, as long as infrequent*

Even Smaller BTB after Tagging



Only hold control-flow instructions (save 80% storage)
Update tag and BTB for new branch after collision

Final 5-stage RISC Datapath & Control

