

18-447 Lecture 16: Cache Design in Context (Uniprocessor)

James C. Hoe
Department of ECE
Carnegie Mellon University

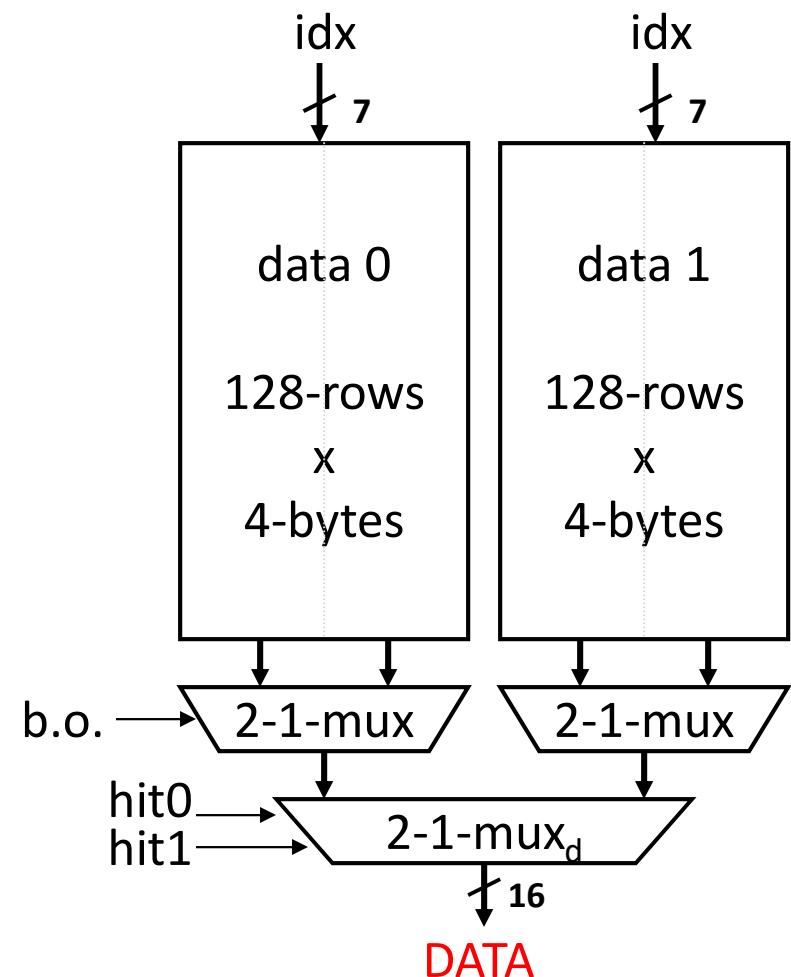


Housekeeping

- Your goal today
 - understand cache design and operation in context
 - focus on uniprocessor for now
- Notices
 - HW 5, due 4/4 (Handout #13)
 - Lab 3, due week 10
 - Midterm 1 regrade due Monday 3/28 noon
- Readings
 - P&H Ch 5

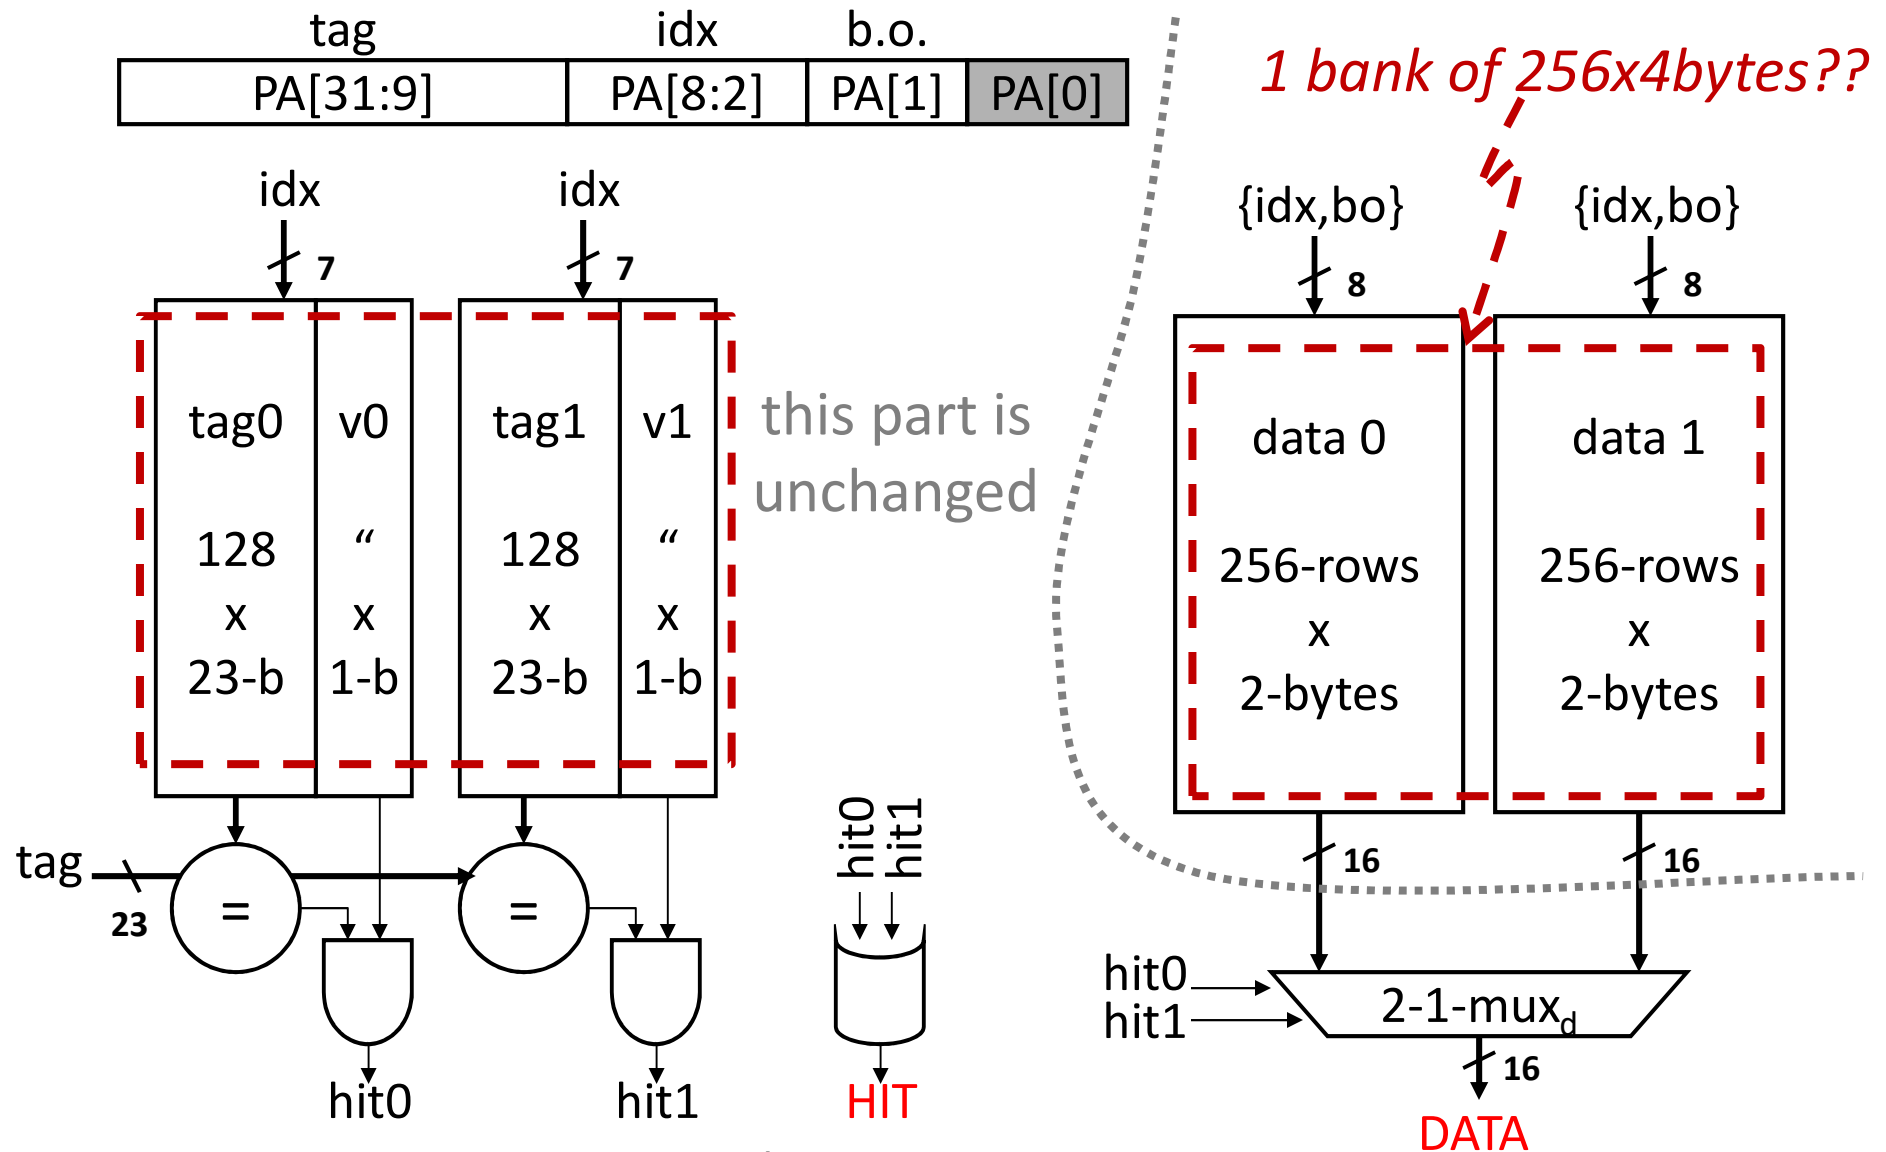
Format of the Midterm

- Covers lectures (L11~L18), HW, labs, assigned readings (from textbooks and papers)
- Types of questions
 - freebies: remember the materials
 - >> **probing: understand the materials** <<
 - applied: apply the materials in original interpretation
- ****70 minutes, 70 points****
 - point values calibrated to time needed
 - closed-book, one 8½x11-in² hand-written cribsheet
 - no electronics
 - use pencil or black/blue ink only
 - ****new rule** no questions in the final 20 10 min**

$M=2^{32}$, $a=2$, $C=1K$, $B=4$, $G=2$: “textbook” solution

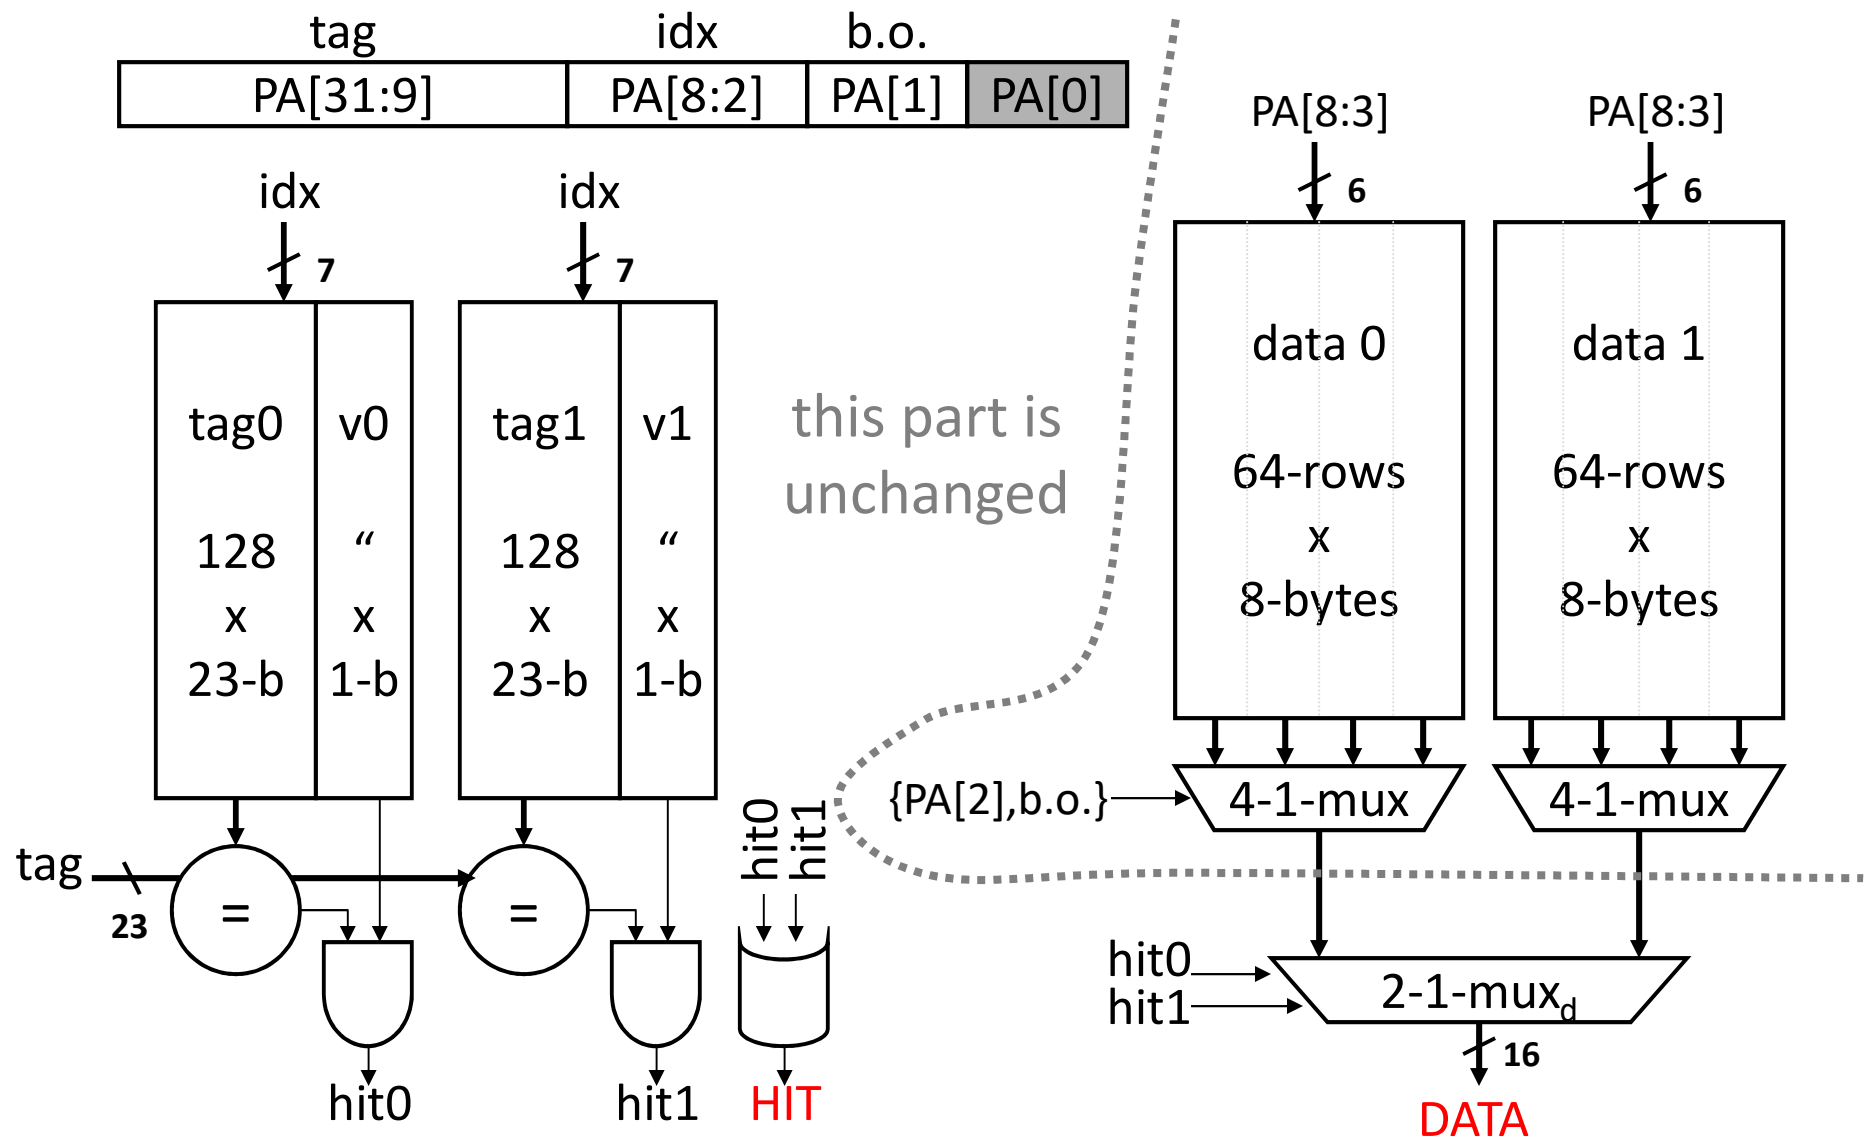


Same cache parameters but tune for “narrower” data SRAM banks



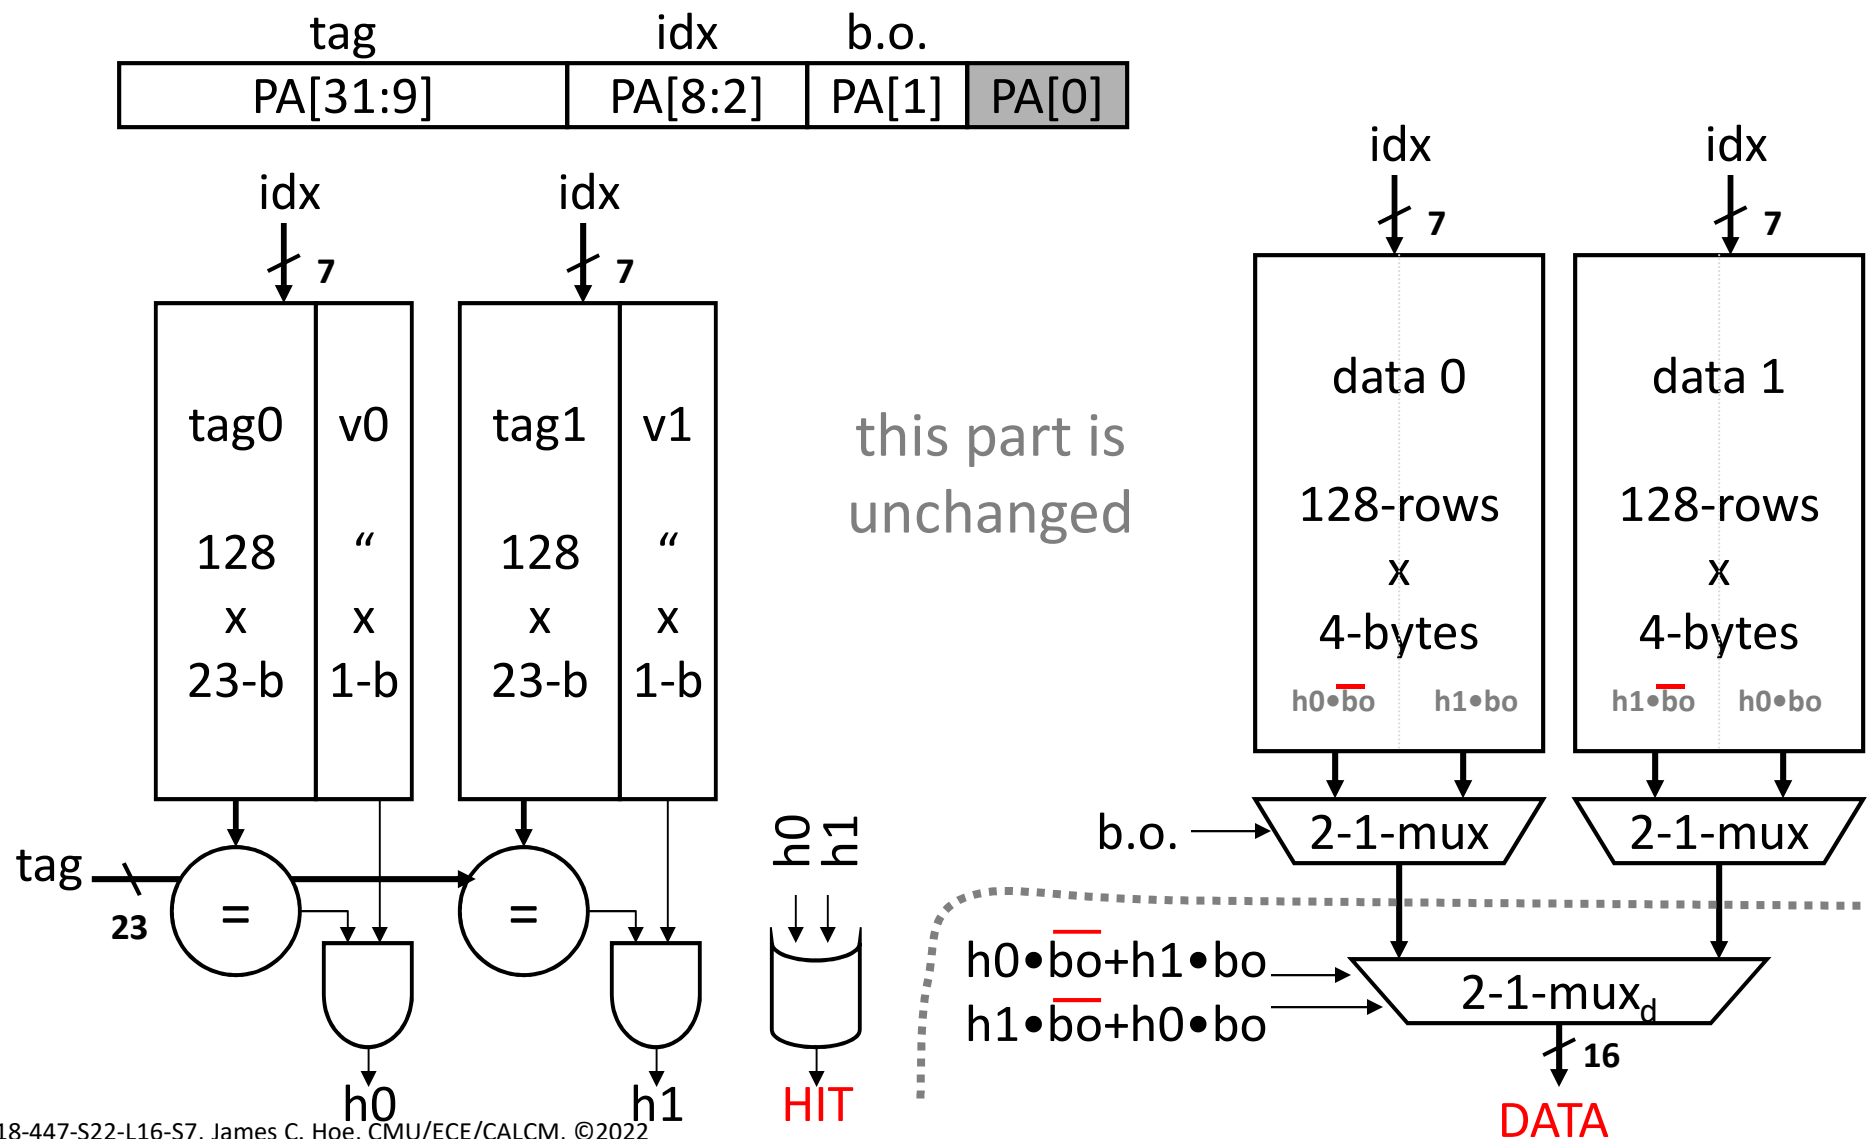
Can you make the tag SRAMs taller/narrower also?

Same cache parameters but tune for “fatter” data SRAM banks






Can you play the same trick on the tag SRAMs?

Same cache parameters but each block frame is interleaved over 2 SRAM banks



aBC Rule of Thumb Cribsheet

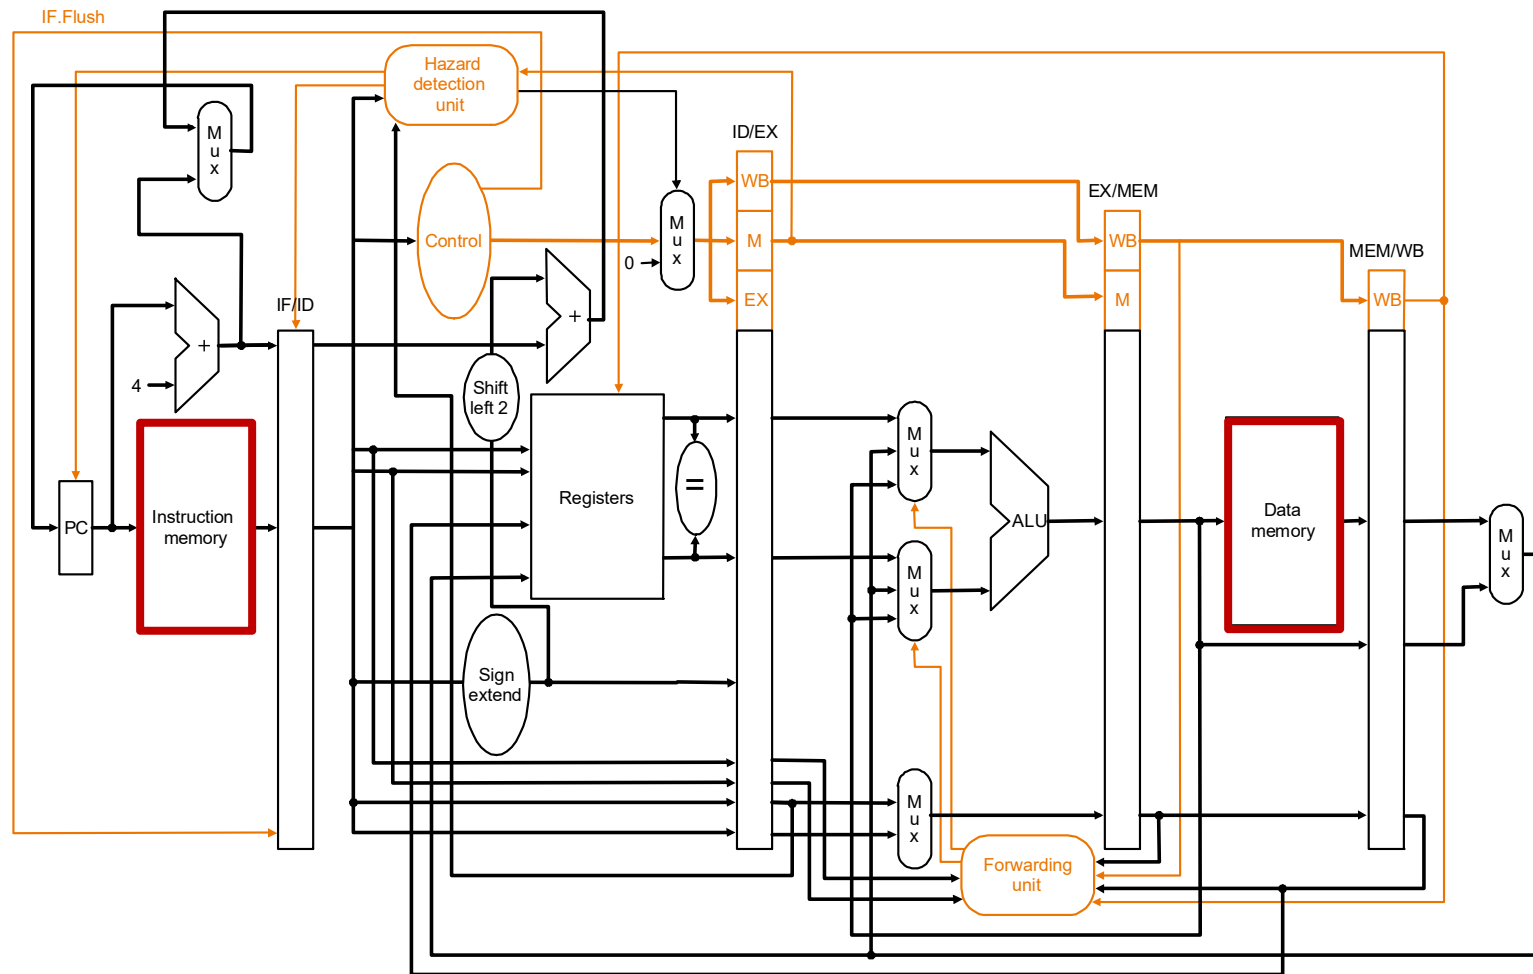


For “typical” programs

The Cache and You


(simple, single core from Lab)

The Context



[Based on original figure from P&H CO&D, COPYRIGHT 2004 Elsevier. ALL RIGHTS RESERVED.]

Cache Interface for Dummies



- Like the magic memory
 - present address, R/W command, etc
 - result or update valid after a short/fixed latency
- Except occasionally, cache needs more time
 - will become valid/ready eventually
 - what to do with pipeline until then? Stall!!

Recall

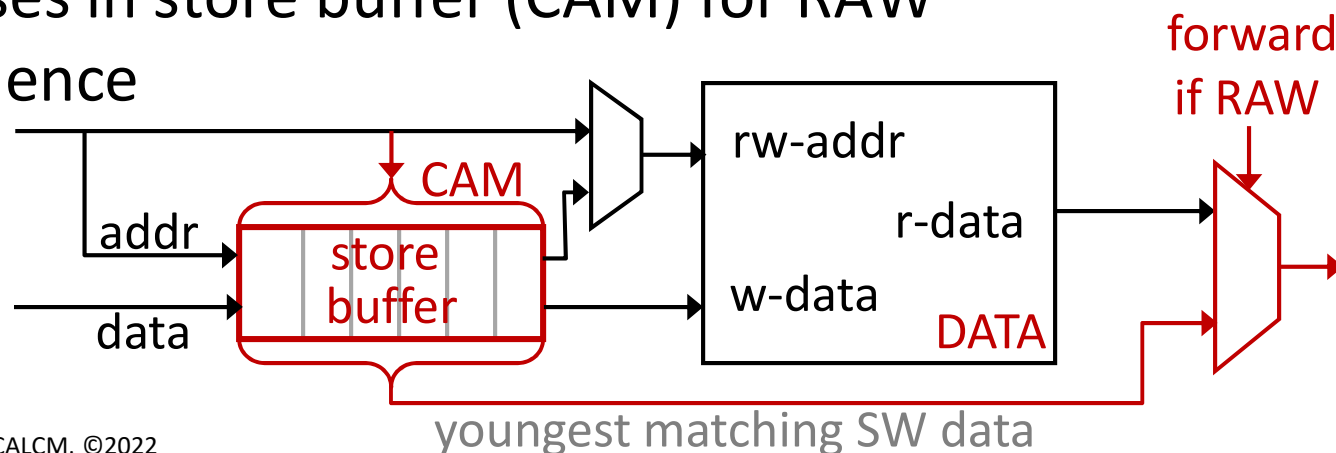
Adding Caches to In-order Pipeline

- On I-fetch and LW assuming 1-cyc SRAM lookup
 - if hit, just like magic memory
 - if miss, stall pipeline until cache ready
- On SW also assuming 1-cycle SRAM lookup
 - if miss, stall pipeline until cache ready (must we??)
 - if hit, ???...
- For SW, need to check tag array to ascertain hit before committing to write data array
 - data array write happens in the next cycle
 - if SW is followed immediately by LW

⇒ **structural hazard** ⇒ **stall**

Store Buffer

- Why stall when memory port is usually free?
- After tag array hit, buffer SW address and data until next free data array cycle (**not used by LW**)
 - allow younger LW to execute (out-of-order)
 - must ensure SW target block not evicted
- Memory dependence and forwarding
 - younger LW must check against pending SW-addresses in store buffer (CAM) for RAW dependence

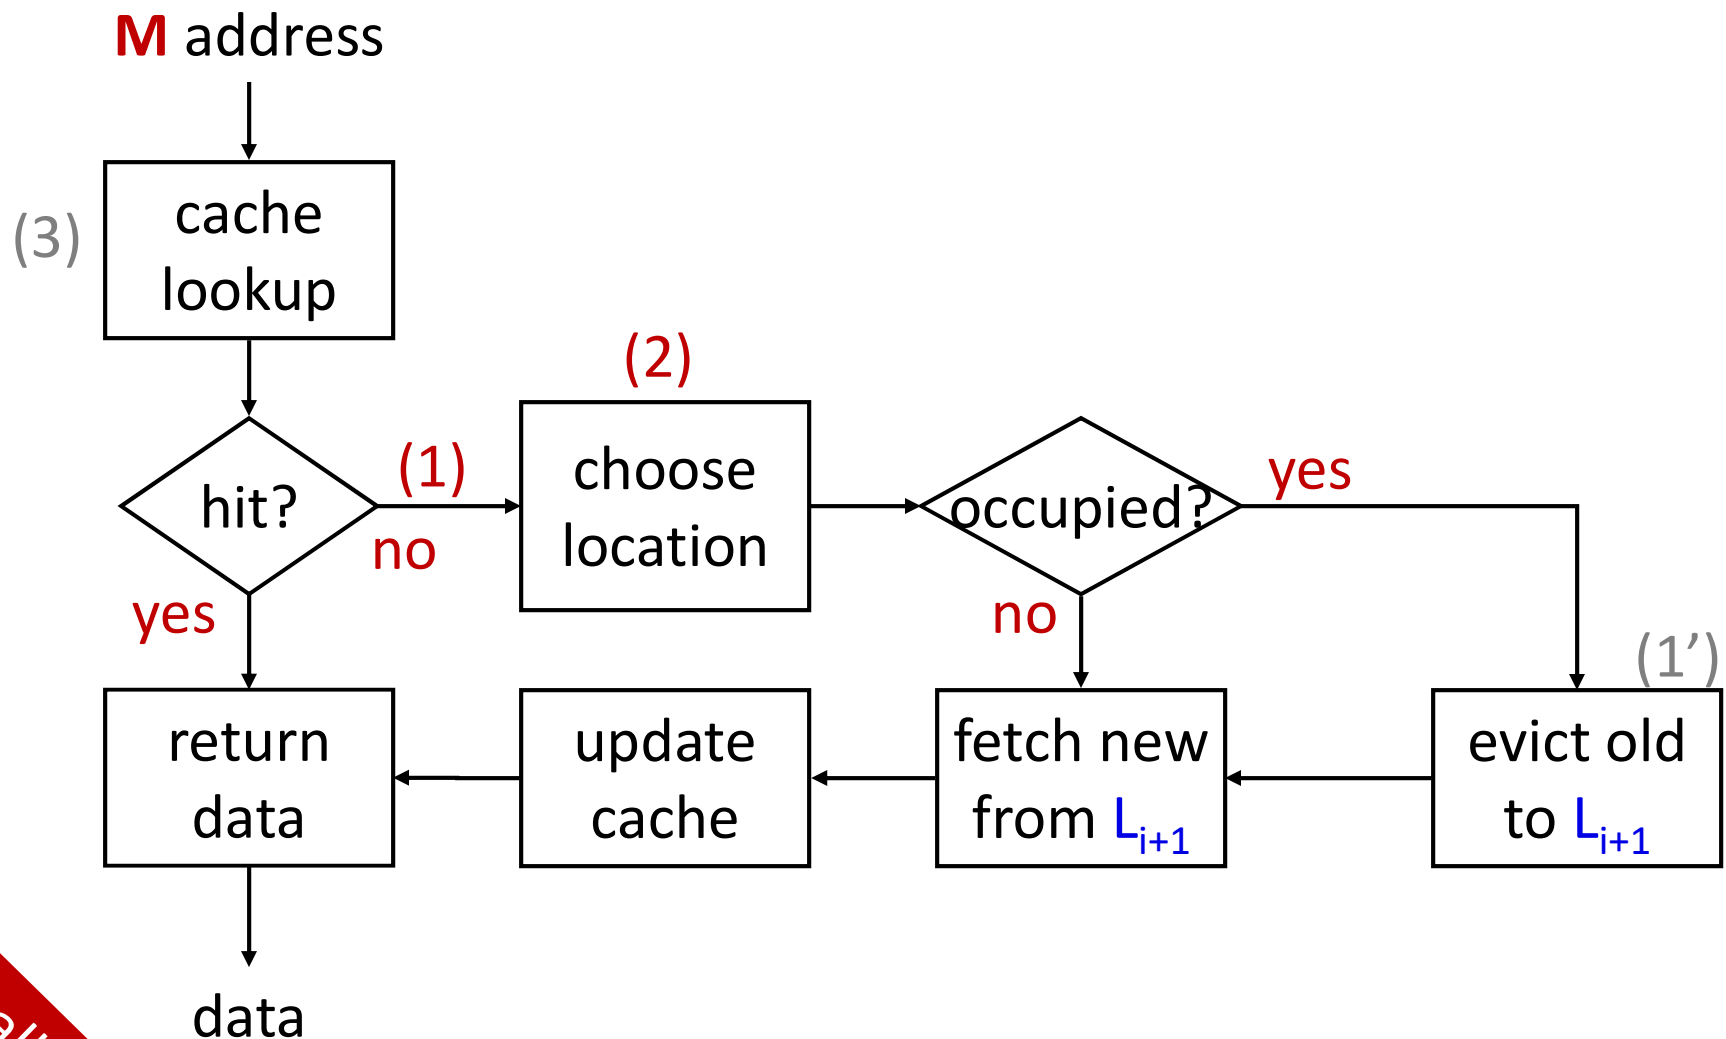


Must wait for a miss? (uniprocessor)

- In-order pipeline must stall for LW-miss
- Younger instructions can move ahead of SW-miss
 - except LW to same address; if so, stall or forward
 - additional SW-misses to same and different addr's can be “completed” from pipeline's view
- Modern out-of-order execution supports non-blocking miss handling for both LW and SW
 - too expensive to stall (CPU/memory speed gap)
 - significant complexity in
 - detecting and resolving memory dependencies
 - constructing precise exception state

Details and more details when building a cache for real

Basic “Cache Controller” (demand-driven version)



Recall

Write-Through Cache

- On write-hit in L_i , should L_{i+1} be updated?
- If yes, write-through
 - simple management (discard on replacement)
 - external agents (DMA and other proc's) see up-to-date values in DRAM
- With write-through, on a write-miss, should a cache block be allocated in L_i (aka write-allocate)?

- Write-through to DRAM not viable today
3.0GHz, IPC=2, 10% SW, ~8byte/SW \Rightarrow ~5GB/sec
L1 (w.o. ECC) write-through to L2 (w. ECC) still useful


Write-Back Cache

- Hold changes in L_i until block is displaced to L_{i+1}
 - on read or write miss, entire block is brought into L_i
 - LWs and SWs hit in L_i until replacement
 - on replacement, L_i copy written back out to L_{i+1}
adds latency to load miss stall
- “Dirty” bit optimization
 - keep per-block status bit to track if a block has been modified since brought into L_i
 - if not dirty, no write-back on replacement
- What if a DMA device wants to read a DRAM location with a dirty cached copy?

How to find out? How to access?

Write-Back Cache and DMA

- DRAM not always up-to-date if write-back
- DMA should see up-to-date value (aka, cache coherent)
- Option 1: SW flushes whole cache or specific blocks before programming DMA
- Option 2: cache monitors snoop bus for external requests
 - ask request to a dirty location to “retry”
 - write out dirty copy before request is repeated




Idempotency and Side-effects

- Loading from real memory location $M[A]$ should return most recent value stored to $M[A]$
 - \Rightarrow writing $M[A]$ once is the same as writing $M[A]$ with same value multiple times in a row
 - \Rightarrow reading $M[A]$ multiple times returns same value

This is why memory caching works!!


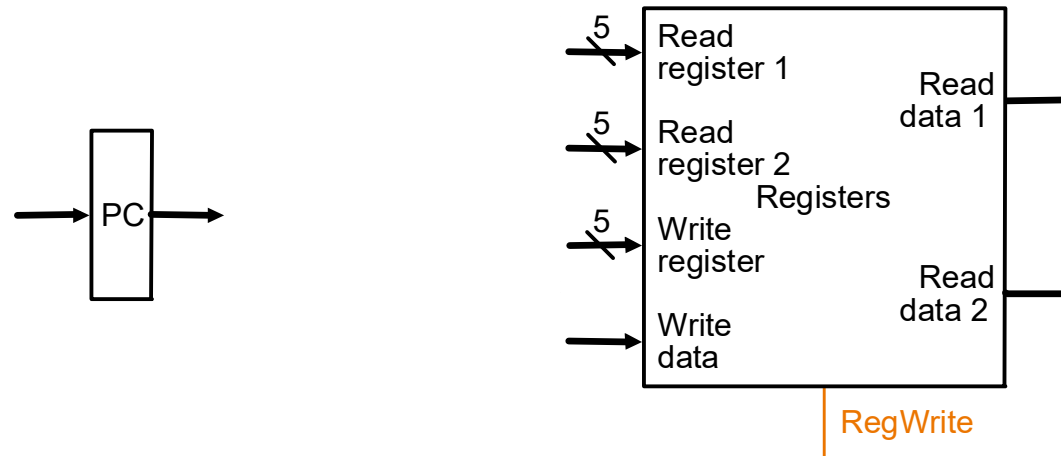
- LW/SW to mmap locations can have side-effects
 - reading/writing mmap location can imply commands and other state changes
 - e.g., a mmap device that is a FIFO
 - SW to 0xffff0000 pushes value
 - LW from 0xffff0000 returns popped value



What happens if 0xffff0000 is cached?

Recall

Program Visible State (aka Architectural State)



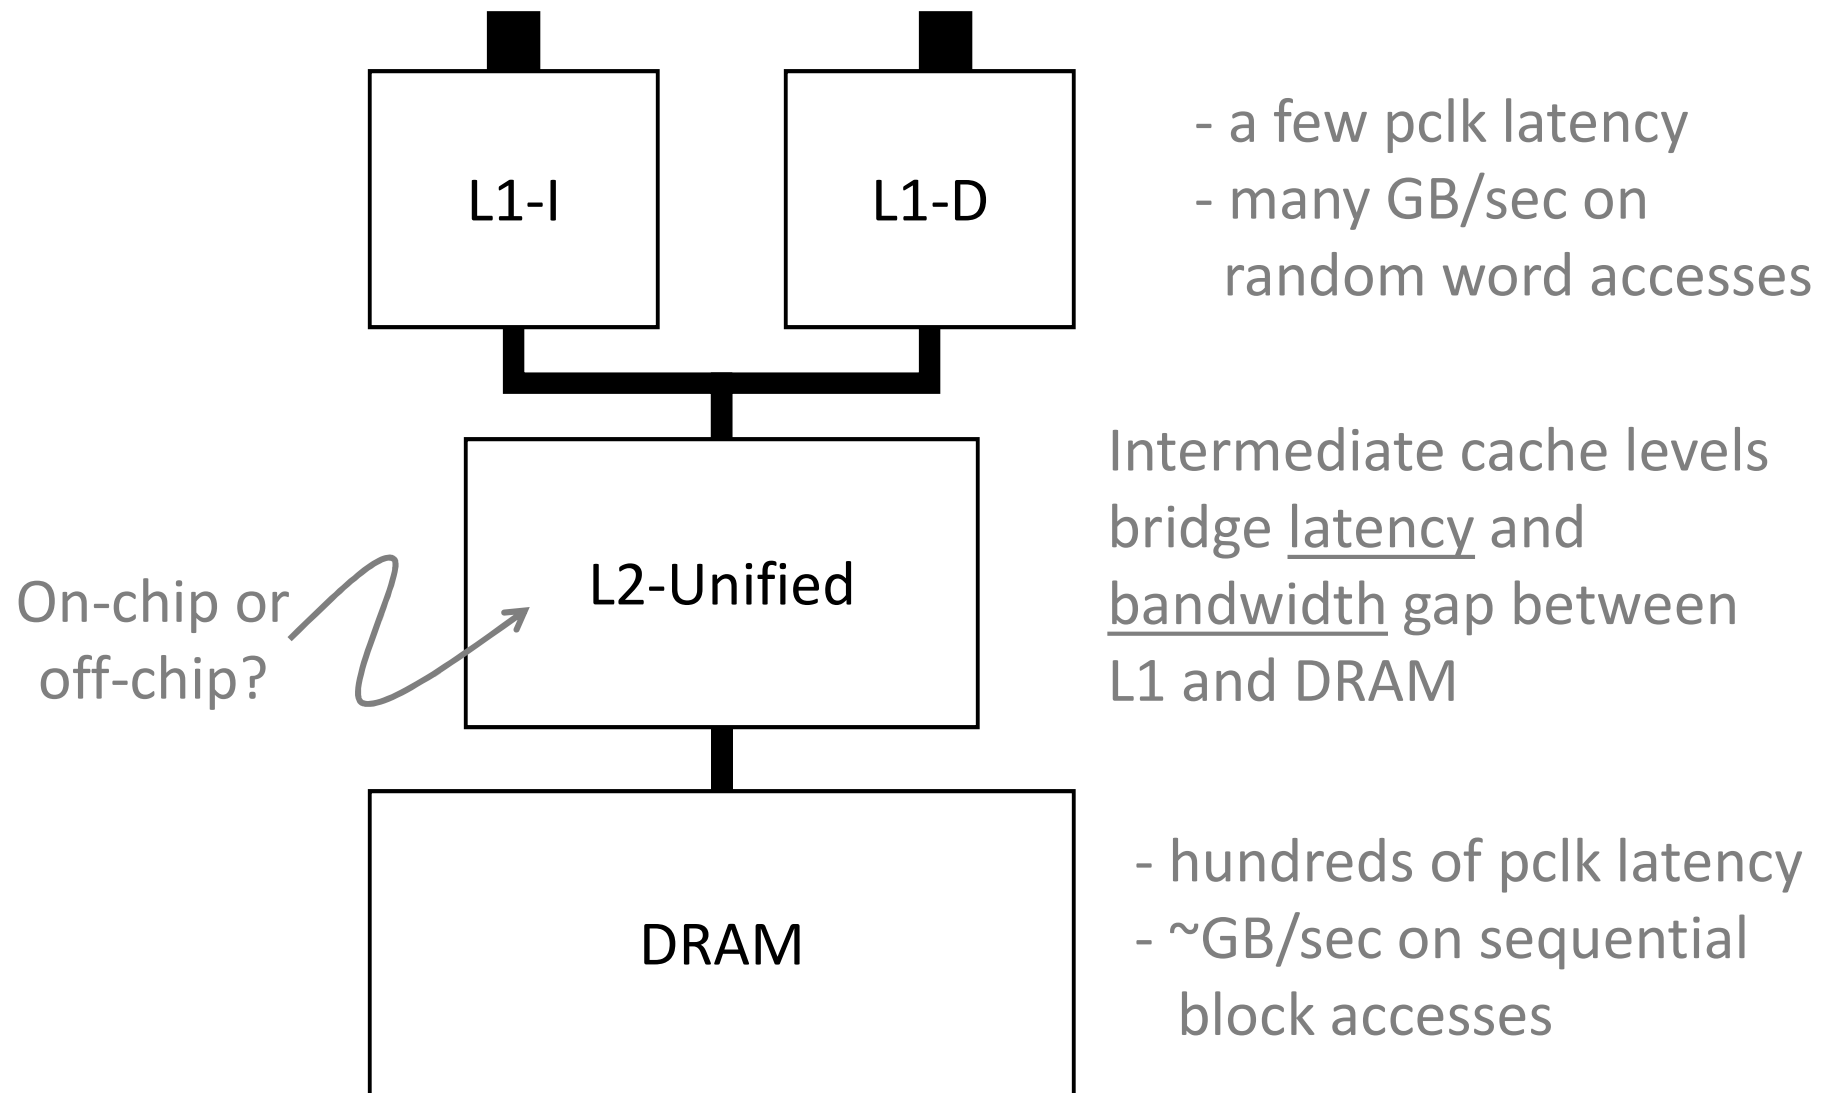
Recall

Harvard vs Princeton Architecture

- Historically
 - “Harvard” referred to Aiken’s Mark series with separate instruction and data memory
 - “Princeton” referred to von Neumann’s unified instruction and data memory
- Contemporary usage: split vs unified “caches”
- L1 I/D caches commonly split and asymmetrical
 - double bandwidth and no-cross pollution on disjoint I and D footprints
 - I-fetch smaller footprint, high-spatial locality and read-only \Rightarrow I-cache smaller, simpler

what about self-modifying code?
- L2 and L3 are unified for simplicity

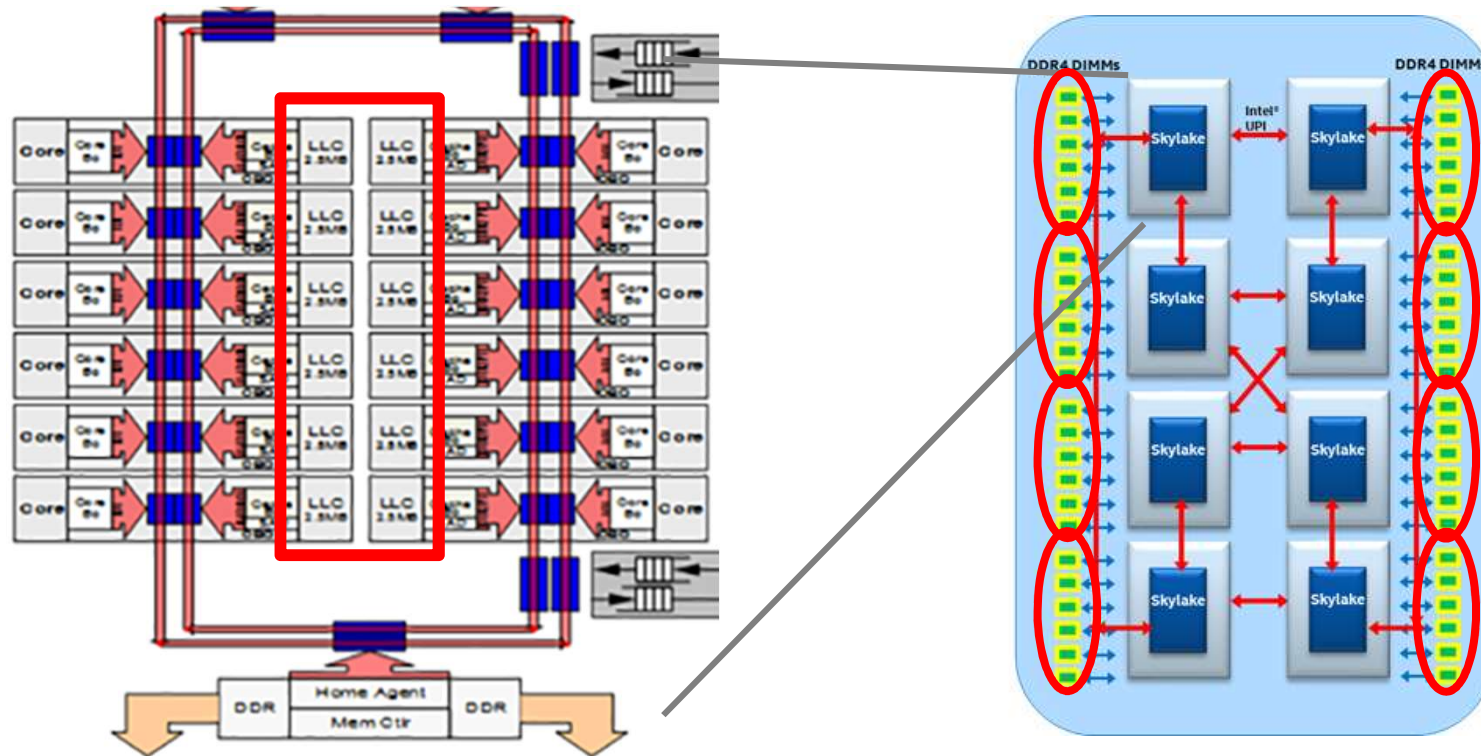
Multi-Level Caches



aBC of Multi-Level Cache Design

- Upper-level caches (L1)
 - small **C**: upper-bound by SRAM access time
 - smallish **B**: upper-bound by **C/B** effects
 - **a**: required to counter **C/B** effects
- Lower-level caches (L2, L3, etc.)
 - large **C**: upper-bound by chip area
 - large **B**: to reduce tag storage overhead
 - **a**: upper bound by complexity and speed
- New very large (10s MB) on-chip caches are distributed structures for multicores
 - same basic notions of ways and sets
 - but they don't look or operate anything like "textbook"

Modern Last-Level Cache (LLC)



[<https://software.intel.com/en-us/articles/intel-xeon-processor-scalable-family-technical-overview>]

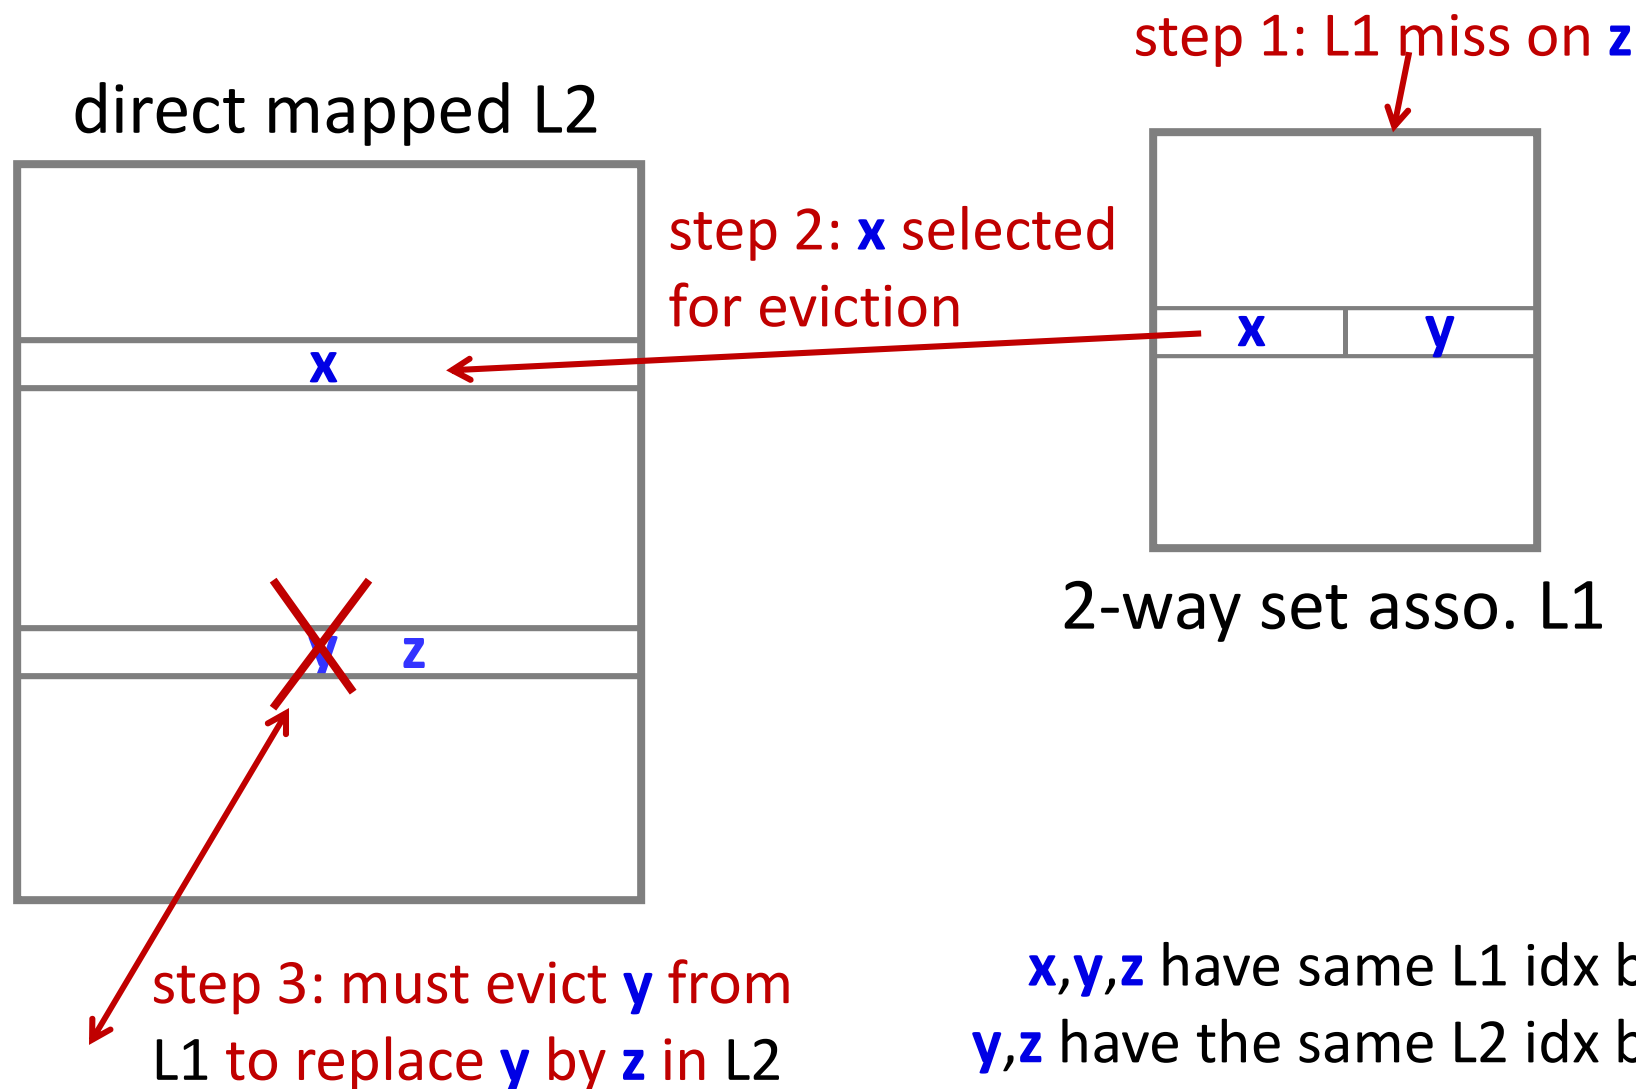
- Disaggregated, asynchronous structure; shared by all cores within a socket
- Hold, fast “coherent” copies of local and remote DRAM locations

Departure from classic uniproc. hierarchy

Inclusion Principle

- Classically, L_i contents is always a subset of L_{i+1}
 - if an address is important enough to be in L_i , it must be important enough to be in L_{i+1}
 - external agents (DMA and other proc's) only have to check the lowest level to know if an address is cached—do not need to consume L1 bandwidth
- Inclusion no longer taken as a given
 - nontrivial to maintain if L_{i+1} has lower associativity
 - too much redundant capacity in multicore with many per-core L_i and shared L_{i+1}

Inclusion Violation Example



x,y,z have same L1 idx bits
y,z have the same L2 idx bits
x,{y,z} have different L2 idx bits

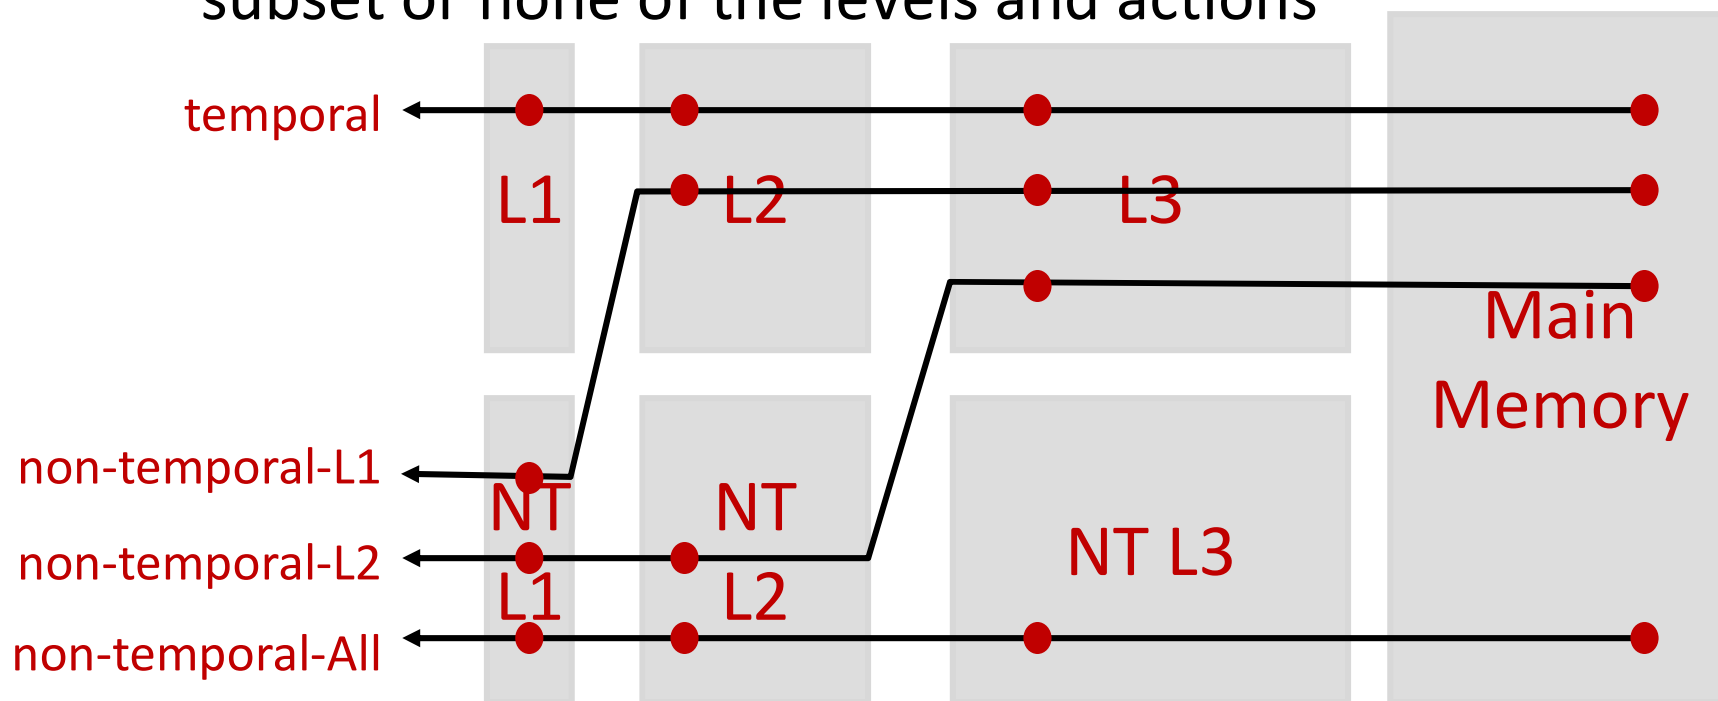
Aside: Victim “Cache”

- High-associativity is an expensive solution to avoid conflicts by a few stray addresses
- Augment a low-associative main cache with a very small but fully associative victim cache
 - blocks evicted from main cache is first held in victim cache
 - if an evicted block is referenced again soon, it is returned to main cache
 - if an evicted block doesn't get referenced again, it will eventually be displaced from victim cache to next level

Plays a different role outside of standard
memory hierarchy stacking

Aside: Software-Assists

- Separate “temporal” vs “non-temporal” hierarchy
 - exposed in the ISA (e.g., Intel IA64 below)
 - load and store instructions include **hints** about where to cache on a cache miss
 - **“hint”** only so implementation could support a subset or none of the levels and actions



Test yourself

Optional Reading: “Measuring Cache and TLB Performance and Their Effect on Benchmark Run Times,” Saavedra and Smith, 1995.

What cache is in your computer?

- How to figure out what cache configuration is in your computer
 - capacity (**C**), associativity (**a**), and block-size (**B**)
 - number of levels
- The presence or lack of a cache should not be detectable by functional behavior of software
- But you could tell if you measured execution time to infer the number of cache misses

Capacity Experiment: assume 2-power **C**

- For increasing **R**ange = 1,2,4,8,16,...
 - allocate a buffer of size **R**
 - repeatedly {read every byte in buffer in sequence}
 - measure average read time in steadystate
- Analysis
 - for small $R \leq C$, expect all reads to hit
 - for large $R > C$, expect reads to miss and detect corresponding jump in memory access time
- If continuing to increase **R**, read time jumps again when buffer size spills out to next cache level

Warning: timing won't be perfect when you try this

Block Size Experiment: knowing **C**

- Allocate a buffer of size **R** \gg **C**
- For increasing **S**=1,2,4,8....,
 - repeatedly {read every **S**'th byte in buffer in sequence}
 - measure average read time in steadystate
- Analysis
 - since **R** \gg **C**, expect first read to a block to miss when revisiting a block
 - reads to same block in same round should hit
 - expect increasing average read time for increasing **S** until **S** \geq **B** (no reuse in block)

Associativity Experiment: knowing **C**

- For increasing **R**, where **R** is a multiple of **C**
 - allocate a buffer of size **R**
 - repeatedly {read every **C**'th byte in buffer in sequence}
- Analysis
 - all **R/C** references map to the same set
 - for small **R** s.t. $(R/C) \leq a$, expect all reads to hit
 - for large **R** s.t. $(R/C) > a$, expect some reads to miss since touching more addresses than ways

note: 100% cache miss if LRU is used

How to detect associativity for lower-level caches?

Know your cache

- What else can you tell?
 - write-back vs write-through/write-allocate
 - unified vs. split design
 - I-cache C, B, a
 - t_i
 - replacement policy of associative caches
- Same mental exercise is required to control cache use in performance tuning

Caveat: experiments may not predict behaviors exactly for modern CPUs with virtual memory, complex hierarchies, and prefetchers