


# **18-447 Lecture 20: ILP to Multicores**

James C. Hoe  
Department of ECE  
Carnegie Mellon University

# Midterm 2 Class Distribution



MINIMUM

**6.0**

MEDIAN

**47.5**

MAXIMUM

**70.0**

MEAN

**45.26**

STD DEV ⓘ

**14.41**

# Midterm 2 Summary Statistics

|          | 1:TLB | 2:except | 3:energy | 4:poll | 5:draw | 6:trace | 7:table | total |
|----------|-------|----------|----------|--------|--------|---------|---------|-------|
| possible | 6     | 8        | 12       | 8      | 19     | 12      | 14      | 70    |
| average  | 4.8   | 3.3      | 5.8      | 4.6    | 8.2    | 9.3     | 9.3     | 45.3  |
| stdev    | 1.8   | 2.8      | 4.1      | 3.0    | 2.1    | 3.3     | 5.6     | 14.4  |
| max      | 6.0   | 8.0      | 12.0     | 8.0    | 10.0   | 12.0    | 14.0    | 70.0  |
| median   | 6.0   | 4.0      | 7.0      | 5.5    | 8.0    | 11.0    | 14.0    | 47.5  |
| min      | 0.0   | 0.0      | 0.0      | 0.0    | 0.0    | 0.0     | 0.0     | 6.0   |

# Housekeeping

- Your goal today
  - transition from sequential to parallel
  - enjoy (only first part, before OOO, on 447 exam)
- Notices
  - get going on Lab 4, now 3 weeks left
  - Handout #17: HW6, due Monday 5/2 noon
- Readings (advanced optional)
  - MIPS R10K Superscalar Microprocessor, Yeager
  - Synthesis Lectures: *Processor Microarchitecture: An Implementation Perspective*, 2010
  - Superscalar Club!!

# Parallelism Defined

- $T_1$  (work measured in time):
  - time to do work with 1 PE
- $T_\infty$  (critical path):
  - time to do work with infinite PEs
  - $T_\infty$  bounded by dataflow dependence
- Average parallelism:

$$P_{avg} = T_1 / T_\infty$$

*let's call  $p$   
concurrency*


- For a system with  $p$  PEs

$$T_p \geq \max\{T_1/p, T_\infty\}$$

- When  $P_{avg} \gg p$

$$T_p \approx T_1/p, \text{ aka "linear speedup"}$$

```
x = a + b;
y = b * 2
z = (x-y) * (x+y)
```



# ILP: Instruction-Level Parallelism

- Average **ILP** =  $T_1 / T_\infty$   
= no. instruction / no. cyc required

code1: **ILP** = 1

i.e., must execute serially

code2: **ILP** = 3

i.e., can execute at the same time

```
code1:  r1 ← r2 + 1
        r3 ← r1 / 17
        r4 ← r0 - r3
```


```
code2:  r1 ← r2 + 1
        r3 ← r9 / 17
        r4 ← r0 - r10
```

# Superscalar Speculative Out-of-Order Execution

# Exploiting **ILP** for Performance

Scalar in-order pipeline with forwarding

- operation latency (**OL**)= **1** base cycle
- peak **IPC** = **1**      *// no concurrency*
- require **ILP**  $\geq 1$  to avoid stall






# Superpipelined Execution

**OL** = **M** minor-cycle; same as **1** base cycle

peak **IPC** = **1** per minor-cycle // *has concurrency though*

required **ILP**  $\geq$  **M**




Achieving full performance requires always finding **M** “independent” instructions in a row

# Superscalar (Inorder) Execution

**OL** = 1 base cycle

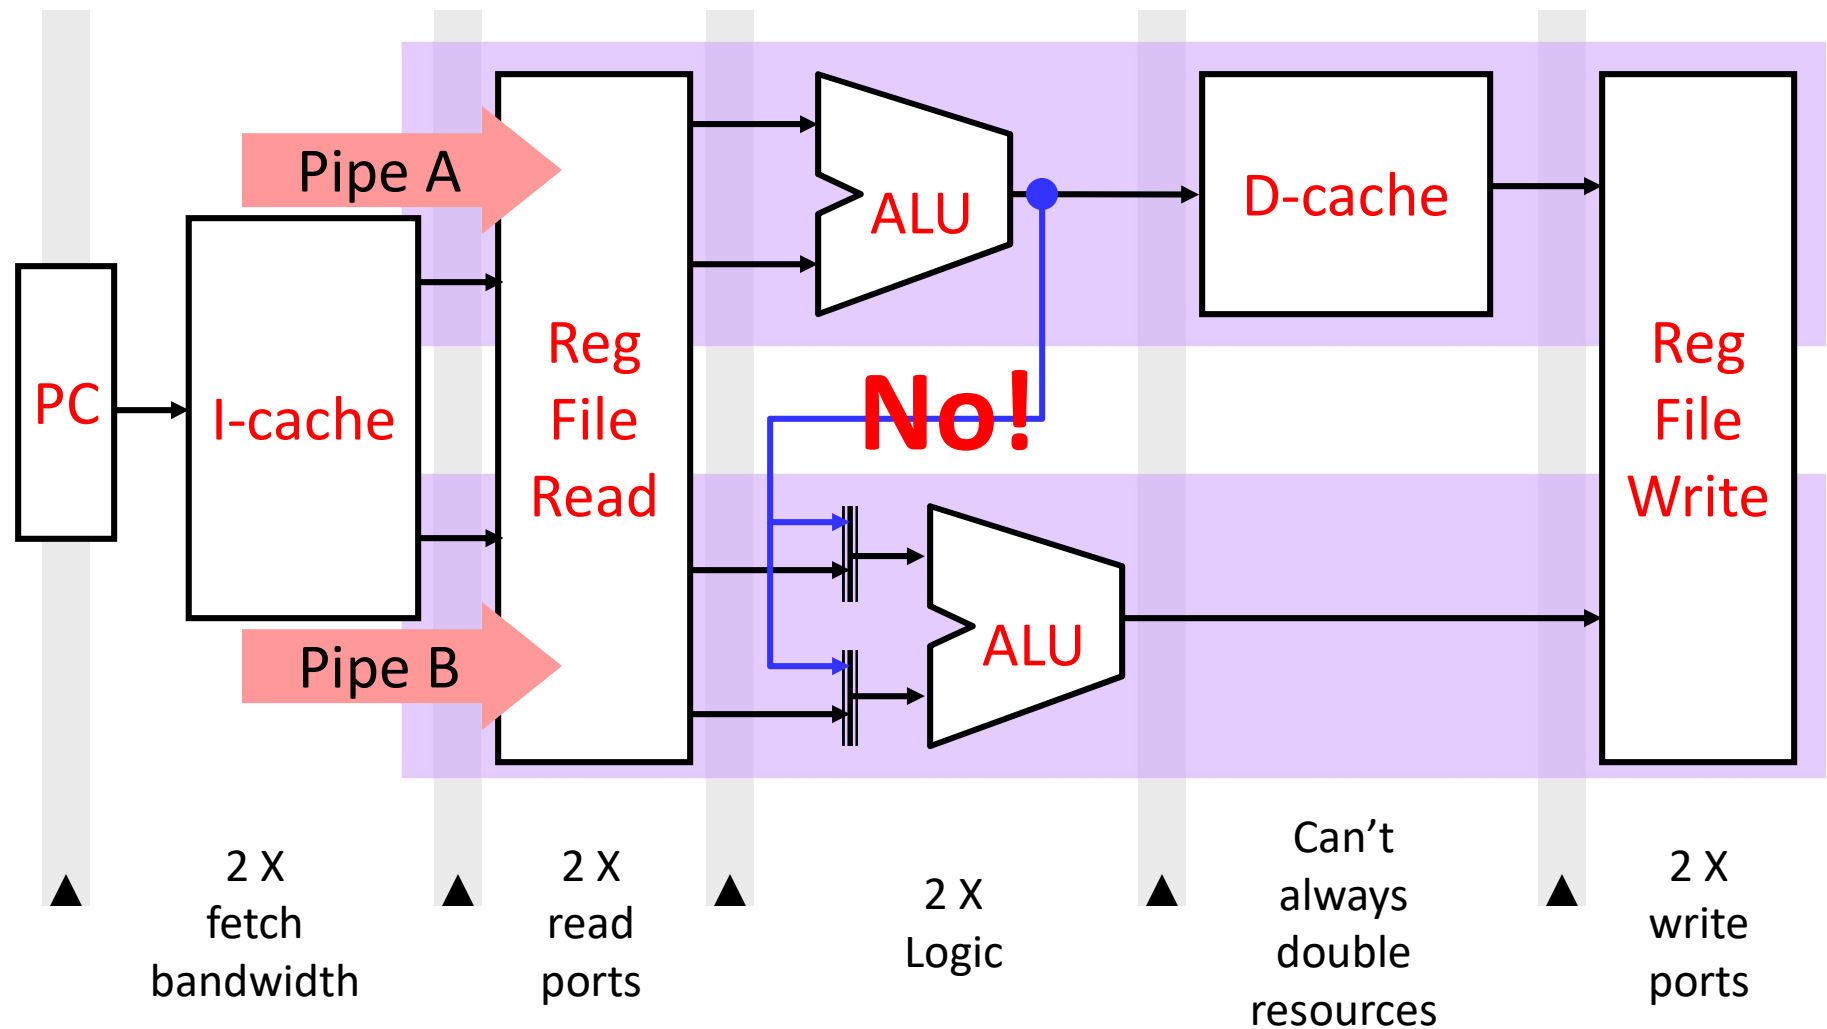
peak **IPC** = **N**

required **ILP**  $\geq$  **N**




Achieving full performance requires finding **N**  
"independent" instructions on every cycle

# Lab 4: 2-way, In-order Superscalar



# Limitations of Inorder Pipeline

- Achieved **IPC** of inorder pipelines degrades rapidly as **NxM** approaches **ILP**
- Despite high concurrency potential, pipeline never full due to frequent dependency stalls!!



# Out-of-Order Execution

- **ILP** is scope dependent

**ILP=1** {  
r1  $\leftarrow$  r2 + 1  
r3  $\leftarrow$  r1 / 17  
r4  $\leftarrow$  r0 - r3  
r11  $\leftarrow$  r12 + 1  
r13  $\leftarrow$  r19 / 17  
r14  $\leftarrow$  r0 - r20 } **ILP=2**

Accessing **ILP=2** requires not only (1) larger scheduling window but also (2) out-of-order execution


# Superscalar Speculative Out-of-Order Execution

# Instruction Micro-Dataflow

- Maintain a buffer of many pending instructions, a.k.a. reservation stations (**RSs**)
  - wait for functional unit to be free
  - wait for required input operands to be available
- Decouple execution order from who is first in line (program order)
  - select inst's in **RS** whose operands are available
  - give preference to older instructions (heuristic)
- A completing instruction (producer) signals dependent instructions (consumer) of operand availability (*producer-push resolution*)

# Tomasulo's Algorithm [IBM 360/91, 1967]

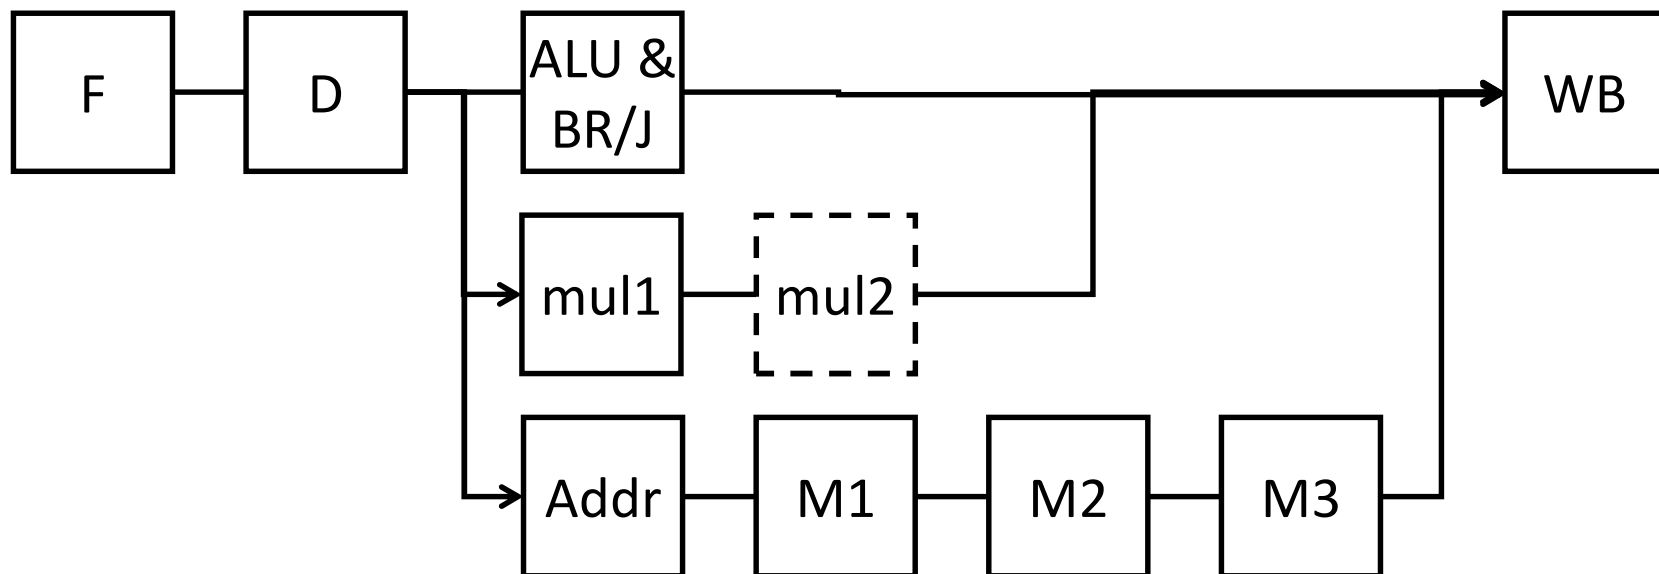
- Dispatch an instruction to a **RS** slot after decode
  - decode received from RF either operand value or placeholder **RS-tag**
  - mark RF dest with **RS-tag** of current inst's **RS** slot
- Inst in **RS** can issue when all operand values ready
- Completing instruction, in addition to updating RF dest, broadcast its **RS-tag** and value to all **RS** slots
- **RS** slot holding matching **RS-tag** placeholder pickup value






# WAW and WAR

- No WAW and WAR before because
  - single write stage
  - write stage at the end (*later than any read stage*)
  - in-order progression in pipeline



# Removing False Dependencies

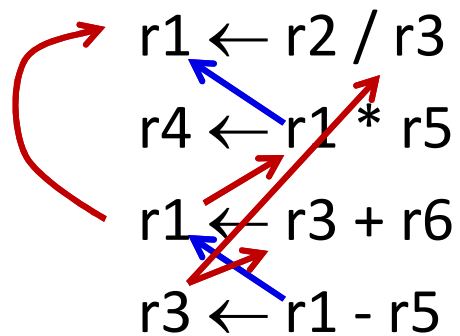
- With out-of-order execution comes WAW and WAR hazards
- Anti and output dependencies are false dependencies on register names rather than data



- With infinite number of registers, anti and output dependencies avoidable by using a new register for each new value

# Register Renaming: Example

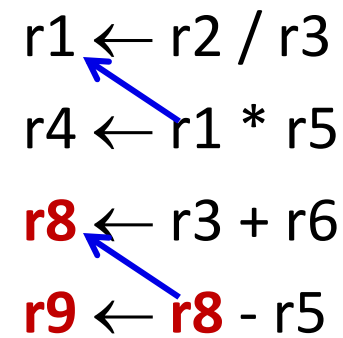
Original



$r1 \leftarrow r2 / r3$   
 $r4 \leftarrow r1 * r5$   
 $r1 \leftarrow r3 + r6$   
 $r3 \leftarrow r1 - r5$

The diagram shows four lines of code. Red arrows indicate the flow of the register 'r1' from the first line to the third line, and from the third line to the fourth line. Blue arrows indicate the flow of the register 'r3' from the first line to the second line, and from the second line to the fourth line.


Renamed



$r1 \leftarrow r2 / r3$   
 $r4 \leftarrow r1 * r5$   
 **$r8 \leftarrow r3 + r6$**   
 **$r9 \leftarrow r8 - r5$**

The diagram shows the same four lines of code as the original, but with renamed registers. The third and fourth lines are highlighted in red. Blue arrows show the flow of 'r3' from the first line to the second line, and from the second line to the third line. A blue arrow also shows the flow of 'r8' from the third line to the fourth line.

# On-the-fly HW Register Renaming



- Maintain mapping from ISA reg. names to physical registers
- When decoding an instruction that updates ' $r_x$ ':
  - allocate unused physical register  $t_y$  to hold inst result
  - set new mapping from ' $r_x$ ' to  $t_y$
  - younger instructions using ' $r_x$ ' as input finds  $t_y$
- De-allocate a physical register for reuse when it is never needed again?

^^^^^when is this exactly?

$r1 \leftarrow r2 / r3$

$r4 \leftarrow r1 * r5$

$r1 \leftarrow r3 + r6$

# Superscalar **Speculative** Out-of-Order Execution

# Control Speculation


- For want of a large window of instructions
    - if 14% of avg. instruction mix is control flow, what is average distance between control flow?
    - instruction fetch must make multiple levels of branch predictions (condition and target) to fetch far ahead of execution and commit
  - Modern CPUs can have over 100 instructions in out-of-order execution scope
- 
- **Question:**
    - **how much more ILP is uncovered with look ahead**
    - **how much useful work is done during look ahead**

**Ans: not much and not much**

# Speculative Out-of-order Execution


- A mispredicted branch after resolution must be rewound and restarted **ASAP!**
- Much trickier than 5-stage pipeline . . .
  - can rewind to an intermediate speculative state
  - a rewind branch could still be speculative and itself be discarded by another rewind!
  - rewind must reestablish both architectural state (register value) and microarchitecture state (e.g., rename table)
  - rewind/restart must be fast (not infrequent)
- Also need to rewind on exceptions . . . .but easier

# Nested Control Flow Speculation





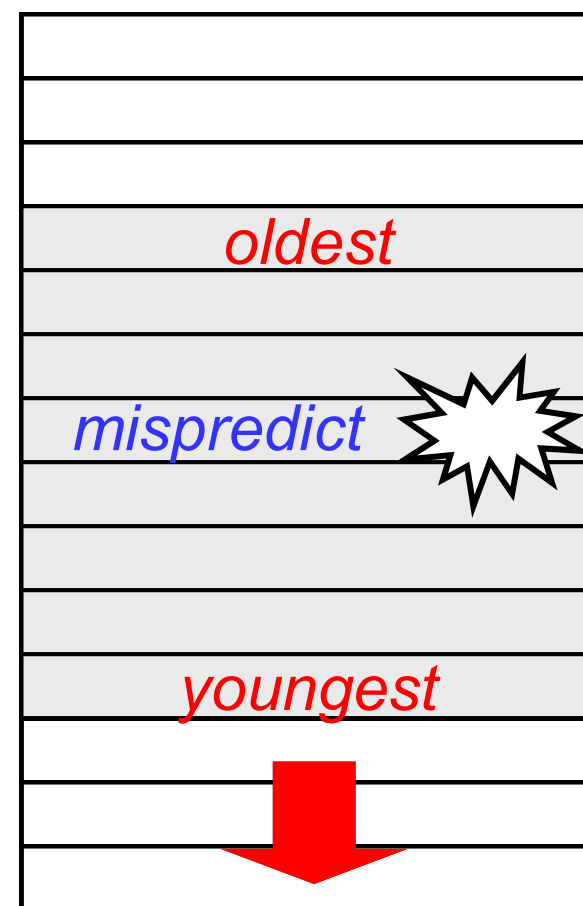
# Mis-speculation Recovery can be Speculative



# Instruction Reorder Buffer (**ROB**)

- Program-order bookkeeping (circular buffer)
  - instructions enter and leave in program order
  - tracks 10s to 100s of in-flight instructions in different stages of execution

- Dynamic juggling of state and dependency
  - oldest finished instruction “commit” architectural state updates on exit
  - all ROB entries considered “speculative” due to potential for exceptions and mispredictions



# In-order vs Speculative State

- In-order state:
  - cumulative architectural effects of all instructions committed in-order so far
  - can never be undone!!
- Speculative state, as viewed by a given inst in **ROB**
  - in-order state + effects of older inst's in **ROB**
  - effects of some older inst's may be pending
- Speculative state effects must be reversible
  - remember both in-order and speculative values for an RF register (may have multiple speculative values)
  - store inst updates memory only at commit time
- Discard younger speculative state to rewind execution to oldest remaining inst in **ROB**


# You have seen this in your pipeline

privileged mode

|     | t <sub>0</sub> | t <sub>1</sub> | t <sub>2</sub> | t <sub>3</sub> | t <sub>4</sub> | t <sub>5</sub> | t <sub>6</sub> | t <sub>7</sub> | t <sub>8</sub> | t <sub>9</sub>   | t <sub>10</sub>  |
|-----|----------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|------------------|------------------|
| IF  | I <sub>0</sub> | I <sub>1</sub> | I <sub>2</sub> | I <sub>3</sub> | I <sub>4</sub> | bub            | bub            | bub            | I <sub>h</sub> | I <sub>h+1</sub> | I <sub>h+2</sub> |
| ID  |                | I <sub>0</sub> | I <sub>1</sub> | I <sub>2</sub> | I <sub>3</sub> | bub            | bub            | bub            | bub            | I <sub>h</sub>   | I <sub>h+1</sub> |
| EX  |                |                | I <sub>0</sub> | I <sub>1</sub> | I <sub>2</sub> | I <sub>3</sub> | bub            | bub            | bub            | bub              | I <sub>h</sub>   |
| MEM |                |                |                | I <sub>0</sub> | I <sub>1</sub> | I <sub>2</sub> | bub            | bub            | bub            | bub              | bub              |
| WB  |                |                |                |                | I <sub>0</sub> | I <sub>1</sub> | I <sub>2</sub> | bub            | bub            | bub              | bub              |

- Flush the pipeline at anytime to start from after last instruction committed
- Decode in ID needs to consider pending “speculative” state updates in EX/MEM/WB, on top of committed RF states

# Superscalar Speculative OOO All Together



# Truth about Superscalar Speculative OOO

- If memory speed kept up with core speed, we would still be building in-order pipelines
- But, by 2005 we were seeing e.g., Intel P4 at 4+GHz
  - 16KB L1 D-cache
    - $t_1 = 4$  cyc int (9 cycle fp)
  - 1024KB L2 D-cache
    - $t_2 = 18$  cyc int (18 cyc fp)
  - Main memory
    - $t_3 = \sim 50$ ns or 180 cyc
- Speculative OOO has really been about
  - finding independent work to do after cache hit&miss
  - getting to future cache misses as early as possible
  - overlapping multiple cache misses for BW

# At the 2005 Peak of Superscalar OOO

|                       | Alpha<br>21364 | AMD<br>Opteron | Intel<br>Xeon | IBM<br>Power5 | MIPS<br>R14000 | Intel<br>Itanium2 |
|-----------------------|----------------|----------------|---------------|---------------|----------------|-------------------|
| clock (GHz)           | 1.30           | 2.4            | <b>3.6</b>    | 1.9           | 0.6            | 1.6               |
| issue rate            | 4              | 3 (x86)        | 3 (rop)       | <b>8</b>      | 4              | 8                 |
| pipeline int/fp       | 7/9            | 9/11           | <b>22/24</b>  | 12/17         | 6              | 8                 |
| inst in flight        | 80             | 72(rop)        | 126 rop       | <b>200</b>    | 48             | inorder           |
| rename reg            | 48+41          | 36+36          | 128           | 48/40         | 32/32          | <b>328</b>        |
| transistor ( $10^6$ ) | 135            | 106            | 125           | 276           | 7.2            | <b>592</b>        |
| power (W)             | <b>155</b>     | 86             | 103           | 120           | 16             | 130               |
| SPECint 2000          | 904            | 1,566          | 1,521         | 1,398         | 483            | <b>1,590</b>      |
| SPECfp 2000           | 1279           | 1,591          | 1,504         | 2,576         | 499            | <b>2,712</b>      |

# At peak minus 5 years

|                       | Alpha<br>21264 | AMD<br>Athlon | Intel<br>P4    | MIPS<br>R12000 | IBM<br>Power3 | HP<br>PA8600 | SUN<br>Ultra3 |
|-----------------------|----------------|---------------|----------------|----------------|---------------|--------------|---------------|
| clock (MHz)           | 833            | 1200          | <b>1500</b>    | 400            | 450           | 552          | 900           |
| issue rate            | 4              | 3 (x86)       | 3 (rop)        | 4              | 4             | 4            | 4             |
| pipeline int/fp       | 7/9            | 9/11          | <b>22/24</b>   | 6              | 7/8           | 7/9          | 14//15        |
| inst in flight        | 80             | 72(rop)       | <b>126</b> rop | 48             | 32            | 56           | inorder       |
| rename reg            | 48+41          | 36+36         | <b>128</b>     | 32+32          | 16+24         | 56           | inorder       |
| transistor ( $10^6$ ) | 15.4           | 37            | 42             | 7.2            | 23            | <b>130</b>   | 29            |
| power (W)             | 75             | <b>76</b>     | 55             | 25             | 36            | 60           | 65            |
| SPECint 2000          | 518            |               | <b>524</b>     | 320            | 286           | 417          | 438           |
| SPECfp 2000           | <b>590</b>     | 304           | 549            | 319            | 356           | 400          | 427           |




# Performance (In)efficiency

- To hit “expected” performance target
  - push frequency harder by deepening pipelines
  - used the 2x transistors to build more complicated microarchitectures so fast/deep pipelines don’t stall (i.e., caches, BP, superscalar, out-of-order)
- The consequence of performance inefficiency is

limit of  
economical  
cooling [ITRS]


Recall



2005, Intel  
P4 Tehas 150W

Figure 8. Power dissipation projections. [Borkar, IEEE Micro, July 1999]

# Efficiency of Parallel Processing



Recall  
Performance per Instruction Trends in Intel®  
Microprocessors, Grochowski et al., 2006]

# At peak plus 1 year

|                          | AMD<br>285  | Intel<br>5160        | Intel<br>965 | Intel<br>Itanium2 | IBM<br>P5+  | MIPS<br>R16000 | SUN<br>Ultra4 |
|--------------------------|-------------|----------------------|--------------|-------------------|-------------|----------------|---------------|
| cores/threads            | <b>2x1</b>  | <b>2x2</b>           | <b>2x2</b>   | <b>2x2</b>        | <b>2x2</b>  | 1x1            | <b>2x1</b>    |
| clock (GHz)              | 2.6         | 3.03                 | <b>3.73</b>  | 1.6               | 2.3         | 0.7            | 1.8           |
| issue rate               | 3 (x86)     | 4 (rop)              | 3 (rop)      | 6                 | <b>8</b>    | 4              | 4             |
| pipeline depth           | 11          | 14                   | <b>31</b>    | 8                 | 17          | 6              | 14            |
| inst in flight           | 72(rop)     | 96(rop)              | 126(rop)     | inorder           | <b>200</b>  | 48             | inorder       |
| on-chip\$ (MB)           | 2x1         | 4                    | 2x2          | <b>2x13</b>       | 1.9         | 0.064          | 2             |
| transistor ( $10^6$ )    | 233         | 291                  | 376          | <b>1700</b>       | 276         | 7.2            | 295           |
| power (W)                | 95          | 80                   | <b>130</b>   | 104               | 100         | 17             | 90            |
| SPECint 2000<br>per core | <b>1942</b> | (1556 <sup>+</sup> ) | 1870         | 1474              | 1820        | 560            | 1300          |
| SPECfp 2000<br>per core  | 2260        | (1694 <sup>+</sup> ) | 2232         | 3017              | <b>3369</b> | 580            | 1800          |

\*3086/+2884 according to [www.spec.org](http://www.spec.org)


18-447-S22-L20-S35, James C. Hoe, CMU/ECE/CALCM, ©2022

Microprocessor Report, Aug 2006

# At peak plus 3 years

|                                | AMD<br>Opteron<br>8360SE | Intel<br>Xeon<br>X7460 | Intel<br>Itanium<br>9050 | IBM<br>P5  | IBM<br>P6 | Fijitsu<br>SPARC 7 | SUN<br>T2  |
|--------------------------------|--------------------------|------------------------|--------------------------|------------|-----------|--------------------|------------|
| cores/threads                  | 4x1                      | <b>6x1</b>             | 2x2                      | 2x2        | 2x2       | 4x2                | <b>8x8</b> |
| clock (GHz)                    | 2.5                      | 2.67                   | 1.60                     | 2.2        | <b>5</b>  | 2.52               | 1.8        |
| issue rate                     | 3 (x86)                  | 4 (rop)                | 6                        | 5          | <b>7</b>  | 4                  | 2          |
| pipeline depth                 | 12/ <b>17</b>            | 14                     | 8                        | 15         | 13        | 15                 | 8/12       |
| out-of-order                   | 72(rop)                  | 96(rop)                | inorder                  | <b>200</b> | limited   | 64                 | inorder    |
| on-chip\$ (MB)                 | 2+2                      | <b>9+16</b>            | 1+12                     | 1.92       | 8         | 6                  | 4          |
| transistor ( $10^6$ )          | 463                      | <b>1900</b>            | 1720                     | 276        | 790       | 600                | 503        |
| power max(W)                   | 105                      | 130                    | 104                      | 100        | >100      | <b>135</b>         | 95         |
| SPECint 2006<br>per-core/total | 14.4/170                 | <b>22/274</b>          | 14.5/1534                | 10.5/197   | 15.8/1837 | 10.5/ <b>2088</b>  | --/142     |
| SPECfp 2006<br>per-core/total  | 18.5/156                 | 22/142                 | 17.3/1671                | 12.9/229   | 20.1/1822 | <b>25.0/1861</b>   | --/111     |

# On to Mainstream Parallelism in Multicores and Manycores



Remember, we got here because we need to compute  
faster while using less energy per operation