

18-447 Lecture 22:

1 Lecture Worth of Parallel Programming Primer

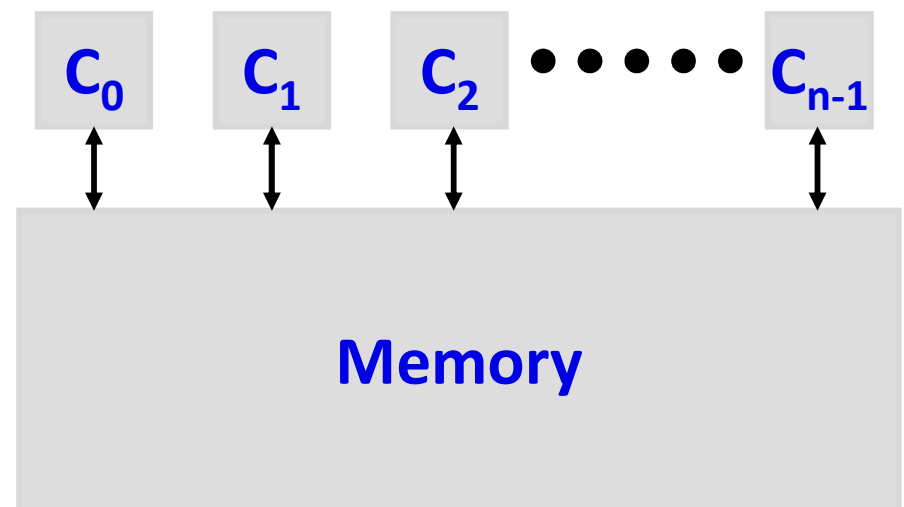
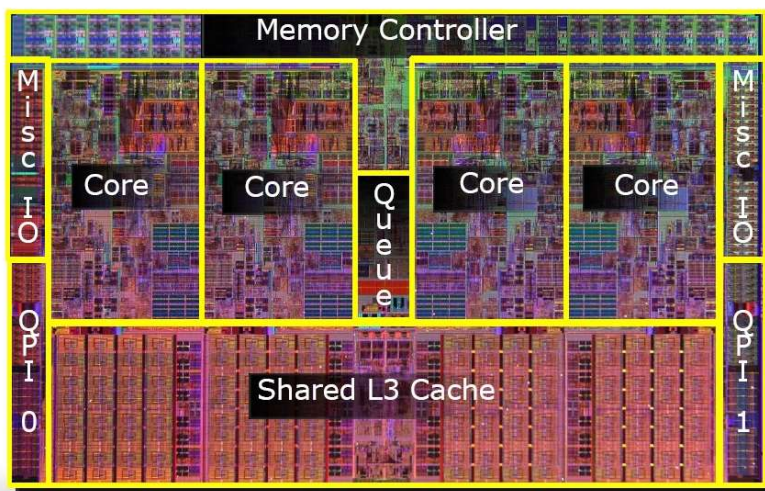
James C. Hoe
Department of ECE
Carnegie Mellon University

Housekeeping

- Your goal today
 - see basic concepts in shared-memory multithreading (context for topics to come)
 - appreciate how easy parallel programming can be
 - appreciate how difficult “good” parallel programming can be
- Notices
 - Lab 4, due week 14
 - HW6, due Monday 5/2 noon
 - Midterm 2 Regrade, due Monday, 4/25
- Readings
 - P&H Ch 6

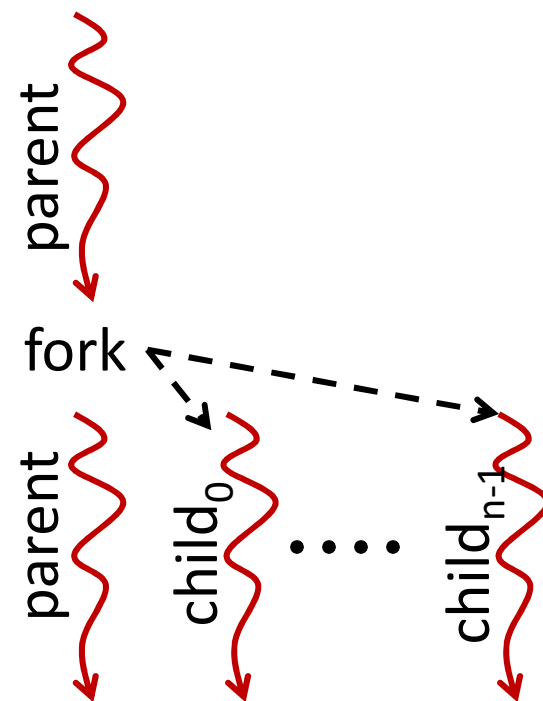
Shared-Memory Multicores

- Today's general-purpose multicore processors are MIMD, symmetric, shared memory
 - individual cores follow classic von Neuman
 - common access to physical address space and mem
 - processes/threads on different cores communicate by writing and reading agreed-upon mem locations



Single Program Multiple Data

- SPMD is MIMD except all threads based on the same program image
- On SMP, SPMD starts as a single-thread process and its memory
- Independent “threads of execution” (think program counters, regfile and stacks) spawned
 - ****same process memory****—same **EA** in different threads refers to shared program and data locations
 - different threads run concurrently (on different cores) or interleaved



SPMD just one of many options; prevalent and easy to start on

E.g., POSIX Threads Create and Join

```

long count=0;           // globals are in memory and shared!!

void *foo(void *arg) { return count = count + (long)arg; }

int main(){
    pthread_t tid[HOWMANY];           // array of thread IDs
    long i;
    void *retval;

    // spawn children threads
    for(i=0; i<HOWMANY; i++ )
        pthread_create( &tid[i],           // ID to be set
                        NULL,               // attribute (default)
                        foo,                // fxn to run by thread
                        (void*)i);          // ptr-size arg to fxn

    // wait for children threads to exit
    for (i=0; i<HOWMANY; i++ )
        pthread_join( tid[i],              // ID to wait on
                      &retval);            // ptr-size return value
}

```

Memory Consistency

- Memory consistency model says for each read which write bound the value to be returned
 - intuitively: a read should return value of “most recent” write to the same address
 - straight forward for a single thread
- In a shared-memory multicore, cores **C1/C2/C3** perform following streams of reads and writes

C1: W(**x**)

C2: W(**x**), W(**x**), W(**y**), R(**x**), R(**y**) . . .

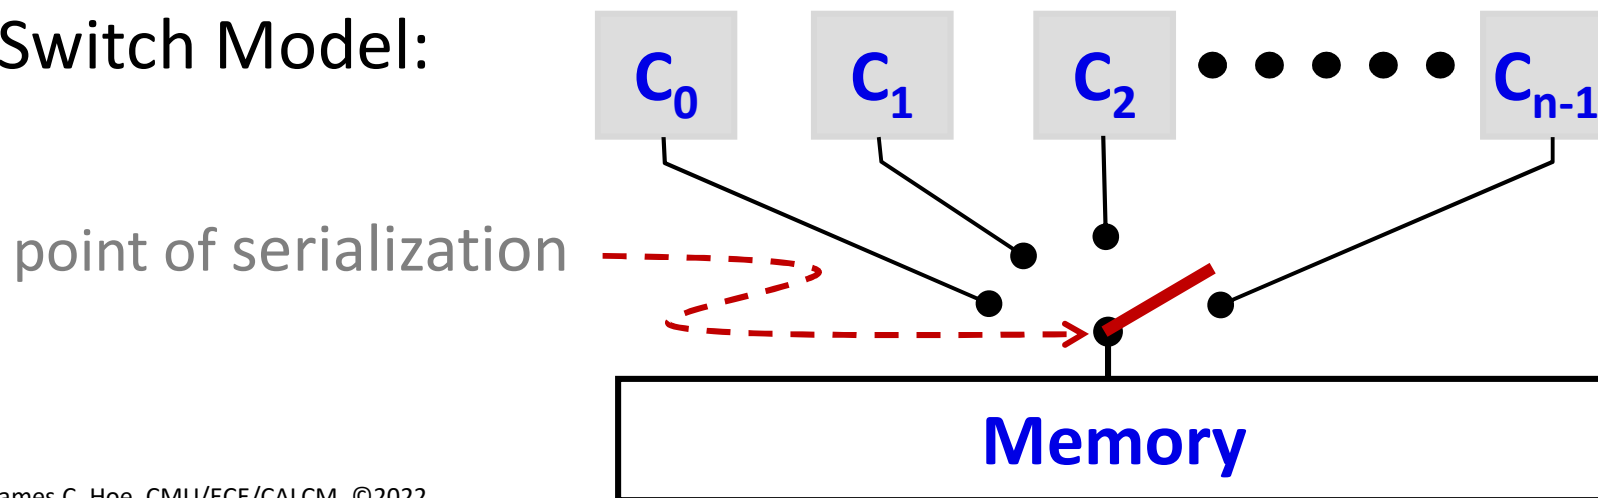
C3: W(**x**), W(**y**), W(**x**) . . .

Which is the last write to **x** before R(**x**) by **C2**?

Ordering determines what can be seen by reads, but
what is observed by reads determines ordering!!

Sequential Consistency (SC)

- A thread perceives its own memory ops in program order (of course)
- Memory ops from threads in program order can be interleaved arbitrarily; different interleaving allowed on different runs, i.e., nondeterminism
- For each run, all threads must not disagree on any orderings observed
- Switch Model:



SC Example: what can and cannot be

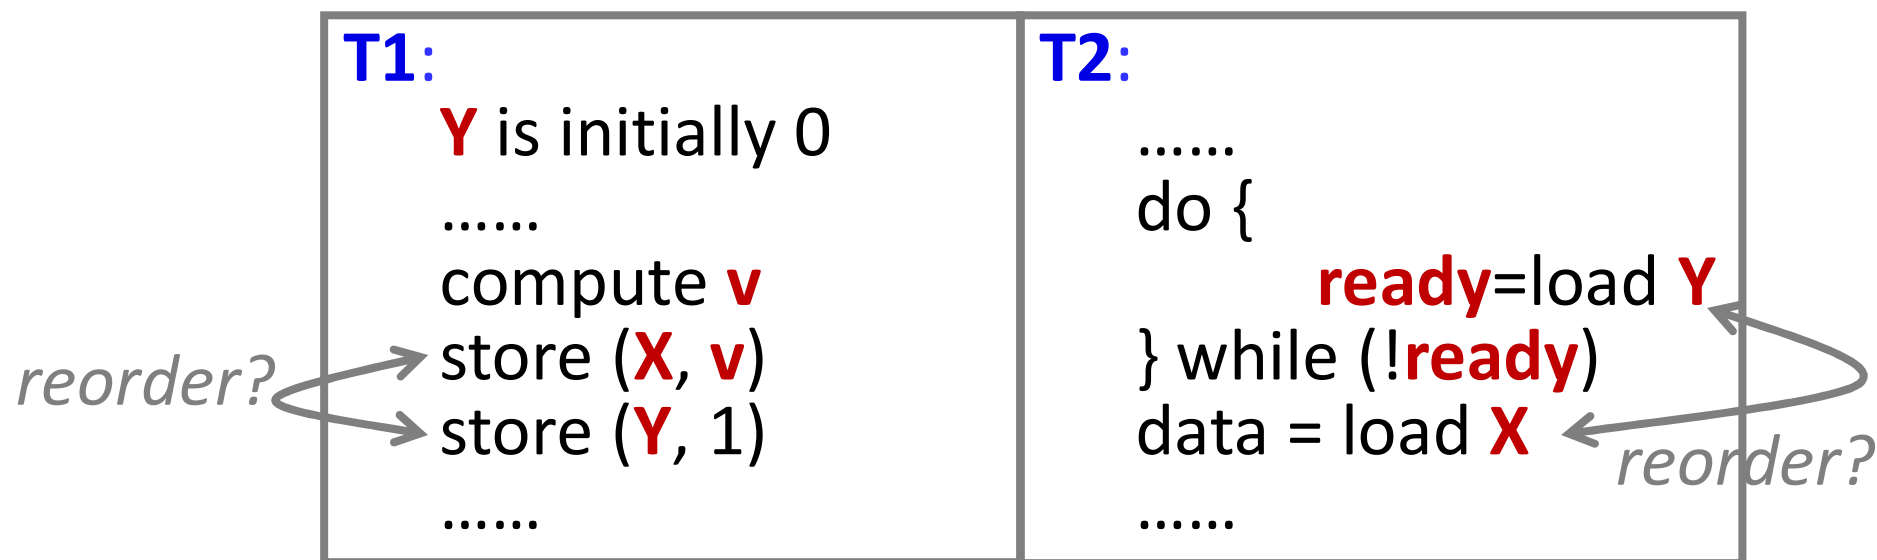
- Threads **T1** and **T2** and shared locations **X** and **Y** (initially **X** = 0, **Y** = 0)

T1: store(X , 1); store(Y , 1);	T2: vy = load(Y); vx = load(X);
------------	--	------------	--

- SC says
 - **vy** and **vx** may get different values from run to run
e.g., (**vy**=0, **vx**=0), (**vy**=0, **vx**=1), or (**vy**=1, **vx**=1)
 - but if **vy** is 1 then **vx** cannot be 0

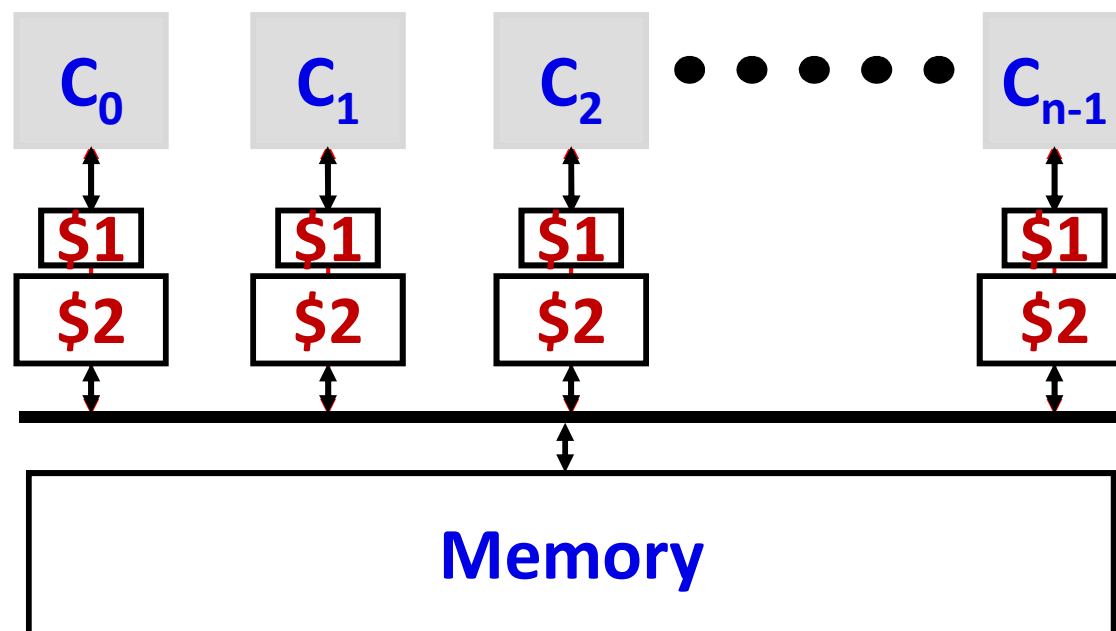
An Useful Example

- Threads **T1** and **T2** communicate via shared memory locations **X** and **Y**
 - **T1** produces result in **X** to be consumed by **T2**
 - **T1** signals readiness to **T2** by setting **Y**



- This works because SC says **T1** and **T2** must see the stores to **X** and **Y** in the same order

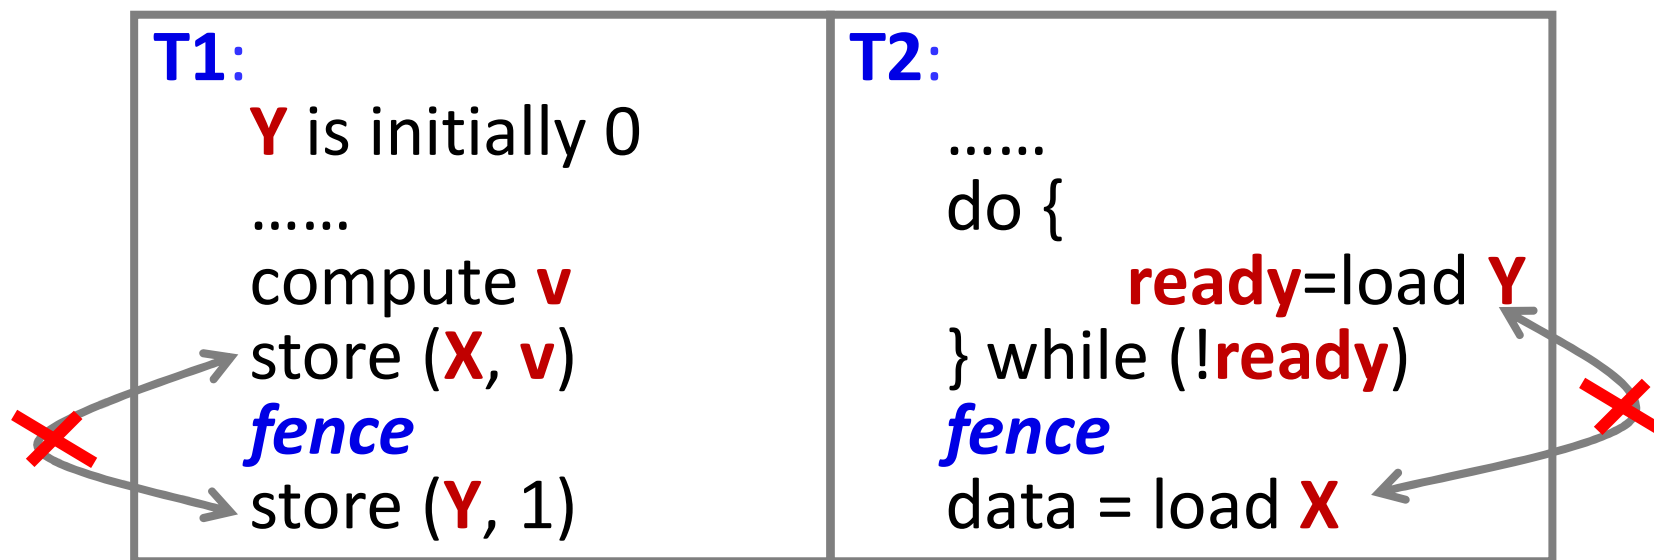
Easy to think about hard to build



- Where is “point of serialization” if memory ops don’t always go to memory or even onto a bus?
- SC restricts many memory reordering optimizations *taken-for-granted* in sequential execution (*e.g., non-blocking miss*)

Weak Consistency (WC)

- WC imposes only uniprocessor memory ordering requirements: $R(x) < W(x)$; $W(x) < R(x)$; $W(x) < W(x)$
- Program inserts explicit memory fence instructions to force serialization when it matters




- If serialization is rare, cheap(hw)/slow fences okay, e.g., completely drain/restart pipeline

Intermediate models exist between SC and WC

Embarrassingly Parallel Processing

- Summing 10,000 numbers from array **A** []
- In sequential algorithm

```
for (i=0; i<10000; i=i+1)
    sum = sum "+" A[i];
```



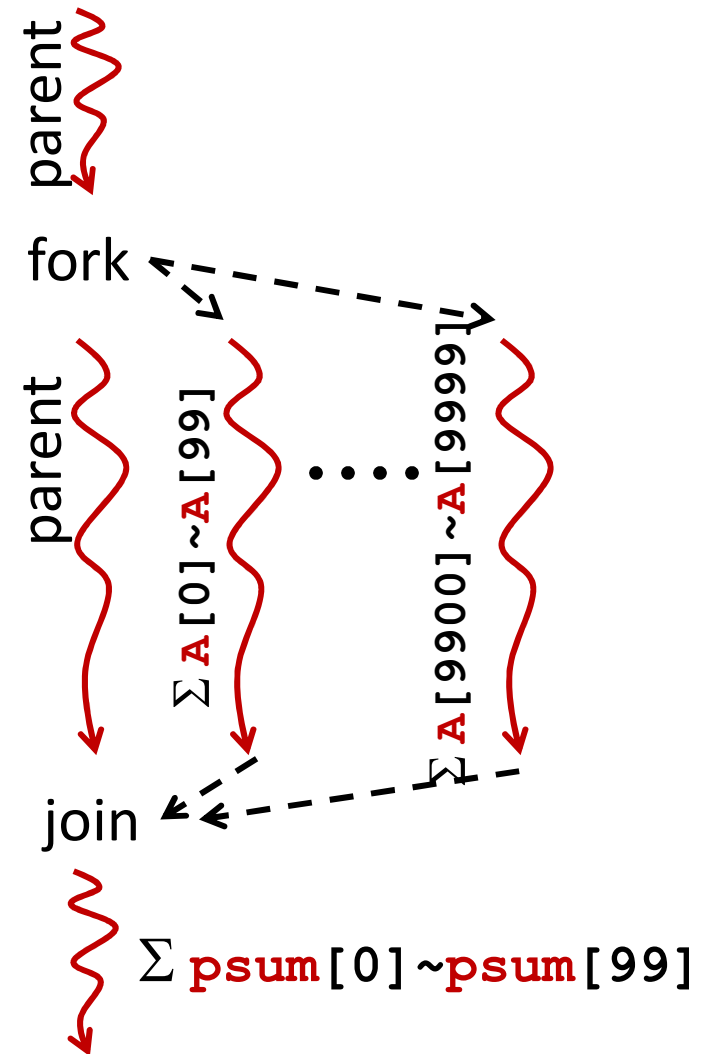
- Assuming “+” is 1 unit-time; everything else free
 - $T_1 = 10,000$
 - $T_\infty = \lceil \log_2 10,000 \rceil = 14$ (*using associativity of “+”*)
 - $P_{avg} = T_1 / T_\infty = 714$
- Ideally, at $p = 100 \ll T_1 / T_\infty$
 expect $T_{100} \approx T_1 / p = 100$ or $S_{100} \approx p = 100$

recall if $T_1 / T_\infty \gg p$ then $S \approx p$

Note **P** vs **p**

Shared-Memory Pthreads Strategy 1

- Fork $p=100$ threads on a p -way shared memory multiprocessor
 - $A[10000]$ is in shared memory
 - $psum[100]$ is also in shared memory
- Child thread- i uses $psum[i]$ to compute its portion of the partial sum
- When all threads finish, parent sums $psum[0] \sim psum[99]$



Children Thread Code

```
double A[ARRAY_SIZE];
double psum[p];

void *sumParallel(void *_id) {
    long id=(long) _id;
    long i;

    psum[id]=0;

    for(i=0;i<(ARRAY_SIZE/p);i++)
        psum[id]+=A[id*(ARRAY_SIZE/p) + i];

    return NULL;
}
```

Parent Code

```
double A[ARRAY_SIZE];
double psum[p];
double sum=0;

int main() {

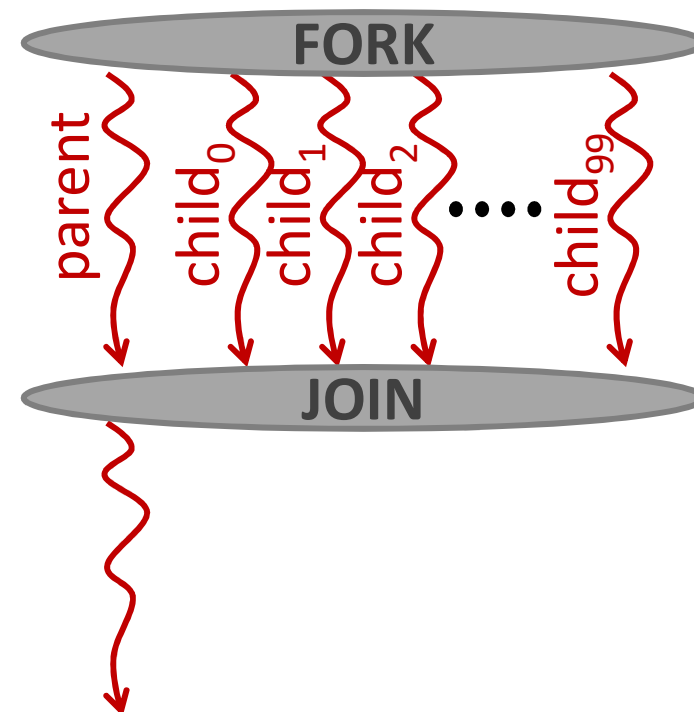
    ... skipped pthreads boilerplate ...

    for(i=0; i<p; i++ )
        pthread_create( &tid[i],
                        NULL,
                        sumParallel,
                        (void*)i );

    for (i=0; i<p; i++ ) {
        pthread_join( tid[i], &retval);
        sum+=psum[i];
    }
}
```

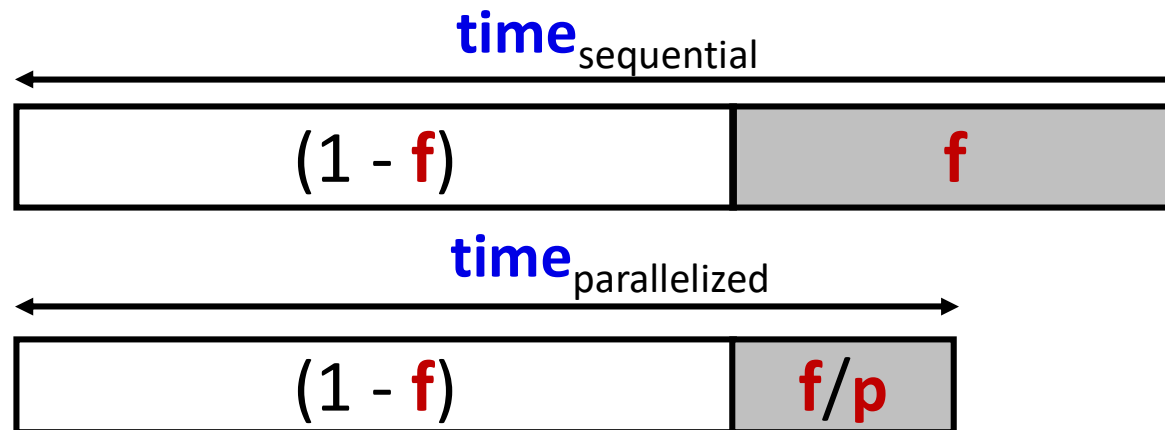
Performance Analysis

- Summing 10,000 numbers on 100 cores
 - 100 threads performs 100 +’s each in parallel
 - parent thread performs 100 +’s sequentially
 - $T_{100} = 100 + 100$
 - $S_{100} = 50$
- If 100,000 num on 100 cores
 - $T_{100} = 1000 + 100$
 - $S_{100} = 90.9$
- If 10,000 num on 10 cores
 - $T_{10} = 1000 + 10$
 - $S_{10} = 9.9$
- Don’t forget,
 - *fork* and *join* are not free
 - moving data (even thru shared memory) not free



The Actual Amdahl's Law

- If only a fraction **f** (by time) is parallelizable by **p**



$$\text{time}_{\text{parallelized}} = \text{time}_{\text{sequential}} \cdot ((1-f) + f/p)$$

$$S_{\text{effective}} = 1 / ((1-f) + f/p)$$

- if **f** is small, **p** doesn't matter
- even when **f** is large, diminishing return on **p**;
eventually “1-**f**” dominates

Strategy 2: parallelizing the reduction

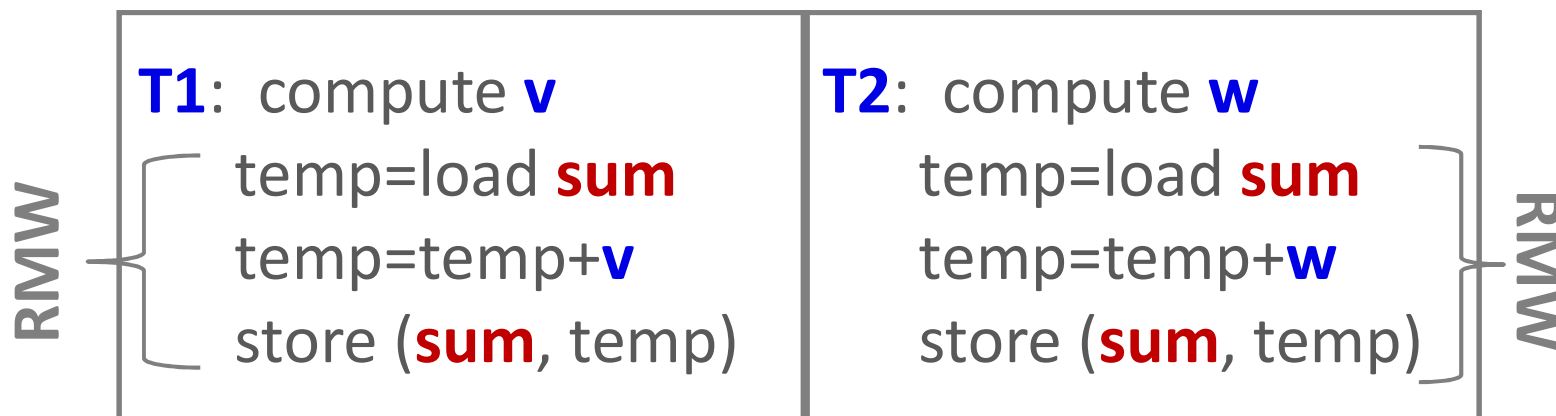
- How about asking each thread to do a bit of the reduction, i.e.,

```
void *sumParallel(void *_id) {  
    long id=(long) _id;  
    long i;  
  
    psum[id]=0;  
  
    for (i=0; i<(ARRAY_SIZE/p); i++)  
        psum[id]+=A[id*ARRAY_SIZE/p+i];  
  
    sum=sum+psum[id];  
  
    return NULL;  
}
```

Assume SC for simplicity

Data Races

- On last slide **sum** is read and updated by all threads at around the same time
- Let's try just 2 threads T1 and T2, **sum** is initially 0



- What are the possible final values of **sum**?
 - **v+w** or **v** or **w** depending on the interleaving of the read/modify/write sequence in **T1** and **T2**
- To work, RMW regions needs to be atomic
i.e., no intervening reads/writes by other threads

Critical Sections

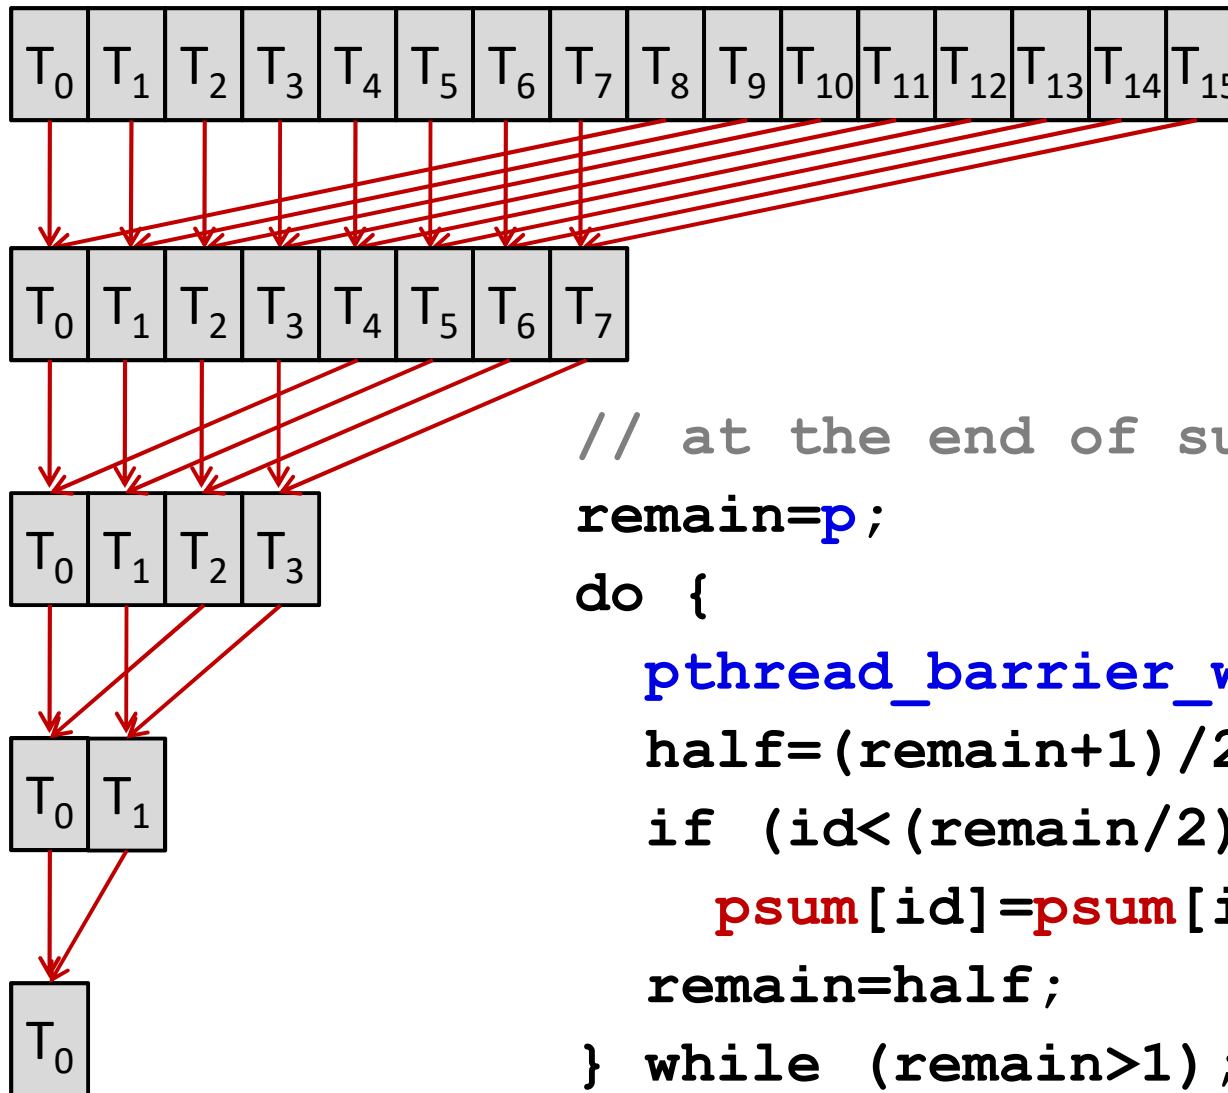
- Special “lock” variables and lock/unlock operators to demarcate a “critical section” that only one thread can enter at a time, e.g.,

```
pthread_mutex_lock(&lockvar);  
sum=sum+psum[id];          // atomic RMW  
pthread_mutex_unlock(&lockvar);
```

- **lock()** blocks until **lockvar** is free or freed (released by previous owner)
- on **unlock()**, if multiple **lock()** pending, only 1 should succeed; the rest keep waiting
- Strategy 2 is now correct but actually slower

*Reduction still sequential plus
extra cost of locking and unlocking*

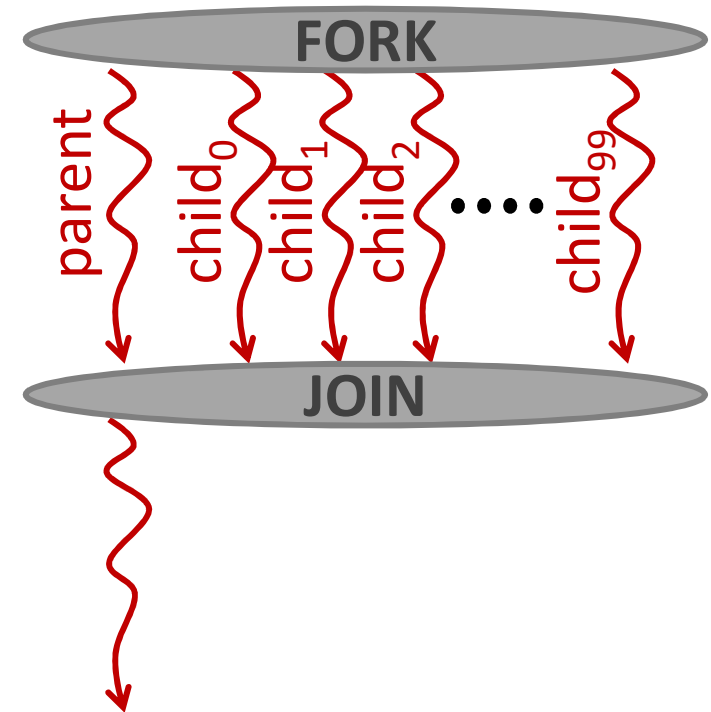
Strategy 3: Parallel Reduction (assume “+” associative and commutative)



```
// at the end of sumParallel()
remain=p;
do {
    pthread_barrier_wait(&barrier);
    half=(remain+1)/2;
    if (id<(remain/2))
        psum[id]=psum[id]+psum[id+half];
    remain=half;
} while (remain>1);
```

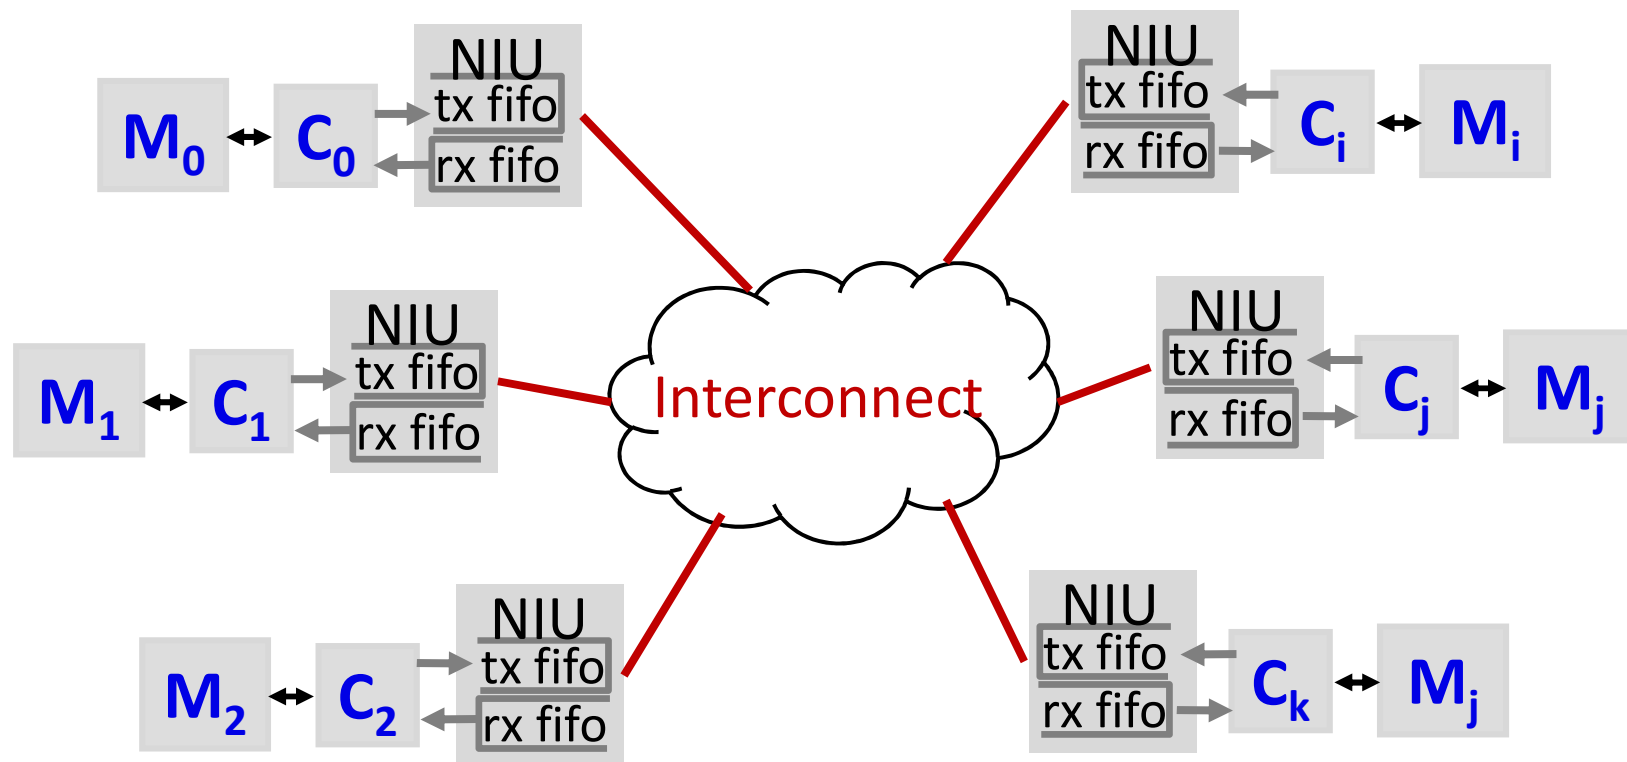
Performance Analysis

- Summing 10,000 on 100 cores
 - 100 threads performs 100 +’s each in parallel, and
 - between 1~7 +’s each in the parallel reduction
 - $T_{100} = 100 + 7$
 - $S_{100} = 93.5$
- If summing 100,000 on 100 cores
 - $T_{100} = 1000 + 7$
 - $S_{100} = 99.3$
- If summing 10,000 on 10 cores
 - $T_{10} = 1000 + 4$
 - $S_{10} = 10.0$



First-order analysis! Don't bet on this.

Message Passing



- Private address space and memory per processor
- Parallel threads on different processors communicate by explicit sending and receiving of messages

Review

Example using Matched Send/Receive

```

if (id==0)                //assume node-0 has A initially
    for (i=1;i<p;i=i+1)
        SEND(i, &A[SHARE*i], SHARE*sizeof(double));
else
    RECEIVE(0,A[])        //receive into local array

sum=0;
for(i=0;i<SHARE;i=i+1) sum=sum+A[i];

remain=p;
do {
    BARRIER();
    half=(remain+1)/2;
    if (id>=half&&id<remain) SEND(id-half,sum,8);
    if (id<(remain/2)) {
        RECEIVE(id+half,&temp);
        sum=sum+temp;
    }
    remain=half;
} while (remain>1);

```

SHARE=HOWMANY/p

[based on P&H Ch 6 example]

Communication Cost

- Communication cost is a part of parallel execution
- Easier to perceive communication cost in message passing
 - overhead: takes time to send and receive data
 - latency: takes time for data to go from A to B
 - gap ($1/\text{bandwidth}$): takes time to push successive data through a finite bandwidth
- Same cost was also there in shared memory

To be continued